



Arrays

- The core data structure in `numpy` is the array. Moving forward, "array" will always refer to a `numpy` array.
- One way to instantiate an array is to pass a list as an argument to the function `np.array`.

```
In [14]: 1 np.array(4, 9, 1, 2)
```

TypeError

Cell In[14], line 1

----> 1 np.array(4, 9, 1, 2)

Traceback (most recent call last)

TypeError: array() takes from 1 to 2 positional arguments but 4 were given

most important
parts are
at the
bottom!

```
In [13]: 1 np.array([4, 9, 1, 2])
```

```
Out[13]: array([4, 9, 1, 2])
```





```
Out[13]: array([4, 9, 1, 2])
```

- Arrays, unlike lists, must be **homogenous** – all elements must be of the

```
In [16]: 1 np.array(['hello', 'hi', 'michigan'])
```

```
Out[16]: array(['hello', 'hi', 'michigan'], dtype='<U8')
```

strings up to
length 8

```
In [17]: 1 np.array(['hi'])
```

```
Out[17]: array(['hi'], dtype='<U2')
```

strings up to
length 2

```
In [15]: 1 # All elements are converted to strings!
2 np.array(['hello', 2025])
```

```
Out[15]: array(['hello', '2025'], dtype='<U21')
```

also often say these operations are **vectorized**.

$$\begin{array}{c|c} \begin{array}{c} 1 \\ 2 \end{array} & * \textcolor{purple}{1.6} = \end{array} \quad \begin{array}{c|c} \begin{array}{c} 1 \\ 2 \end{array} & * \begin{array}{c} 1.6 \\ 1.6 \end{array} = \end{array} \quad \begin{array}{c|c} \begin{array}{c} 1.6 \\ 3.2 \end{array} & \end{array}$$

```
In [18]: 1 temps = [68, 72, 65, 64, 62, 61, 59, 64, 64, 63, 65, 62]
          2 temps
```

Out[18]: [68, 72, 65, 64, 62, 61, 59, 64, 64, 63, 65, 62]

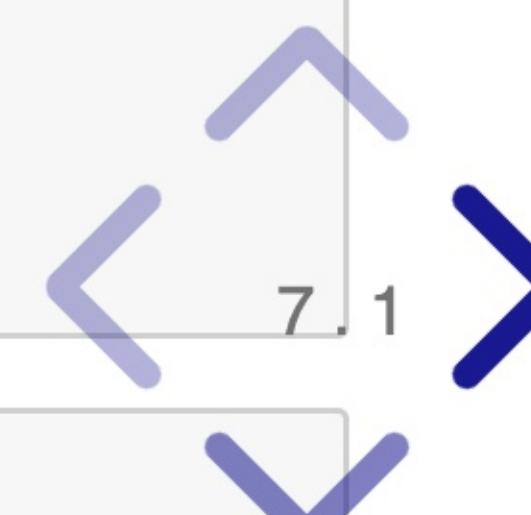
```
In [ ]: 1 temp_array = np.array(temps) no [ ] !
```

```
In [ ]: 1 # Increase all temperatures by 3 degrees.
          2 ...
```

temps is
already a
list :

```
In [ ]: 1 # Halve all temperatures.
          2 ...
```

```
In [ ]: 1 # Convert all temperatures to Celsius
```



operations are much quicker than if we used a vanilla Python `for`-loop.

Also, the fact that arrays must be homogenous lend themselves to more efficient representations in memory.

- We can time code in a Jupyter Notebook. Let's try and square a long sequence of integers and see how long it takes with a Python loop:

In [36]:

```
1 %%timeit
2 squares = []
3 for i in range(1_000_000):
4     squares.append(i * i)
```

31.3 ms ± 639 µs per loop (mean ± std. dev. of 7 runs, 10 loops each)

- In vanilla Python, this takes about 0.03 seconds per loop.

In numpy:

30x speedup!

In [37]:

```
1 %%timeit
2 squares = np.arange(1_000_000) ** 2
```

1.2 ms ± 28.6 µs per loop (mean ± std. dev. of 7 runs, 1,000 loops each)

- We can apply arithmetic operations to multiple arrays, provided they have the same length.

- The result is computed **element-wise**, which means that the arithmetic operation is applied to one pair of elements from each array at a time.

$$\begin{array}{c} 1 \\ 2 \end{array} - \begin{array}{c} 1 \\ 1 \end{array} = \begin{array}{c} 0 \\ 1 \end{array}$$

$$\begin{array}{c} 1 \\ 2 \end{array} * \begin{array}{c} 1 \\ 2 \end{array} = \begin{array}{c} 1 \\ 4 \end{array}$$

$$\begin{array}{c} 1 \\ 2 \end{array} / \begin{array}{c} 1 \\ 2 \end{array} = \begin{array}{c} 1 \\ 1 \end{array}$$

```
In [38]: 1 a = np.array([4, 5, -1])
          2 b = np.array([2, 3, 2])
```

```
In [39]: 1 a + b
```

Out[39]: array([6, 8, 1])

```
In [40]: 1 a / b
```

Out[40]: array([2. , 1.67, -0.5])

```
In [41]: 1 a ** 2 + b ** 2
```

Out[41]: array([20, 34, 5])





Array methods

- Arrays come equipped with several handy methods; some examples are below, but you can read about them all [here](#).

```
In [44]: 1 arr = np.array([3, 8, 4, -3.2])
```

```
In [45]: 1 arr.mean()
```

```
Out[45]: 2.95
```

```
In [46]: 1 arr.sum()
```

```
Out[46]: 11.8
```

```
In [47]: 1 arr.max()
```

```
Out[47]: 8.0
```

```
In [48]: 1 arr.argmax()
```

```
Out[48]: 1
```

```
In [ ]: 1 ...
```

```
In [ ]: 1 # An attribute, not a method.
```

```
2 ...
```





Array methods

- Arrays come equipped with several handy methods; some examples are below, but you can read about them all [here](#).

```
In [44]: 1 arr = np.array([3, 8, 4, -3.2])
```

```
In [45]: 1 arr.mean()
```

```
Out[45]: 2.95
```

```
In [46]: 1 arr.sum()
```

```
Out[46]: 11.8
```

```
In [47]: 1 arr.max()
```

```
Out[47]: 8.0
```

```
In [48]: 1 arr.argmax()
```

```
Out[48]: 1
```

one array.

```
In [55]: 1 (arr ** 2 + 2 * arr + np.log1p(arr ** 2)).min()
```

```
Out[55]: 6.25947884465547
```

```
In [ ]: 1 # An attribute, not a method.  
2 ...
```



Activity

Congrats! 🎉 You won the lottery 💰. Here's how your payout works: on the first day of September, you are paid \$0.01. Every day thereafter, your pay doubles, so on the second day you're paid \$0.02, on the third day you're paid \$0.04, on the fourth day you're paid \$0.08, and so on.

September has 30 days.

Write a **one-line expression** that uses the numbers `2` and `30`, along with the function `np.arange` and at least one array method, that computes the total amount **in dollars** you will be paid in September. No `for`-loops or list comprehensions allowed!

We have a [walkthrough video](#) of this problem (also posted on the course website under "videos" for Lecture 3), but don't watch it until you've tried it yourself!

```
In [62]: 1 np.arange(30)
```

```
Out[62]: array([ 0,  1,  2, ..., 27, 28, 29])
```

```
In [ ]:
```

1 2 4 8 16 32 ... - - - - - 29

```
In [ ]:
```

1 2⁰ + 2¹ + 2² + 2³ + 2⁴ . . . - - - - - + 2

```
In [ ]:
```

1





Boolean filtering

- Comparisons with arrays yield **Boolean** arrays! These can be used to answer questions about the values in an array.

In [67]: 1 temp_array

```
Out[67]: array([20. , 22.22, 18.33, 17.78, 16.67, 16.11, 15. , 17.78, 17.78,  
17.22, 18.33, 16.67])
```

```
In [69]: 1 temp_array >= 18
```

```
Out[69]: array([ True,  True,  True, False, False, False, False, False,  
   False,  True, False])
```

- How many values are greater than or equal to 18?

```
In [70]: 1 (temp_array >= 18).sum()
```

Out[70]: 4

In []:



```
Out[67]: array([20. , 22.22, 18.33, 17.78, 16.67, 16.11, 15. , 17.78, 17.78,  
17.22, 18.33, 16.67])
```

```
In [69]: 1 temp_array >= 18
```

```
Out[69]: array([ True,  True,  True, False, False, False, False, False,  
False,  True, False])
```

- How many values are greater than or equal to 18?

```
In [70]: 1 (temp_array >= 18).sum()
```

```
Out[70]: 4
```

```
In [71]: 1 np.count_nonzero(temp_array >= 18)
```

```
Out[71]: 4
```

- What fraction of values are greater than or equal to 18?

```
In [72]: 1 (temp_array >= 18).mean()
```

```
Out[72]: 0.3333333333333333
```

Mean = $\frac{\text{sum}}{n} = \frac{\# \text{ true}}{\# \text{ total}}$.

Out[71]: 4

- What fraction of values are greater than or equal to 18?

In [72]: 1 (temp_array >= 18).mean()

Out[72]: 0.3333333333333333

- Which values are greater than or equal to 18?

In [74]: 1 temp_array

Out[74]: array([20., 22.22, 18.33, 17.78, 16.67, 16.11, 15. , 17.78, 17.78, 17.22, 18.33, 16.67])

In [75]: 1 temp_array >= 18

Out[75]: array([True, True, True, False, False, False, False, False, True, False])

In [73]: 1 temp_array[temp_array >= 18]

Out[73]: array([20. , 22.22, 18.33, 18.33])

treats boolean array
as mask

```
In [71]: np.count_nonzero(temp_array) -> 10
```

Out[71]: 4

- What fraction of values are greater than or equal to 18?

```
In [72]: 1 | (temp_array >= 18).mean()
```

```
Out[72]: 0.3333333333333333
```

- Which values are greater than or equal to 18?

```
In [74]: 1 | temp_array
```

```
Out[74]: array([20., 22.22, 18.33, 17.78, 16.67, 16.11, 15. , 17.78, 17.78,  
    17.22, 18.33, 16.67])
```

```
In [75]: 1 | temp_array >= 18
```

```
Out[75]: array([True, True, True, False, False, False, False,  
    False, True, False])
```

```
In [73]: 1 | temp_array[temp_array >= 18]
```

```
Out[73]: array([20. , 22.22, 18.33, 18.33])
```

treats boolean array
as mask



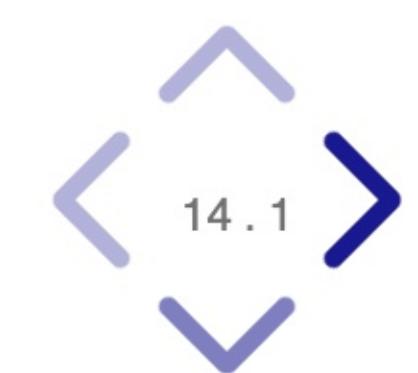
Note: & and | vs. and and or

- In Python, the standard symbols for "and" and "or" are, literally, `and` and `or`.

```
In [78]: 1 if (5 > 3 and 'h' + 'i' == 'hi') or (-2 > 0):  
2     print('success')
```

```
success
```

True *True* *False*





Note: & and | vs. and and or

- In Python, the standard symbols for "and" and "or" are, literally, `and` and `or`.

```
In [78]: 1 if (5 > 3 and 'h' + 'i' == 'hi') or (-2 > 0):  
2     print('success')
```

success

- But, when taking the **element-wise** and/or of two arrays, the standard operators don't work.
Instead, use the **bitwise** operators: `&` for "and", `|` for "or".

```
In [81]: 1 temp_array
```

```
Out[81]: array([20., 22.22, 18.33, 17.78, 16.67, 16.11, 15., 17.78, 17.78,  
                 17.22, 18.33, 16.67])
```

```
In [82]: 1 # Don't forget parentheses when using multiple conditions!  
2 temp_array[(temp_array % 2 == 0) | (temp_array == temp_array.min())]
```

```
Out[82]: array([20., 15.])
```

or •



```
Out[83]: array([[5, 1, 9, 7],  
 [9, 8, 2, 3],  
 [2, 5, 0, 4]])
```

```
In [84]: 1 # nums has 3 rows and 4 columns.  
2 nums.shape
```

```
Out[84]: (3, 4)
```

- In addition to creating 2D arrays from scratch, we can also create 2D arrays by *reshaping* other arrays.

```
In [87]: 1 nums
```

```
Out[87]: array([[5, 1, 9, 7],  
 [9, 8, 2, 3],  
 [2, 5, 0, 4]])
```

```
In [86]: 1 nums.reshape((2, 6))
```

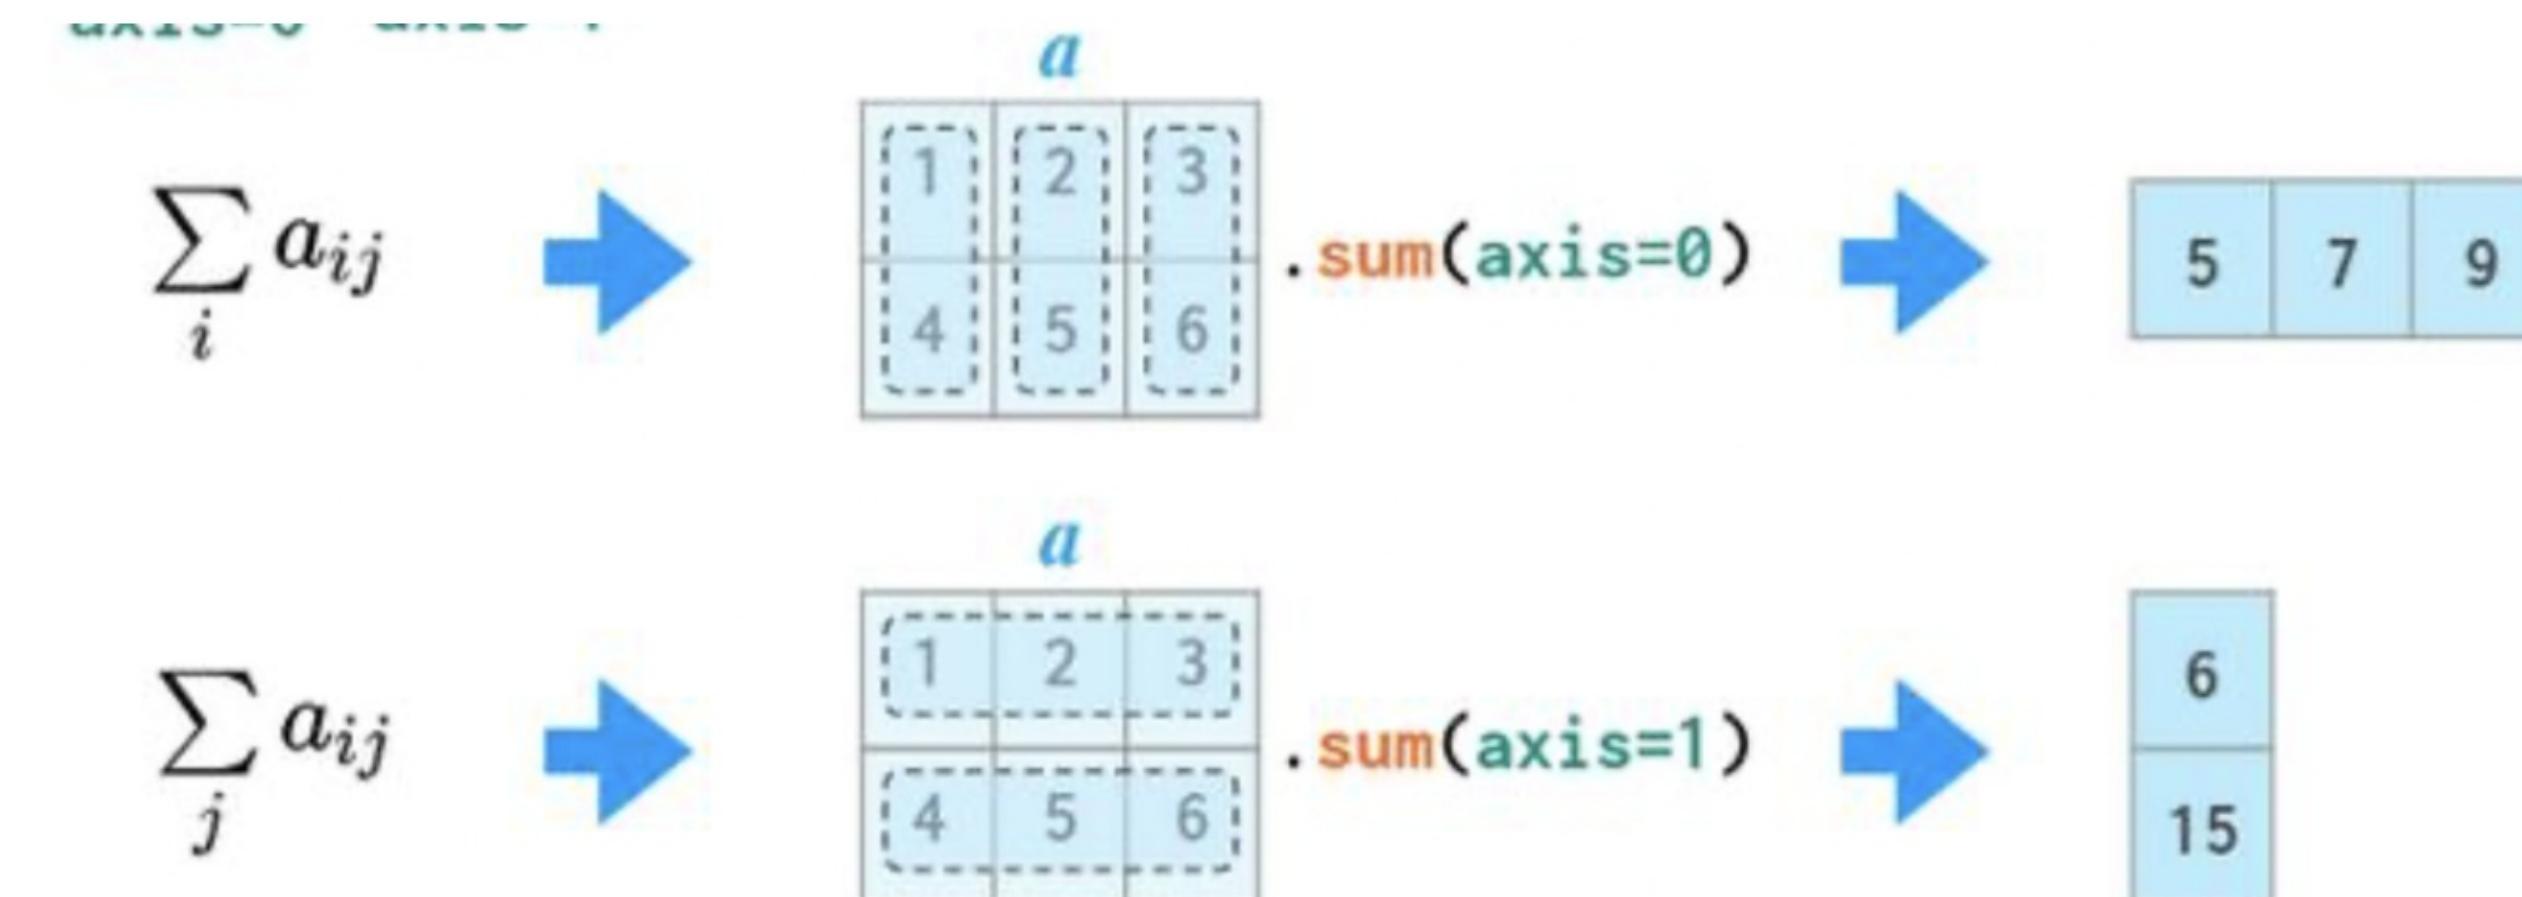
```
Out[86]: array([[5, 1, 9, 7, 9, 8],  
 [2, 3, 2, 5, 0, 4]])
```

reshape to 2 rows,
6 columns.

```
In [ ]: 1
```

```
In [ ]: 1
```

```
In [ ]: 1
```



```
In [89]: 1 a = np.arange(1, 7).reshape((2, 3))  
2 a
```

```
Out[89]: array([[1, 2, 3],  
                 [4, 5, 6]])
```

- If we specify `axis=0`, `a.sum` will "compress" along axis 0.

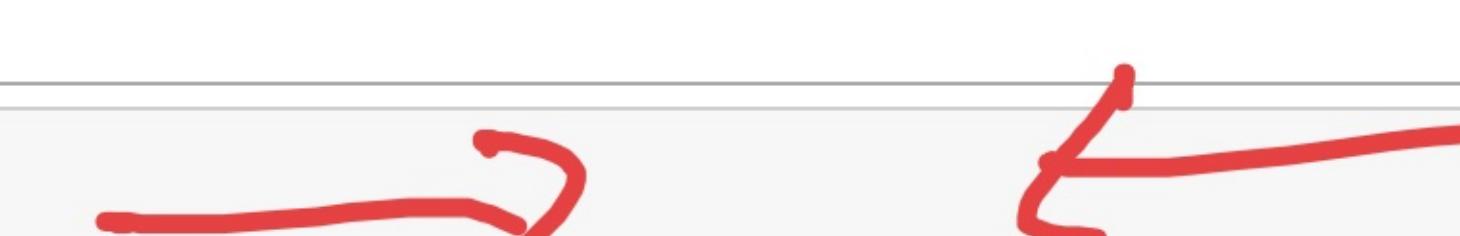


sum of columns

```
In [90]: 1 a.sum(axis=0)
```

```
Out[90]: array([5, 7, 9])
```

- If we specify `axis=1`, `a.sum` will "compress" along axis 1.



sum of rows

```
In [91]: 1 a.sum(axis=1)
```

```
Out[91]: array([ 6, 15])
```



- You can use [square brackets] to **slice** rows and columns out of an array, too.

- The general convention is:

array[<row positions>, <column positions>]

```
In [94]: 1 | a
```

```
Out[94]: array([[1, 2, 3],  
                 [4, 5, 6]])
```

```
In [95]: 1 | # Accesses row 0 and all columns.  
2 | a[0, :]
```

```
Out[95]: 2
```

```
In [98]: 1 | a[0, :]
```

```
Out[98]: array([1, 2, 3])
```

```
In [99]: 1 | a[0]
```

```
Out[99]: array([1, 2, 3])
```

```
In [100]: 1 | a[:, 1]
```

```
Out[100]: array([2, 5])
```





5 rows,
3 columns

Activity

Suppose we run the cell below.

```
s = (5, 3)
grid = np.ones(s) * 2 * np.arange(1, 16).reshape(s)
grid[-1, 1:].sum()
```

What is the output of the cell? Try and answer without writing any code. See the annotated slides for the solution.

In [1]: 1

In [1]: 1

grid = $2 \times [14, 15] \rightarrow [28, 30] \rightarrow \boxed{\sum 8}$

$\begin{array}{ccc} 1 & 2 & 3 \\ 4 & 5 & 6 \\ 7 & 8 & 9 \\ 10 & 11 & 12 \\ 13 & 14 & 15 \end{array}$

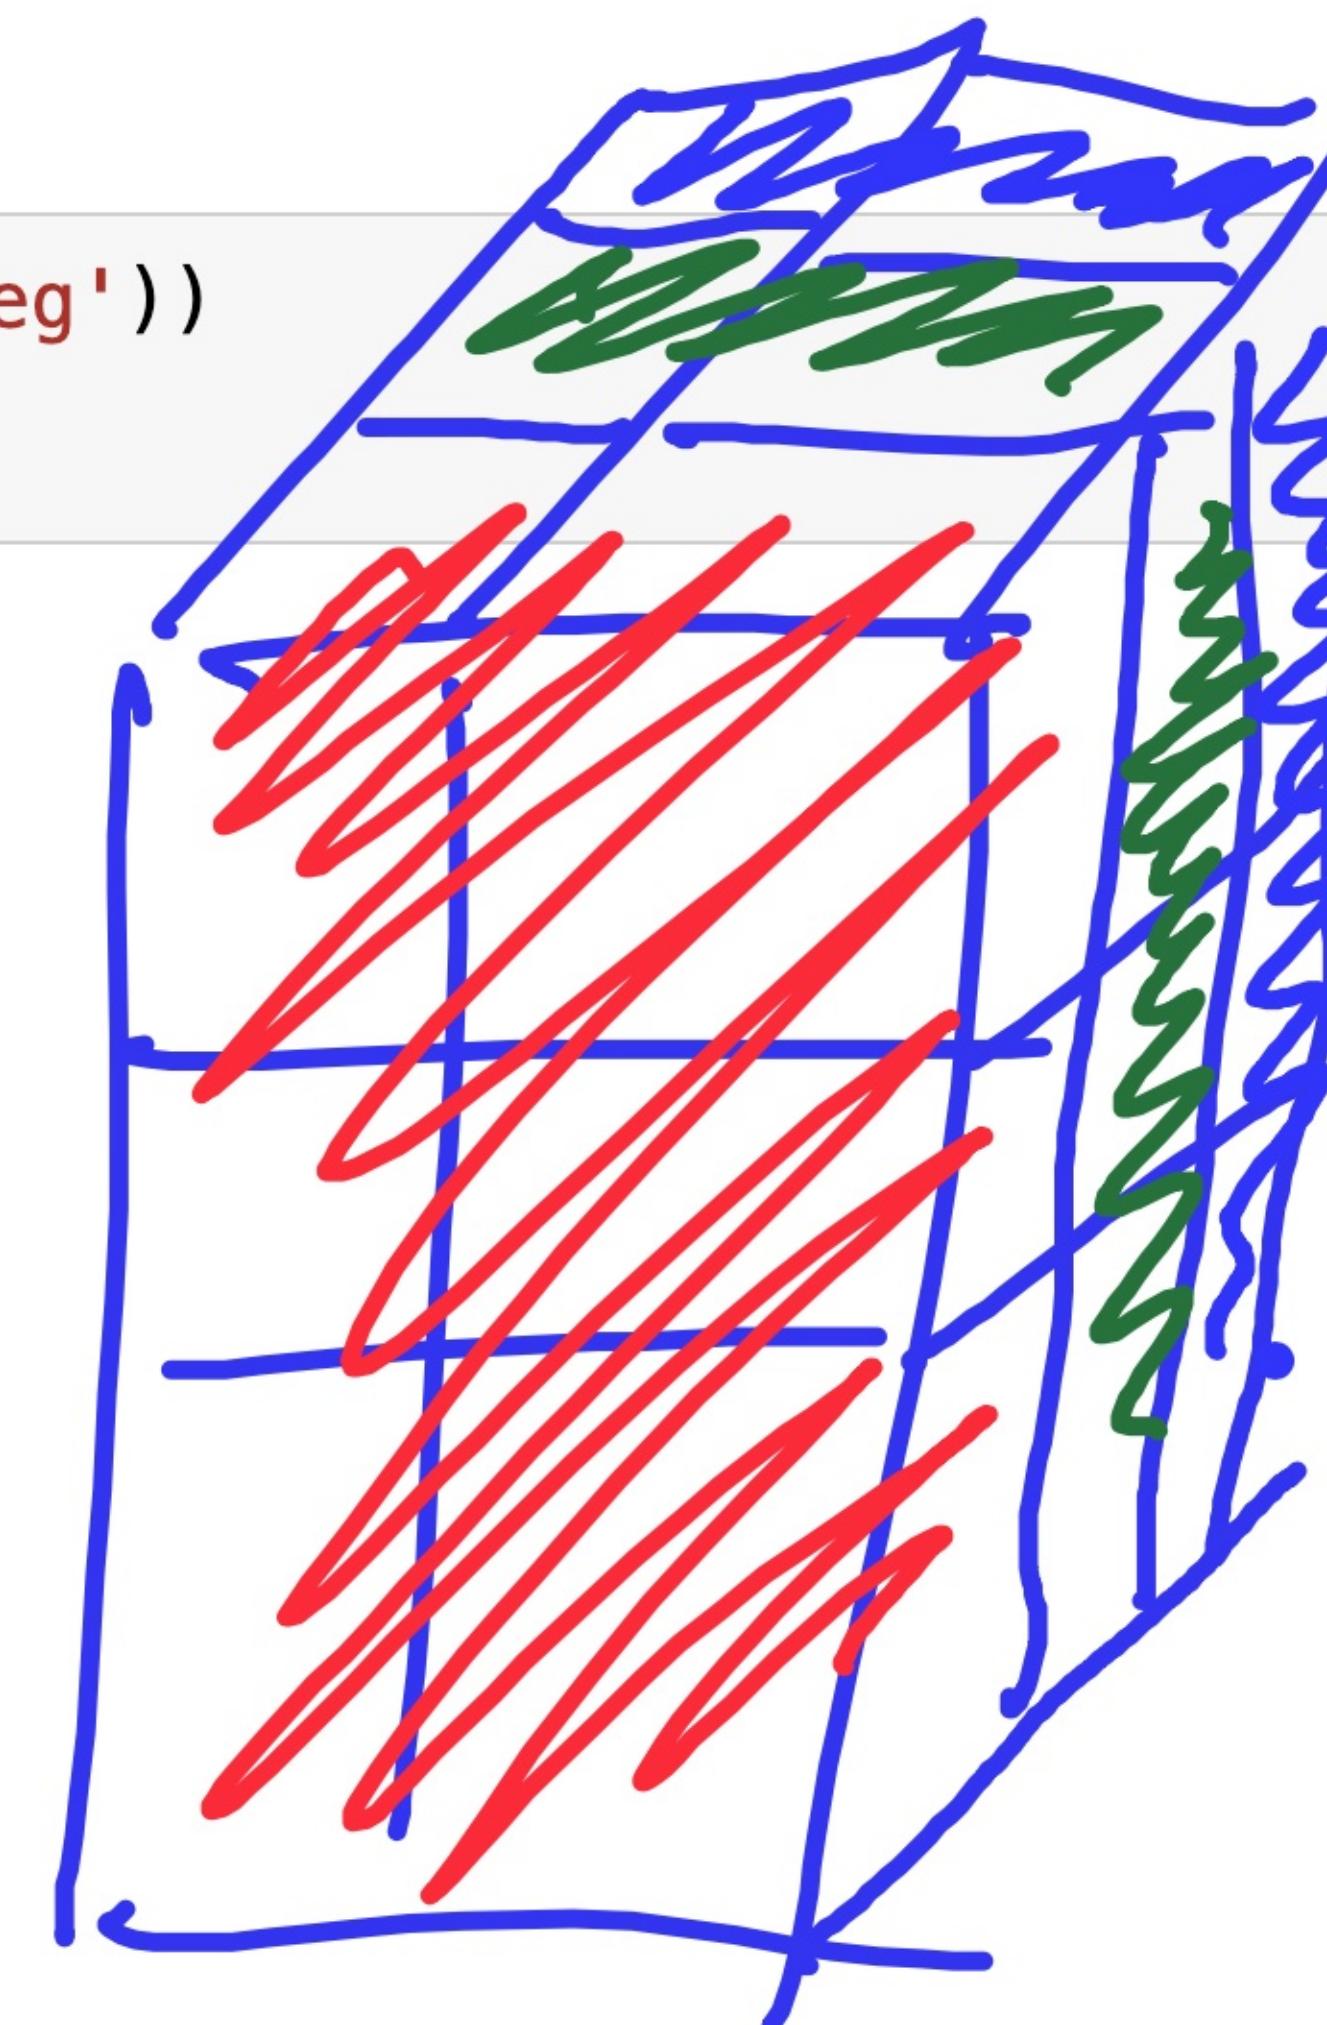
row cols $[-1, -2]$

value, green value, and blue value. Each of these can vary from 0 (least intensity) to 255 (most intensity).

Experiment with RGB colors [here](#).

```
In [102]: 1 img = np.asarray(Image.open('imgs/junior.jpeg'))  
2 img
```

```
Out[102]: array([[ [ 98,   62,   46],  
      [ 89,   56,   39],  
      [ 88,   56,   41],  
      ...,  
      [123,   78,   59],  
      [125,   78,   60],  
      [128,   81,   63]],  
  
[[ 96,   60,   44],  
 [ 89,   56,   39],  
 [ 89,   57,   42],  
 ...,  
 [124,   79,   60],  
 [125,   78,   60],  
 [127,   80,   62]],  
  
[[ 94,   58,   42],  
 [ 89,   56,   39],  
 [ 89,   57,   42],  
 ...,  
 [125,   78,   60],  
 [125,   78,   60],  
 [126,   79,   61]],  
  
...,  
  
[[ 89,   50,   11],  
 [ 85,   46,    7],  
 [ 81,   42,    3],
```





Matrix multiplication

- In the coming weeks, we'll start to rely more and more on tools from linear algebra.

You'll need this in Homework 2!

- Suppose the matrix A and vectors \vec{x} and \vec{y} are defined as follows:

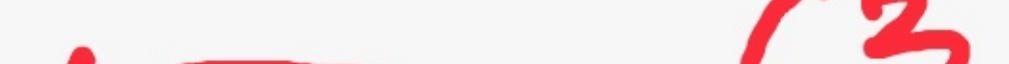
$$A = \begin{bmatrix} 2 & -5 & 1 \\ 0 & 3 & 2 \end{bmatrix} \quad \vec{x} = \begin{bmatrix} 1 \\ -1 \\ 4 \end{bmatrix} \quad \vec{y} = \begin{bmatrix} 3 \\ -2 \end{bmatrix}$$

2x3

3x1

2x1

```
In [ ]: 1 A = np.array([[2, -5, 1],  
2                 [0, 3, 2]])  
3 x = np.array([[1],  
4                 [-1],  
5                 [4]])  
6 y = np.array([[3],  
7                 [-2]])
```



(3, 1)

[1, -1, 4]

- Suppose the matrix A and vectors \vec{x} and \vec{y} are defined as follows:

$$A = \begin{bmatrix} 2 & -5 & 1 \\ 0 & 3 & 2 \end{bmatrix} \quad \vec{x} = \begin{bmatrix} 1 \\ -1 \\ 4 \end{bmatrix} \quad \vec{y} = \begin{bmatrix} 3 \\ -2 \end{bmatrix}$$

```
In [115]: 1 A = np.array([[2, -5, 1],
2                      [0, 3, 2]])
3 x = np.array([[1],
4                  [-1],
5                  [4]])
6 y = np.array([[3],
7                  [-2]])
```

- We can use `numpy` to compute various quantities involving A , \vec{x} , and \vec{y} .

For instance, what is the result of the product $A\vec{x}$?

See the annotated slides for the math worked out.

```
In [116]: 1 A @ x
Out[116]: array([11,
[-5]])
```

$$(2)(1) + (-5)(-1) + 1(4) \\ = 2 + 5 + 4 = 11$$



- Suppose the matrix A and vectors \vec{x} and \vec{y} are defined as follows:

$$A = \begin{bmatrix} 2 & -5 & 1 \\ 0 & 3 & 2 \end{bmatrix} \quad \vec{x} = \begin{bmatrix} 1 \\ -1 \\ 4 \end{bmatrix} \quad \vec{y} = \begin{bmatrix} 3 \\ 8 \\ -2 \end{bmatrix}$$

```
In [115]: 1 A = np.array([[2, -5, 1],
2                  [0, 3, 2]])
3 x = np.array([[1],
4                  [-1],
5                  [4]])
6 y = np.array([[3],
7                  [-2]])
```

- We can use `numpy` to compute various quantities involving A , \vec{x} , and \vec{y} .

For instance, what is the result of the product $A\vec{x}$?

See the annotated slides for the math worked out.

```
In [118]: 1 A * y
```

```
Out[118]: array([[ 6, -15,   3],
[ 0, -6,  -4]])
```



not matrix multiplication.



- The sequence of Fibonacci numbers,

1, 1, 2, 3, 5, 8, 13, 21, 34, 55, 89, ...

can be computed using matrix multiplication!

- It can be shown (with induction!) that if $f_1 = 1, f_2 = 1$, and $f_n = f_{n-1} + f_{n-2}, n \geq 2$, then:

$$\begin{bmatrix} f_{n+1} & f_n \\ f_n & f_{n-1} \end{bmatrix} = \begin{bmatrix} 1 & 1 \\ 1 & 0 \end{bmatrix}^n$$

```
In [125]: 1 fib = np.array([[1, 1],  
2 [1, 0]])
```

```
In [128]: 1 fib ** 8
```

```
Out[128]: array([[1, 1],  
 [1, 0]])
```

```
In [127]: 1 np.linalg.matrix_power(fib, 8)
```

```
Out[127]: array([[34, 21],  
 [21, 13]])
```

not the same!



Random sampling

- `np.random.choice` and `np.random.multinomial` allow us to draw random samples.
- `np.random.choice` returns randomly selected element(s) from the provided sequence.
By default, this is done **with** replacement, but it can be done without replacement, too.

```
In [227]: 1 unique_names = np.unique(np.load('data/sp25-names.npy', allow_pickle=True))
2 unique_names
```

```
Out[227]: array(['Abhinav', 'Alvin', 'Andrew', ..., 'Toby', 'Trent', 'Yu'],
                 dtype=object)
```

```
In [233]: 1 (unique_names == 'Suraj').sum()
```

```
Out[233]: 1
```

```
In [232]: 1 np.random.choice(unique_names, 5)
```

```
Out[232]: array(['Suraj', 'Suraj', 'Andrew', 'Andrew', 'Abhinav'], dtype=object)
```

```
In [ ]: 1 ...
```

by default,
sample WITH
REPLACEMENT.





Random sampling

- `np.random.choice` and `np.random.multinomial` allow us to draw random samples.
- `np.random.choice` returns randomly selected element(s) from the provided sequence.
By default, this is done **with** replacement, but it can be done without replacement, too.

```
In [227]: 1 unique_names = np.unique(np.load('data/sp25-names.npy', allow_pickle=True))
2 unique_names
```

```
Out[227]: array(['Abhinav', 'Alvin', 'Andrew', ..., 'Toby', 'Trent', 'Yu'],
                 dtype=object)
```

```
In [233]: 1 (unique_names == 'Suraj').sum()
```

```
Out[233]: 1
```

```
In [232]: 1 np.random.choice(unique_names, 5)
```

```
Out[232]: array(['Suraj', 'Suraj', 'Andrew', 'Andrew', 'Abhinav'], dtype=object)
```

```
In [ ]: 1 ...
```

by default,
sample WITH
REPLACEMENT





Example: Coin flipping

- **Question:** What is the probability that I see between 40 and 50 heads, inclusive, when I flip a fair coin 100 times?
- To estimate this probability, we need to:
 - Flip 100 fair coins and write down the number of heads,
 - and repeat that process many, many times.

```
In [291]: 1 np.random.multinomial(100, [0.5, 0.5])
```

```
Out[291]: array([50, 50])
```

```
In [292]: 1 np.random.multinomial(100, [0.5, 0.5], 100_000)
```

```
Out[292]: array([[43, 57],  
                 [50, 50],  
                 [53, 47],  
                 ...,  
                 [54, 46],  
                 [55, 45],  
                 [55, 45]])
```

10 million coins flipped!! .

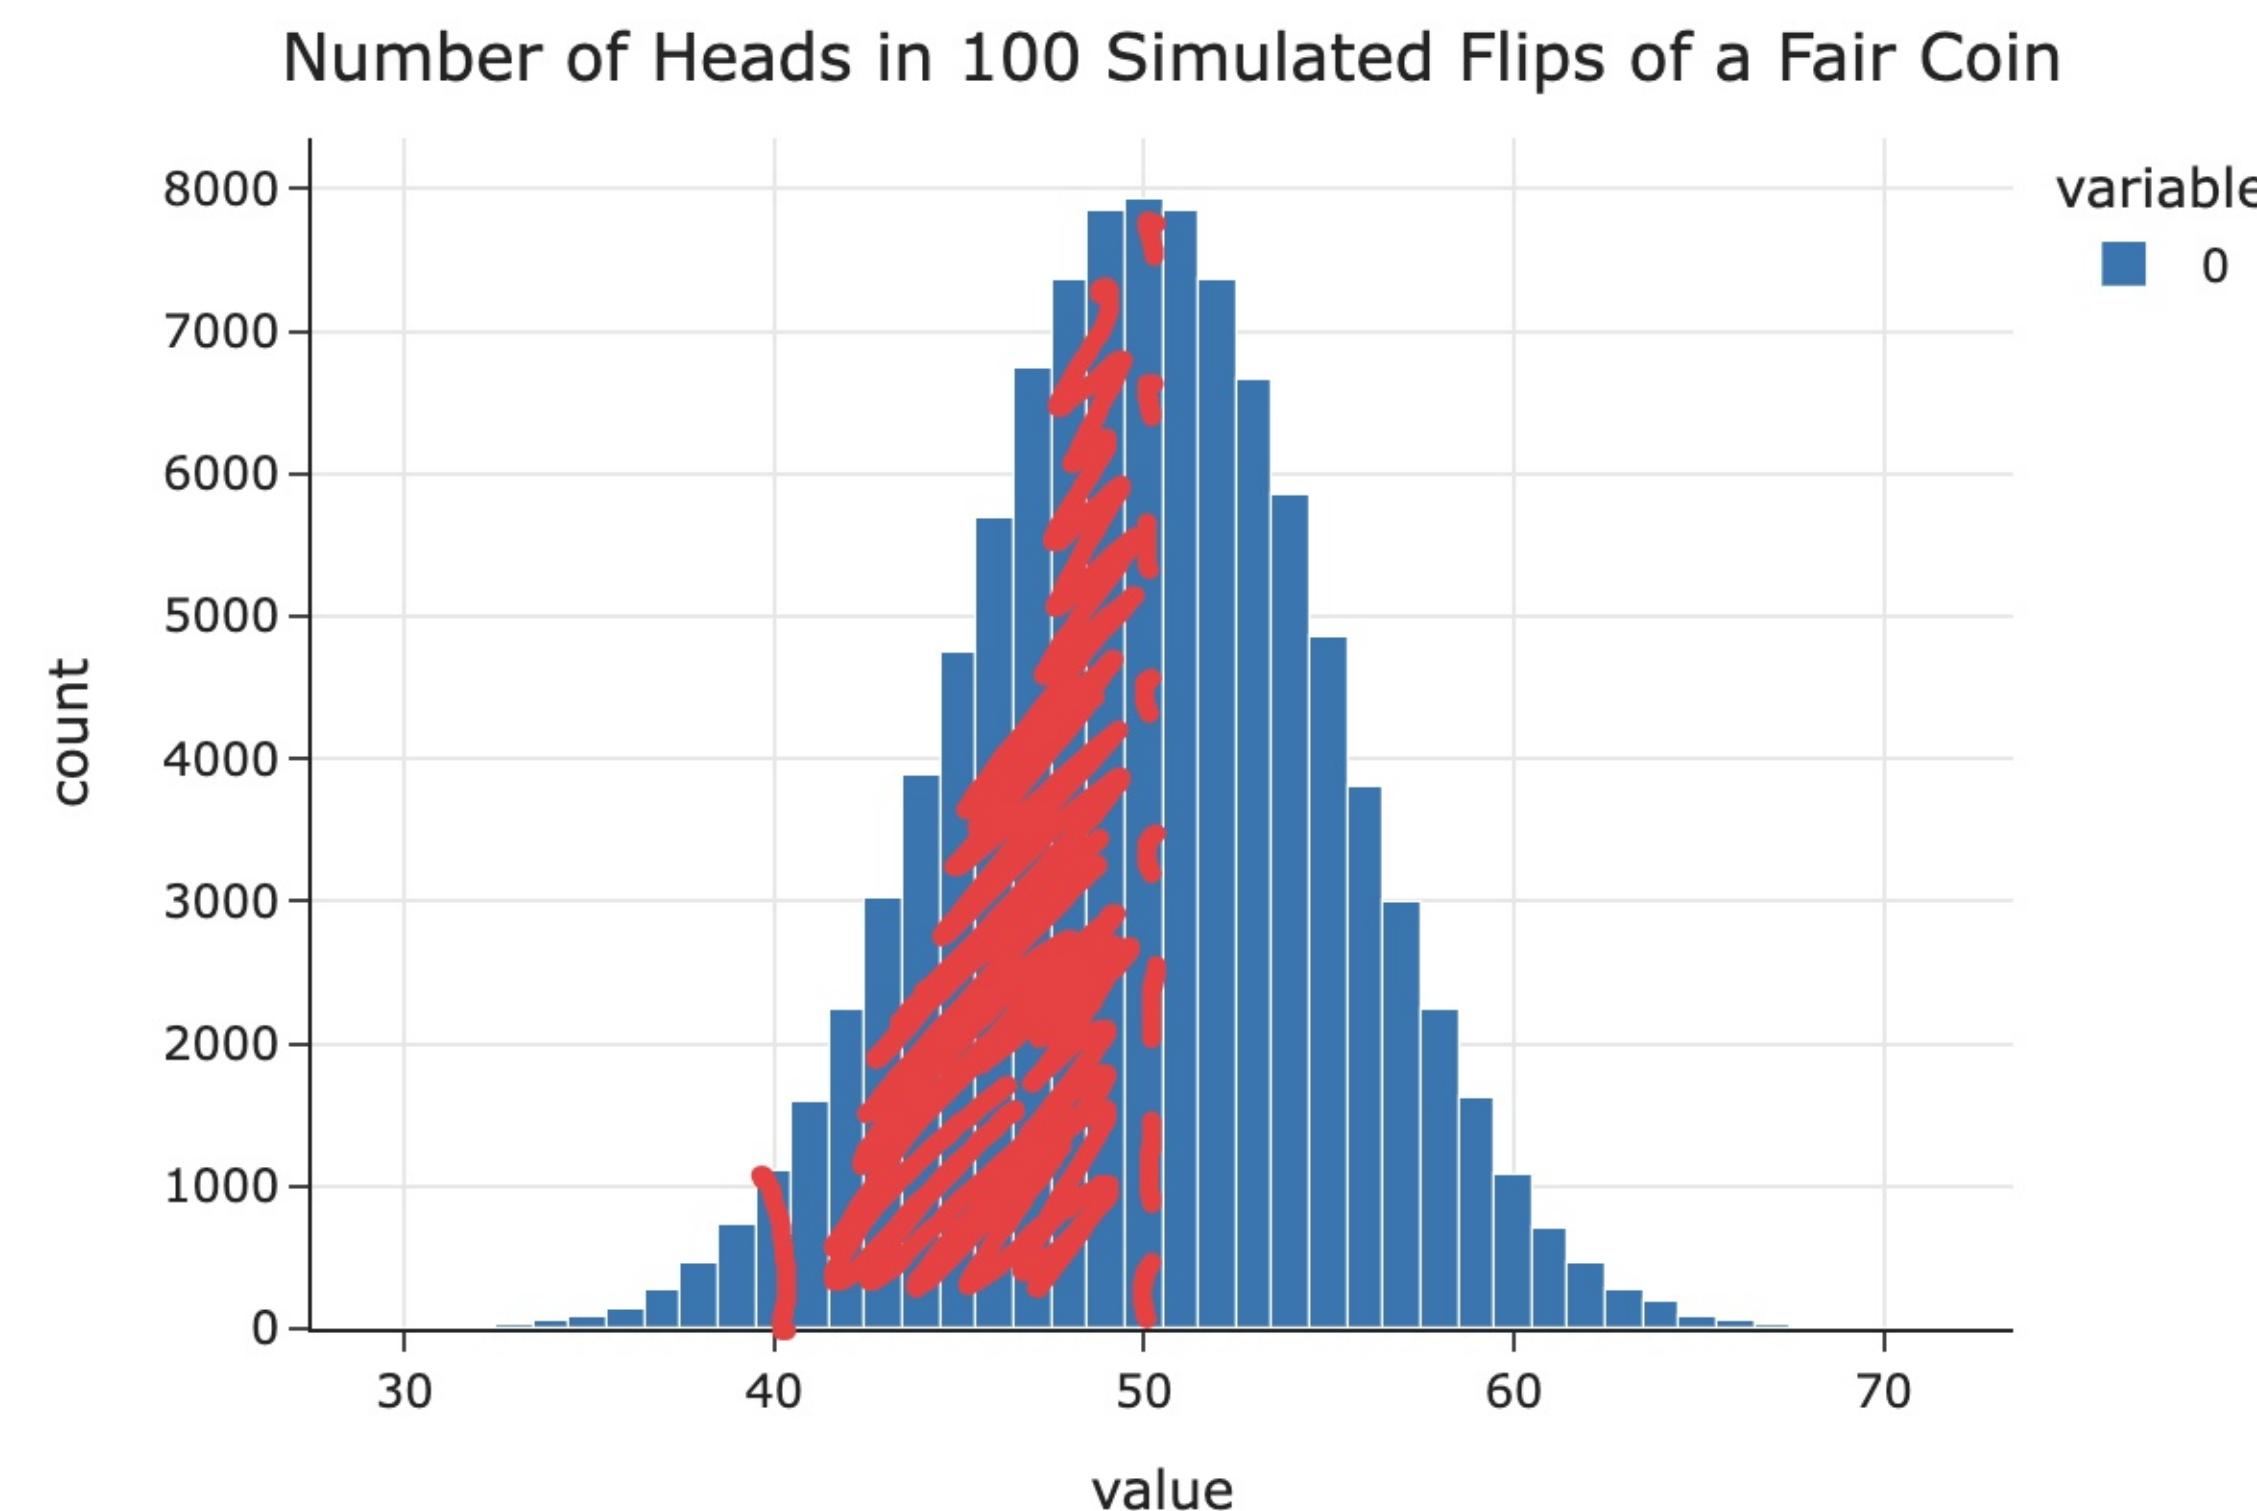
```
In [ ]: 1 # outcomes is an array with 100,000 elements,  
2 # each of which is the number of heads in 100 simulated flips of a fair coin.  
3 outcomes = ...  
4 outcomes
```





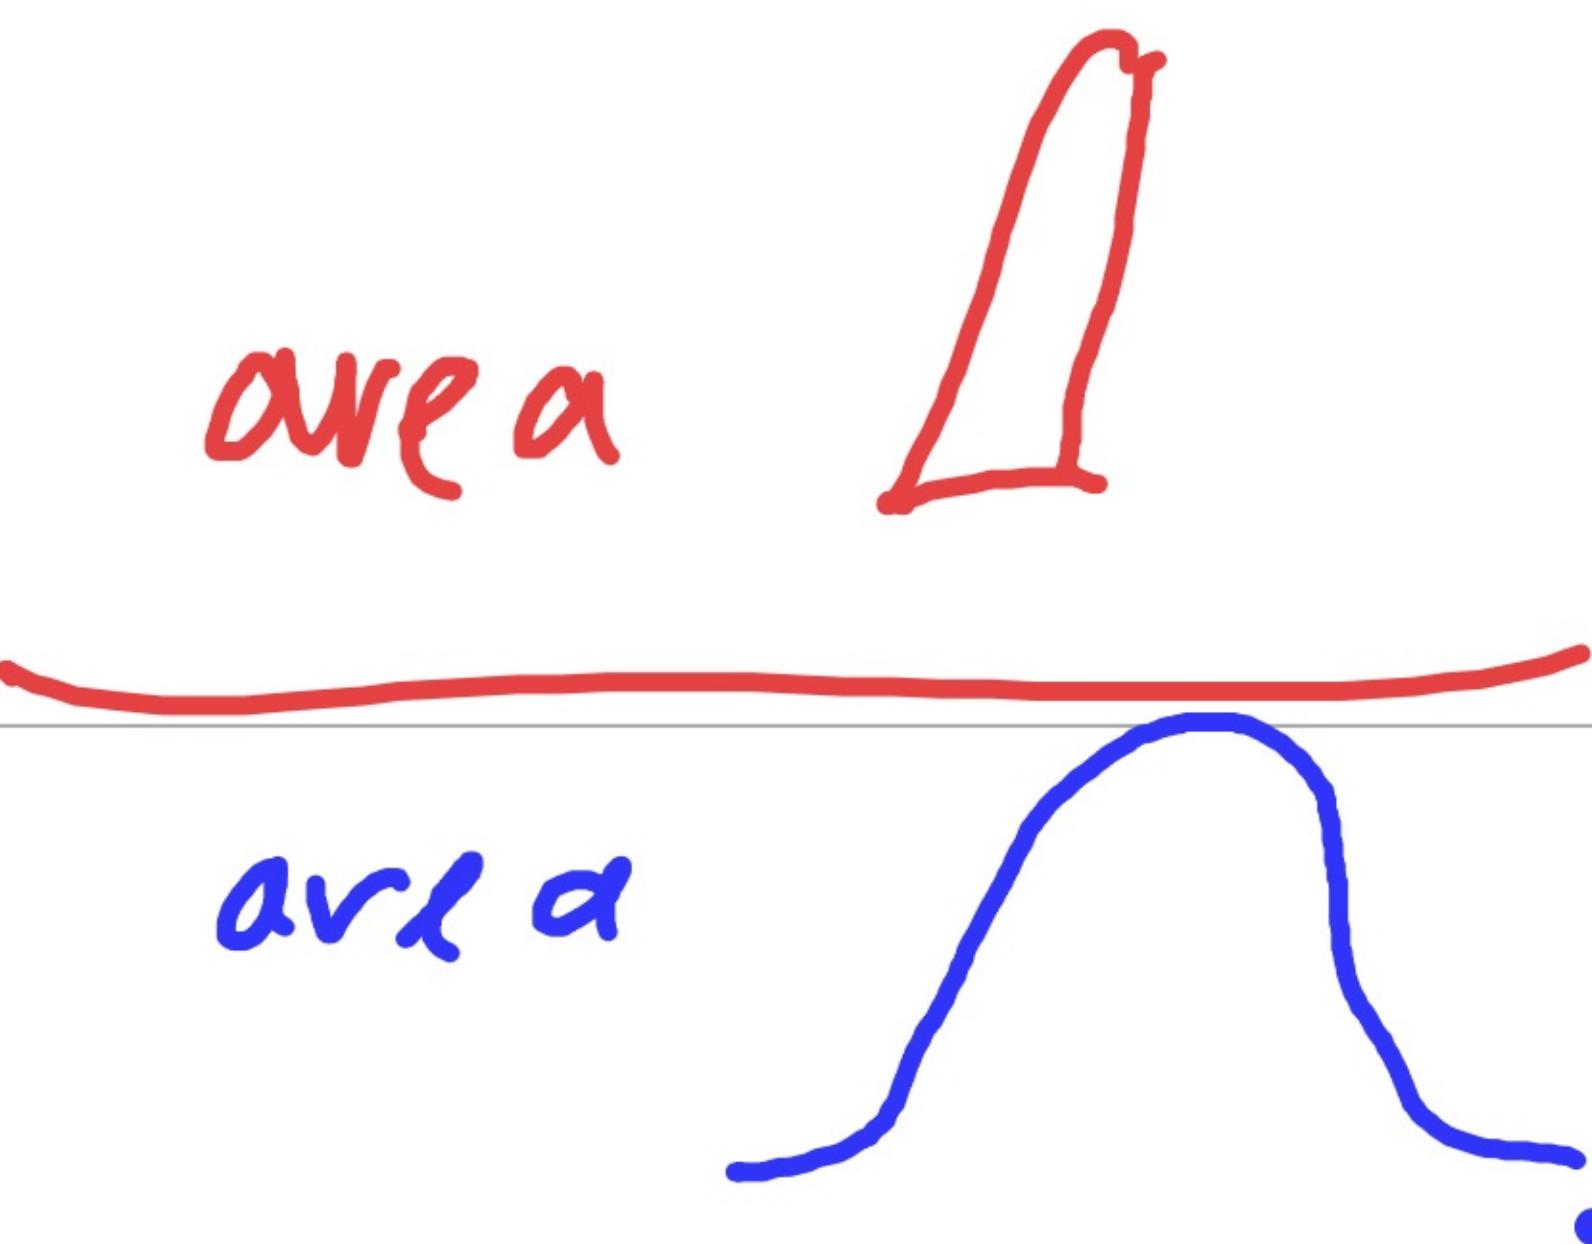
Estimating a probability from empirical results

```
In [301]: 1 px.histogram(outcomes, title='Number of Heads in 100 Simulated Flips of a Fair Coin')
```



$$P(40 \leq \text{heads} \leq 58)$$

are a



ark a





Example: Airplane seats

- A **permutation** of a sequence is a reshuffling of its elements.
`np.random.permutation` returns a permutation of the specified sequence.
- Suppose a flight on Wolverine Airlines is scheduled for n passengers, all of whom have an assigned seat.
- The airline loses track of seat assignments and everyone sits in a random seat.

What is the probability that *nobody* is in their originally assigned seat?

$$\frac{1}{e}.$$

