

15 rows x 7 columns

```
In [11]: 1 dogs.tail(2)
```

Out [11]:

	breed	kind	lifetime_cost	longevity	size	weight	height
40	Mastiff	working	13581.0	6.50	large	175.0	30.0
41	Saint Bernard	working	20022.0	7.78	large	155.0	26.5

- To sort by a column, use the `sort_values` method.

`ascending=False` is a **keyword argument**, meaning you need to specify the name of the argument to use it. You've seen some examples of this in the `plotly` part of Homework 1.

```
In [13]: 1 # Note that the index is no longer 0, 1, 2, ...!  
2 dogs.sort_values('lifetime_cost')
```

Out [13]:

	breed	kind	lifetime_cost	longevity	size	weight	height
40	Mastiff	working	13581.0	6.50	large	175.0	30.0
38	Bloodhound	hound	13824.0	6.75	large	85.0	25.0
39	Bullmastiff	working	13936.0	7.57	large	115.0	25.5
...
9	German Shorthaired Pointer	sporting	25842.0	11.46	large	62.5	24.0
7	Chihuahua	toy	26250.0	16.50	small	5.5	5.0
29	Giant Schnauzer	working	26686.0	10.00	large	77.5	25.5

42 rows x 7 columns

by default, sorts in ascending order

index no longer sorted!

In [34]: 1 dogs

Out[34]:

	kind	lifetime_cost	longevity	size	weight	height
breed						
Brittany	sporting	22589.0	12.92	medium	35.0	19.0
Cairn Terrier	terrier	21992.0	13.84	small	14.0	10.0
English Cocker Spaniel	sporting	18993.0	11.66	medium	30.0	16.0
...
Bullmastiff	working	13936.0	7.57	large	115.0	25.5
Mastiff	working	13581.0	6.50	large	175.0	30.0
Saint Bernard	working	20022.0	7.78	large	155.0	26.5

42 rows x 6 columns

In [35]: 1 # Returns a Series. Note the index appears again on the left!
2 dogs['height']

Out[35]: breed
Brittany 19.0
Cairn Terrier 10.0
English Cocker Spaniel 16.0
...
Bullmastiff 25.5
Mastiff 30.0
Saint Bernard 26.5
Name: height, Length: 42, dtype: float64

index stays

this is a Series

In []: 1 # Returns a DataFrame.

Bullmastiff 25.5
Mastiff 30.0
Saint Bernard 26.5
Name: height, Length: 42, dtype: float64

In [37]: 1 dogs[['height', 'longevity', 'lifetime_cost']]

Out[37]:

	height	longevity	lifetime_cost
breed			
Brittany	19.0	12.92	22589.0
Cairn Terrier	10.0	13.84	21992.0
English Cocker Spaniel	16.0	11.66	18993.0
...
Bullmastiff	25.5	7.57	13936.0
Mastiff	30.0	6.50	13581.0
Saint Bernard	26.5	7.78	20022.0

42 rows x 3 columns

one value → Series
list → DataFrame.

In [36]: 1 # Returns a DataFrame.
2 dogs[['height']]

Out[36]:

	height
breed	
Brittany	19.0
Cairn Terrier	10.0
English Cocker Spaniel	16.0
...	...
Bullmastiff	25.5

File index.pyx:181, in pandas._

File pandas/_libs/hashtable_class_helper.pxi:7080, in pandas._libs.hashtable.PyObjectHashTable.get_item()

File pandas/_libs/hashtable_class_helper.pxi:7088, in pandas._libs.hashtable.PyObjectHashTable.get_item()

KeyError: 'breed'

The above exception was the direct cause of the following exception:

KeyError Traceback (most recent call last)

Cell In[39], line 2

```
1 # Breeds are stored in the index, which is not a column!
----> 2 dogs['breed']
```

File ~/miniforge3/envs/pds/lib/python3.10/site-packages/pandas/core/frame.py:3896, in DataFrame.__getitem__(self, key)

```
3894 if self.columns.nlevels > 1:
3895     return self._getitem_multilevel(key)
-> 3896 indexer = self.columns.get_loc(key)
3897 if is_integer(indexer):
3898     indexer = [indexer]
```

File ~/miniforge3/envs/pds/lib/python3.10/site-packages/pandas/core/indexes/base.py:3797, in Index.get_loc(self, key)

```
3792 if isinstance(casted_key, slice) or (
3793     isinstance(casted_key, abc.Iterable)
3794     and any(isinstance(x, slice) for x in casted_key)
3795 ):
3796     raise InvalidIndexError(key)
-> 3797 raise KeyError(key) from err
3798 except TypeError:
3799     # If we have a listlike key, _check_indexing_error will raise
3800     # InvalidIndexError. Otherwise we fall through and re-raise
3801     # the TypeError.
3802     self._check_indexing_error(key)
```

KeyError: 'breed'

look at bottom

In []: 1 ...

Cairn Terrier	terrier	21992.0	13.84	small	14.0	10.0
English Cocker Spaniel	sporting	18993.0	11.66	medium	30.0	16.0
...
Bullmastiff	working	13936.0	7.57	large	115.0	25.5
Mastiff	working	13581.0	6.50	large	175.0	30.0
Saint Bernard	working	20022.0	7.78	large	155.0	26.5

42 rows x 6 columns

```
In [43]: 1 # What are the unique kinds of dogs?
        2 dogs['kind'].unique()
```

```
Out[43]: array(['sporting', 'terrier', 'herding', 'working', 'toy', 'non-sporting',
               'hound'], dtype=object)
```

```
In [44]: 1 # How many unique kinds of dogs are there?
        2 dogs['kind'].nunique()
```

```
Out[44]: 7
```

```
In [45]: 1 # What's the distribution of kinds?
        2 # value_counts is super useful – and I love asking exam questions about it!
        3 dogs['kind'].value_counts()
```

```
Out[45]: kind
sporting    12
terrier     8
working     7
toy         6
hound       5
herding     2
non-sporting 2
Name: count, dtype: int64
```

unique values

frequencies

In [58]: 1 show_chaining_slides()

```
breed
Mastiff          large
Saint Bernard   large
Newfoundland     large
...
Rhodesian Ridgeback large
Giant Schnauzer  large
Clumber Spaniel  medium
Name: size, Length: 10, dtype: object
```

```
size
large      9
medium     1
Name: count, dtype: int64
```

```
dogs.sort_values("weight", ascending=False).head(10)["size"].value_counts().index[0]
type: Series
```




Series support vectorized operations

- Series operations are vectorized, just like with arrays.
- When performing elementwise-operations involving multiple Series, **pandas** aligns the Series by their **index**.

```
In [61]: 1 x = pd.Series({'a': 1, 'b': 2})  
        2 x
```

```
Out[61]: a    1  
        b    2  
        dtype: int64
```

values in a Series
are shown on the right!

```
In [62]: 1 x * 5
```

```
Out[62]: a    5  
        b   10  
        dtype: int64
```

```
In [ ]: 1 y = pd.Series({'b': 5, 'c': -1, 'a': 10})  
        2 y
```

```
In [ ]: 1 # If x and y were regular numpy arrays, this would error because of the size mismatch.  
        2 x + y
```



with 2D arrays.

In [94]: 1 dogs

Out[94]:

	kind	lifetime_cost	longevity	size	weight	height
breed						
Brittany	sporting	22589.0	12.92	medium	35.0	19.0
Cairn Terrier	terrier	21992.0	13.84	small	14.0	10.0
English Cocker Spaniel	sporting	18993.0	11.66	medium	30.0	16.0
...
Bullmastiff	working	13936.0	7.57	large	115.0	25.5
Mastiff	working	13581.0	6.50	large	175.0	30.0
Saint Bernard	working	20022.0	7.78	large	155.0	26.5

42 rows x 6 columns

In [98]: 1 *# Try removing the iloc and see what happens!*
2 dogs.iloc[:2, 2:]

Out[98]:

	longevity	size	weight	height
breed				
Brittany	12.92	medium	35.0	19.0
Cairn Terrier	13.84	small	14.0	10.0



```
Out[114]: breed
Brittany False
Cairn Terrier True
English Cocker Spaniel False
...
Bullmastiff False
Mastiff False
Saint Bernard False
Name: kind, Length: 42, dtype: bool
```

```
In [115]: 1 dogs.loc[dogs['kind'] == 'terrier']
```

Out[115]:

	kind	lifetime_cost	longevity	size	weight	height
breed						
Cairn Terrier	terrier	21992.0	13.84	small	14.0	10.0
Miniature Schnauzer	terrier	20087.0	11.81	small	15.5	13.0
Norfolk Terrier	terrier	24308.0	13.07	small	12.0	9.5
...
Scottish Terrier	terrier	17525.0	10.69	small	20.0	10.0
Kerry Blue Terrier	terrier	17240.0	9.40	medium	36.5	18.5
Bull Terrier	terrier	18490.0	10.21	medium	60.0	21.5

8 rows x 6 columns

keeps just
the rows
where Series
has True



- Example: Among all breeds with 'Retriever' in the name, which is the second tallest?

```
In [149]: 1 (
          2     dogs.loc[dogs.index.str.contains('Retriever'), 'height']
          3     .sort_values(ascending=False)
          4     .index[1]
          5 )
```

Out[149]: 'Labrador Retriever'

```
In [145]: 1 # Since we're selecting both rows AND columns, we do need loc here.
          2 (
          3     dogs[dogs.index.str.contains('Retriever')]
          4     .sort_values('height', ascending=False)
          5     .index[1]
          6 )
```

Out[145]: 'Labrador Retriever'

understand the difference!

- Example: Among all breeds with 'Retriever' in the name, which is the second tallest?

```
In [149]: 1 (
          2     dogs.loc[dogs.index.str.contains('Retriever'), 'height']
          3     .sort_values(ascending=False)
          4     .index[1]
          5 )
```

Out[149]: 'Labrador Retriever'

```
In [145]: 1 # Since we're selecting both rows AND columns, we do need loc here.
          2 (
          3     dogs[dogs.index.str.contains('Retriever')]
          4     .sort_values('height', ascending=False)
          5     .index[1]
          6 )
```

Out[145]: 'Labrador Retriever'

understand the difference!

In []: 1

Find the average height of medium dogs.

```
In [174]: 1 dogs.loc[dogs['size'] == 'small', 'height'].mean()
```

```
Out[174]: 10.883333333333333
```

```
In [172]: 1 dogs.loc[dogs['size'] == 'medium', 'height'].mean()
```

```
Out[172]: 19.0
```

```
In [175]: 1 dogs.loc[dogs['size'] == 'large', 'height'].mean()
```

```
Out[175]: 25.8
```

```
In [178]: 1 dogs.groupby('size')['height'].mean()
```

```
Out[178]: size
large      25.80
medium     19.00
small      10.88
Name: height, dtype: float64
```

In []: 1