

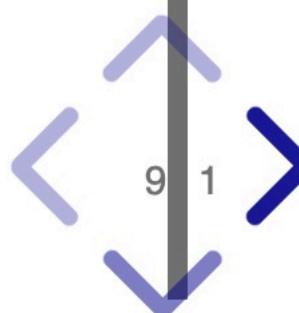
In [9]:

```
1 from sklearn.preprocessing import PolynomialFeatures
2 # fit_transform fits and transforms the same input.
3 # We tell it not to add a column of 1s, because
4 # LinearRegression() does this automatically later on. because Lin. Reg.
5 # already
5 d2 = PolynomialFeatures(3, include_bias=False)
6 d2.fit_transform(np.array([1, 2, 3, 4, -2]).reshape(-1, 1))
```

Out[9]: array([[1., 1., 1.],
 [2., 4., 8.],
 [3., 9., 27.],
 [4., 16., 64.],
 [-2., 4., -8.]]) . . .

$x^1 \quad x^2 \quad x^3$

because Lin. Reg.
already
adds
intercept!

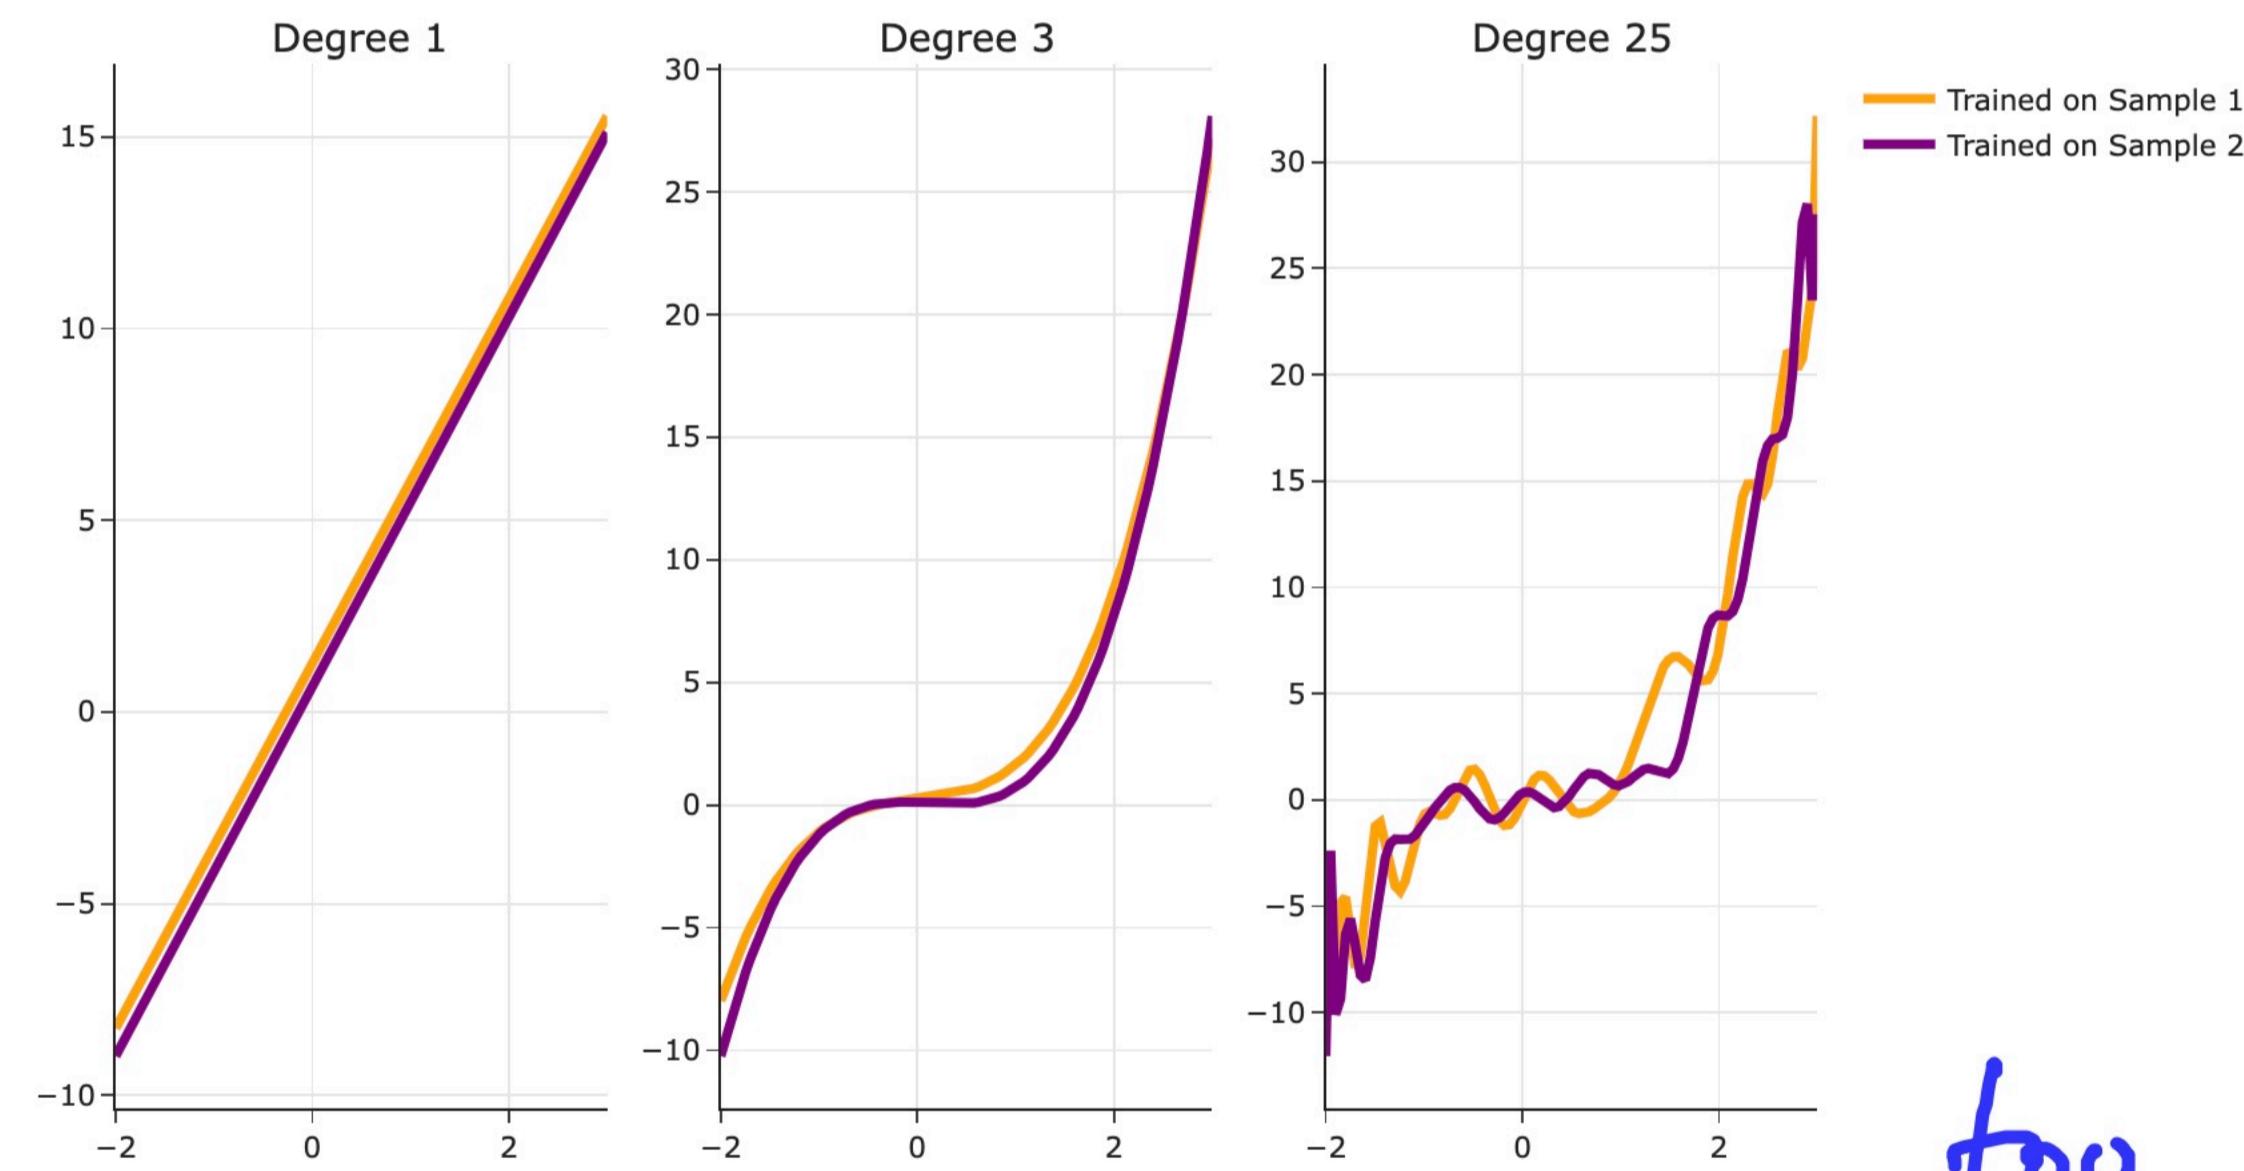




- Model variance: The variance of a model's predictions, across all datasets.

- In other words, for a given \vec{x}_i , how much does $H^*(\vec{x}_i)$ vary across all datasets?

In [13]: 1 fig



too flexible .

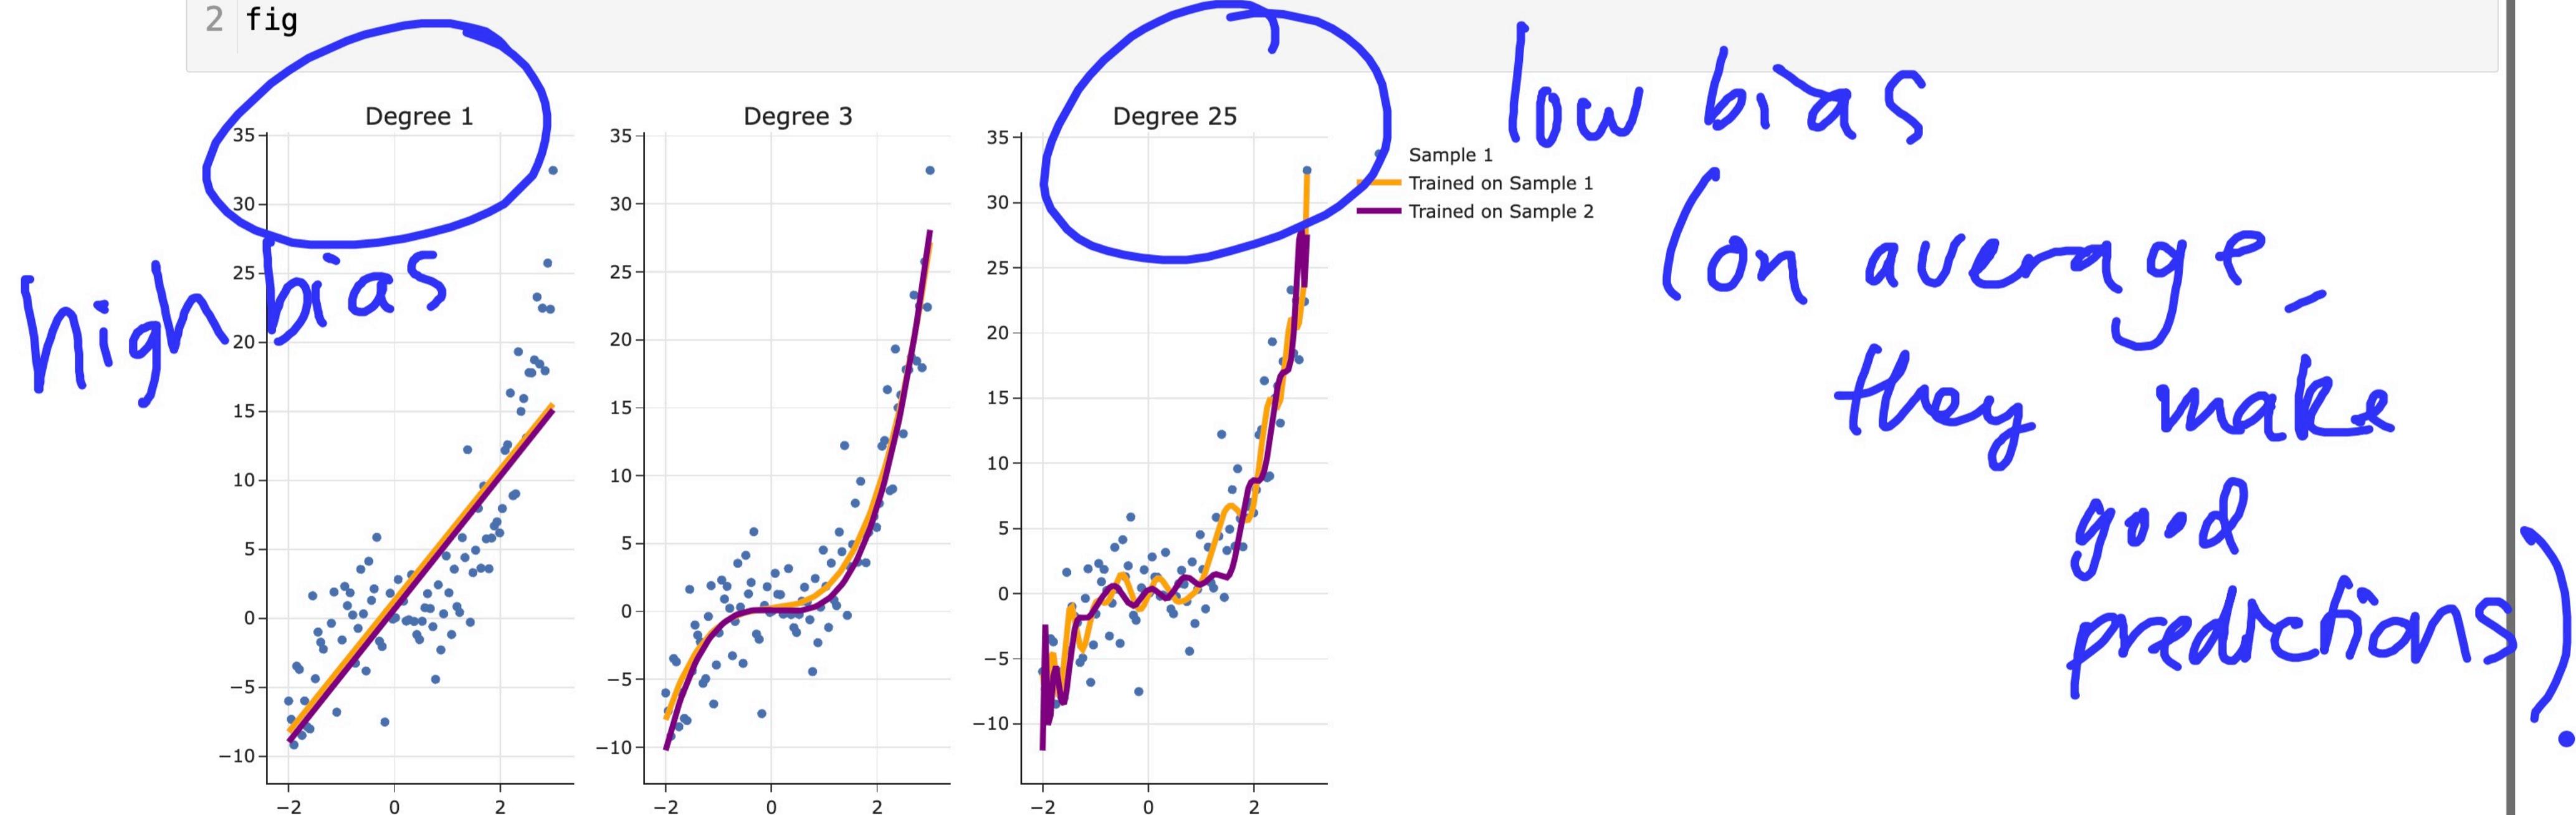
- Low model variance is good! ✓
- High model variance is a sign of **overfitting**, i.e. that our model is too **complicated** and is prone to fitting to the noise in our training data.



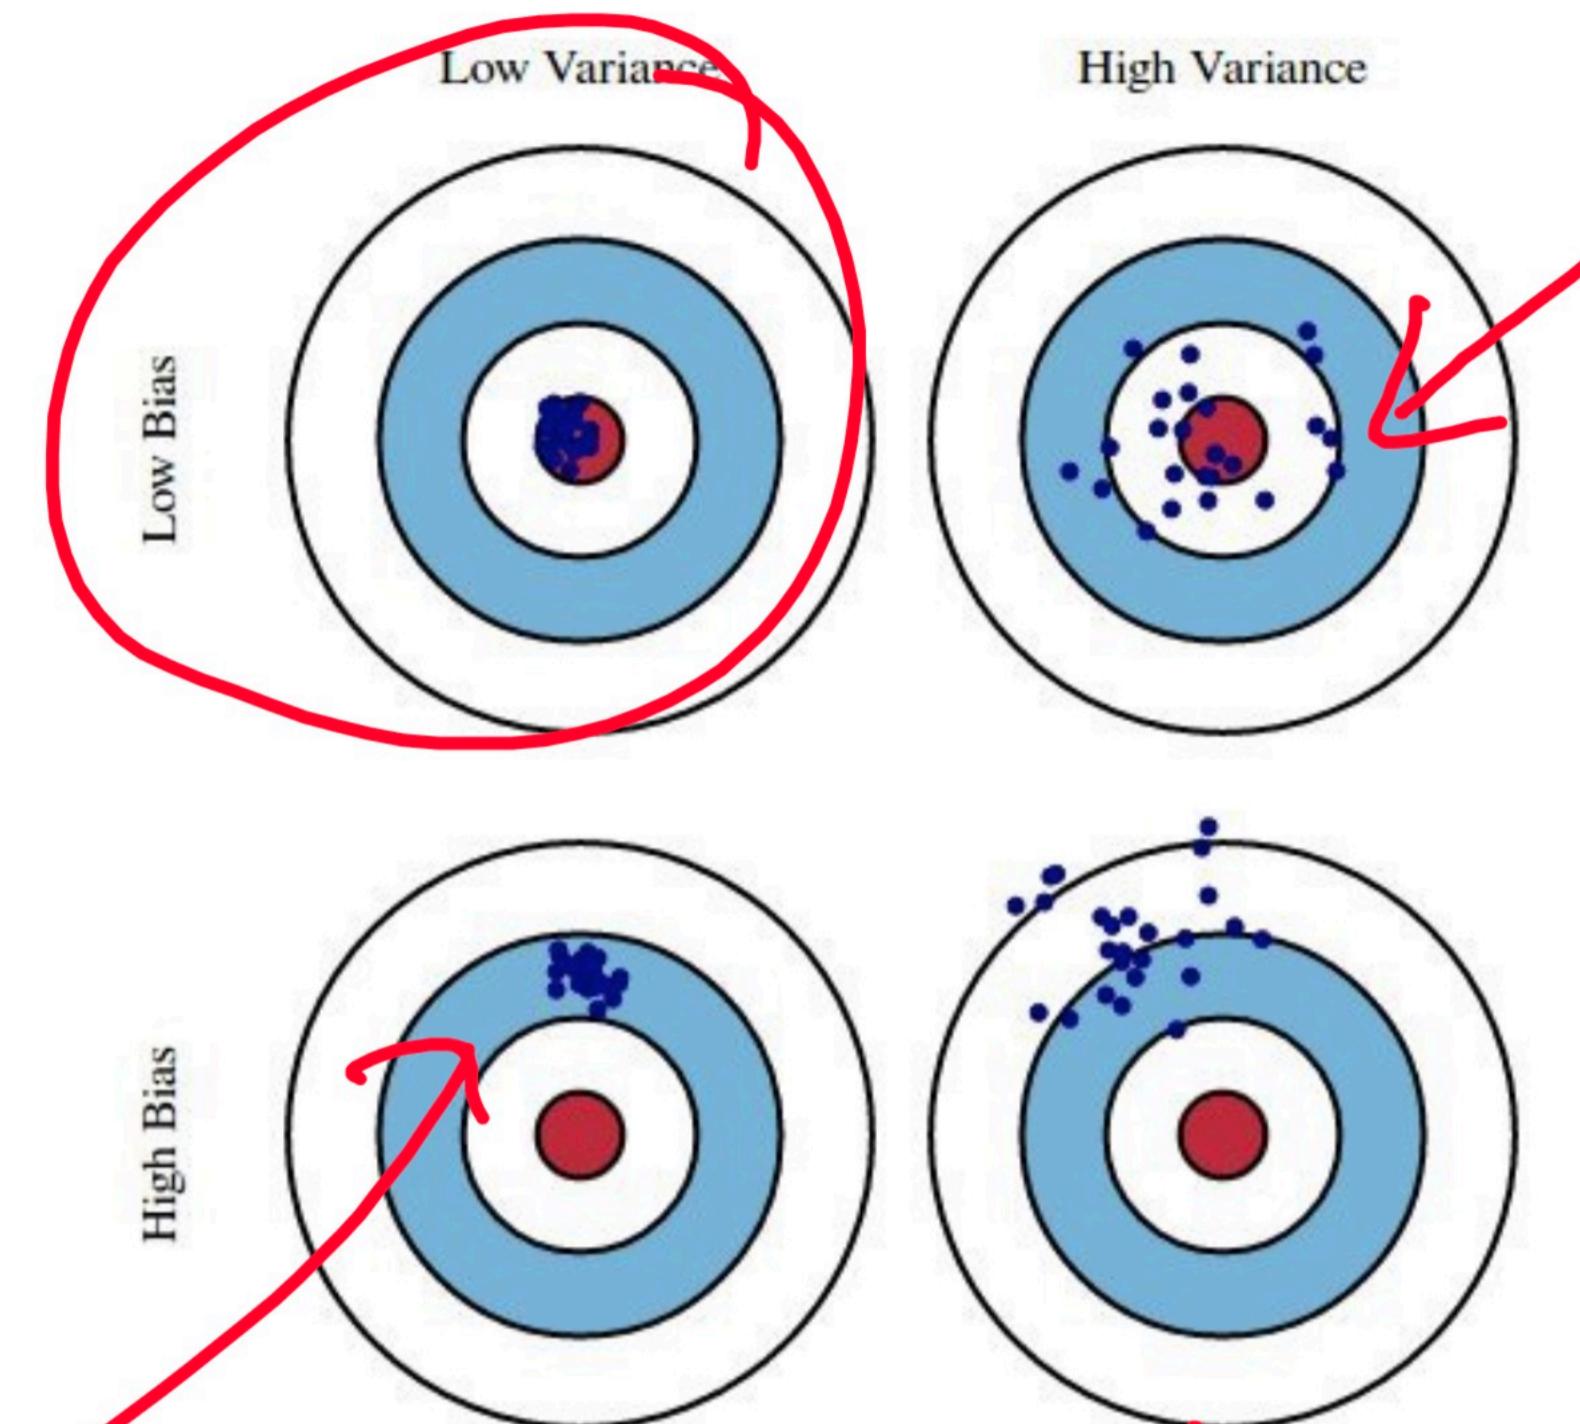


- **Model bias:** The averaged deviation between a predicted value and an actual value, across all datasets.
- In other words, for a given \vec{x}_i , how far is $H^*(\vec{x}_i)$ from the true y_i , on average?

```
In [14]: 1 fig = util.plot_multiple_models(sample_1, sample_2, degs=[1, 3, 25], data=True)
2 fig
```

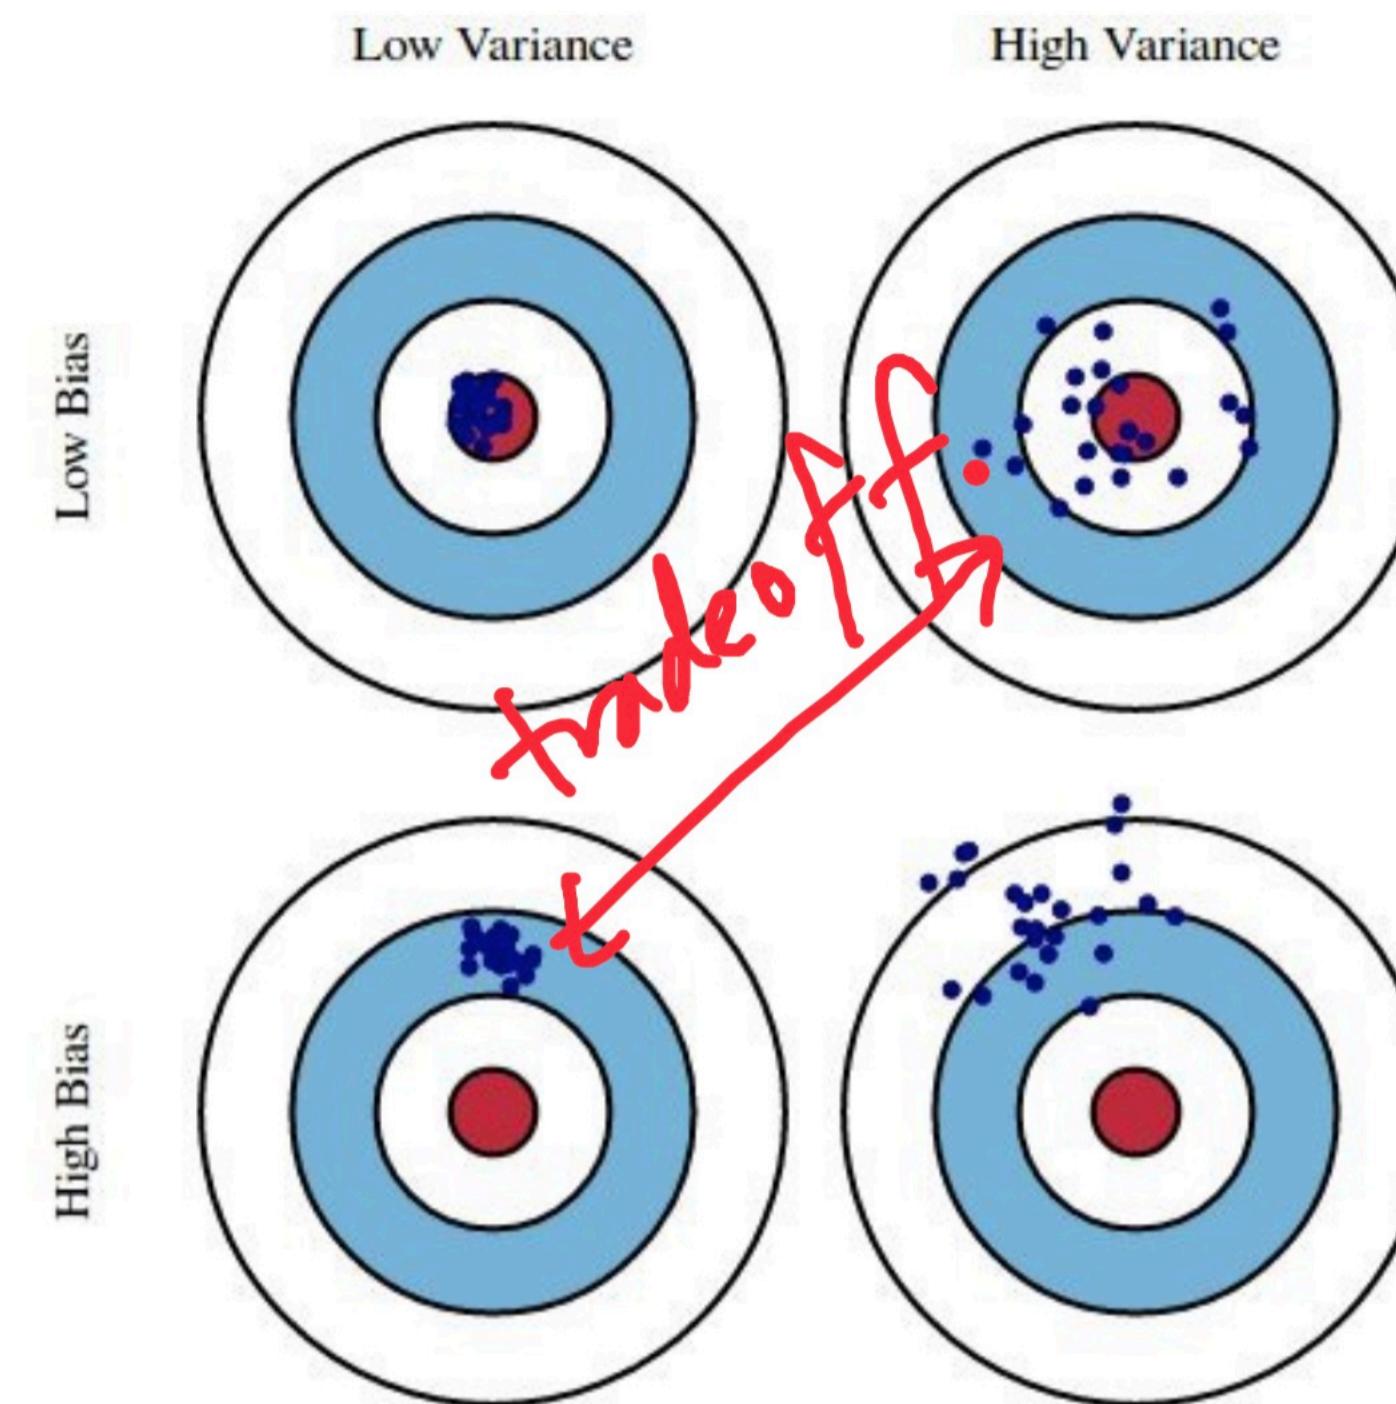


ideal



on average,
predictions
are good,
but each one
can be
very wrong
&
different

- Here, suppose:
 - The **red bulls-eye** represents your **true weight and height** 🚶.
 - The **dark blue darts** represent **predictions of your weight and height** using different models that were fit using different samples drawn from the same population.



- Here, suppose:
 - The **red bulls-eye** represents your **true weight and height** 🧑.
 - The **dark blue darts** represent **predictions of your weight and height** using different models that were fit using different samples drawn from the same population.
- We'd like our models to be in the top left, but in practice that's hard to achieve!



Risk vs. empirical risk

aka "average loss":

- Since Lecture 11, we've minimized **empirical risk** to find optimal model parameters \vec{w}^* :

$$\vec{w}^* = \underset{\vec{w}}{\operatorname{argmin}} \frac{1}{n} \sum_{i=1}^n (y_i - H(\vec{x}_i))^2$$

$$\vec{w} = \underset{\vec{w}}{\operatorname{arg\,min}} -\frac{1}{n} \sum_{i=1}^n (y_i - \Pi(\vec{x}_i))^2$$

- **Key idea:** A model that works well on past data should work well on future data, if future data looks like past data.
- What we really want is for the:
 - **expected** loss for a new data point $(\vec{x}_{\text{new}}, y_{\text{new}})$,
 - drawn from the same population as the training set, to be small.

That is, we want to minimize **risk**:

$$\text{risk} = \mathbb{E}[y_{\text{new}} - H(\vec{x}_{\text{new}})]^2$$

expected value of a random variable

- In general, we don't know the entire population distribution of xs and ys, so we can't compute risk exactly.

That's why we compute **empirical** risk!

empirical risk \approx risk.

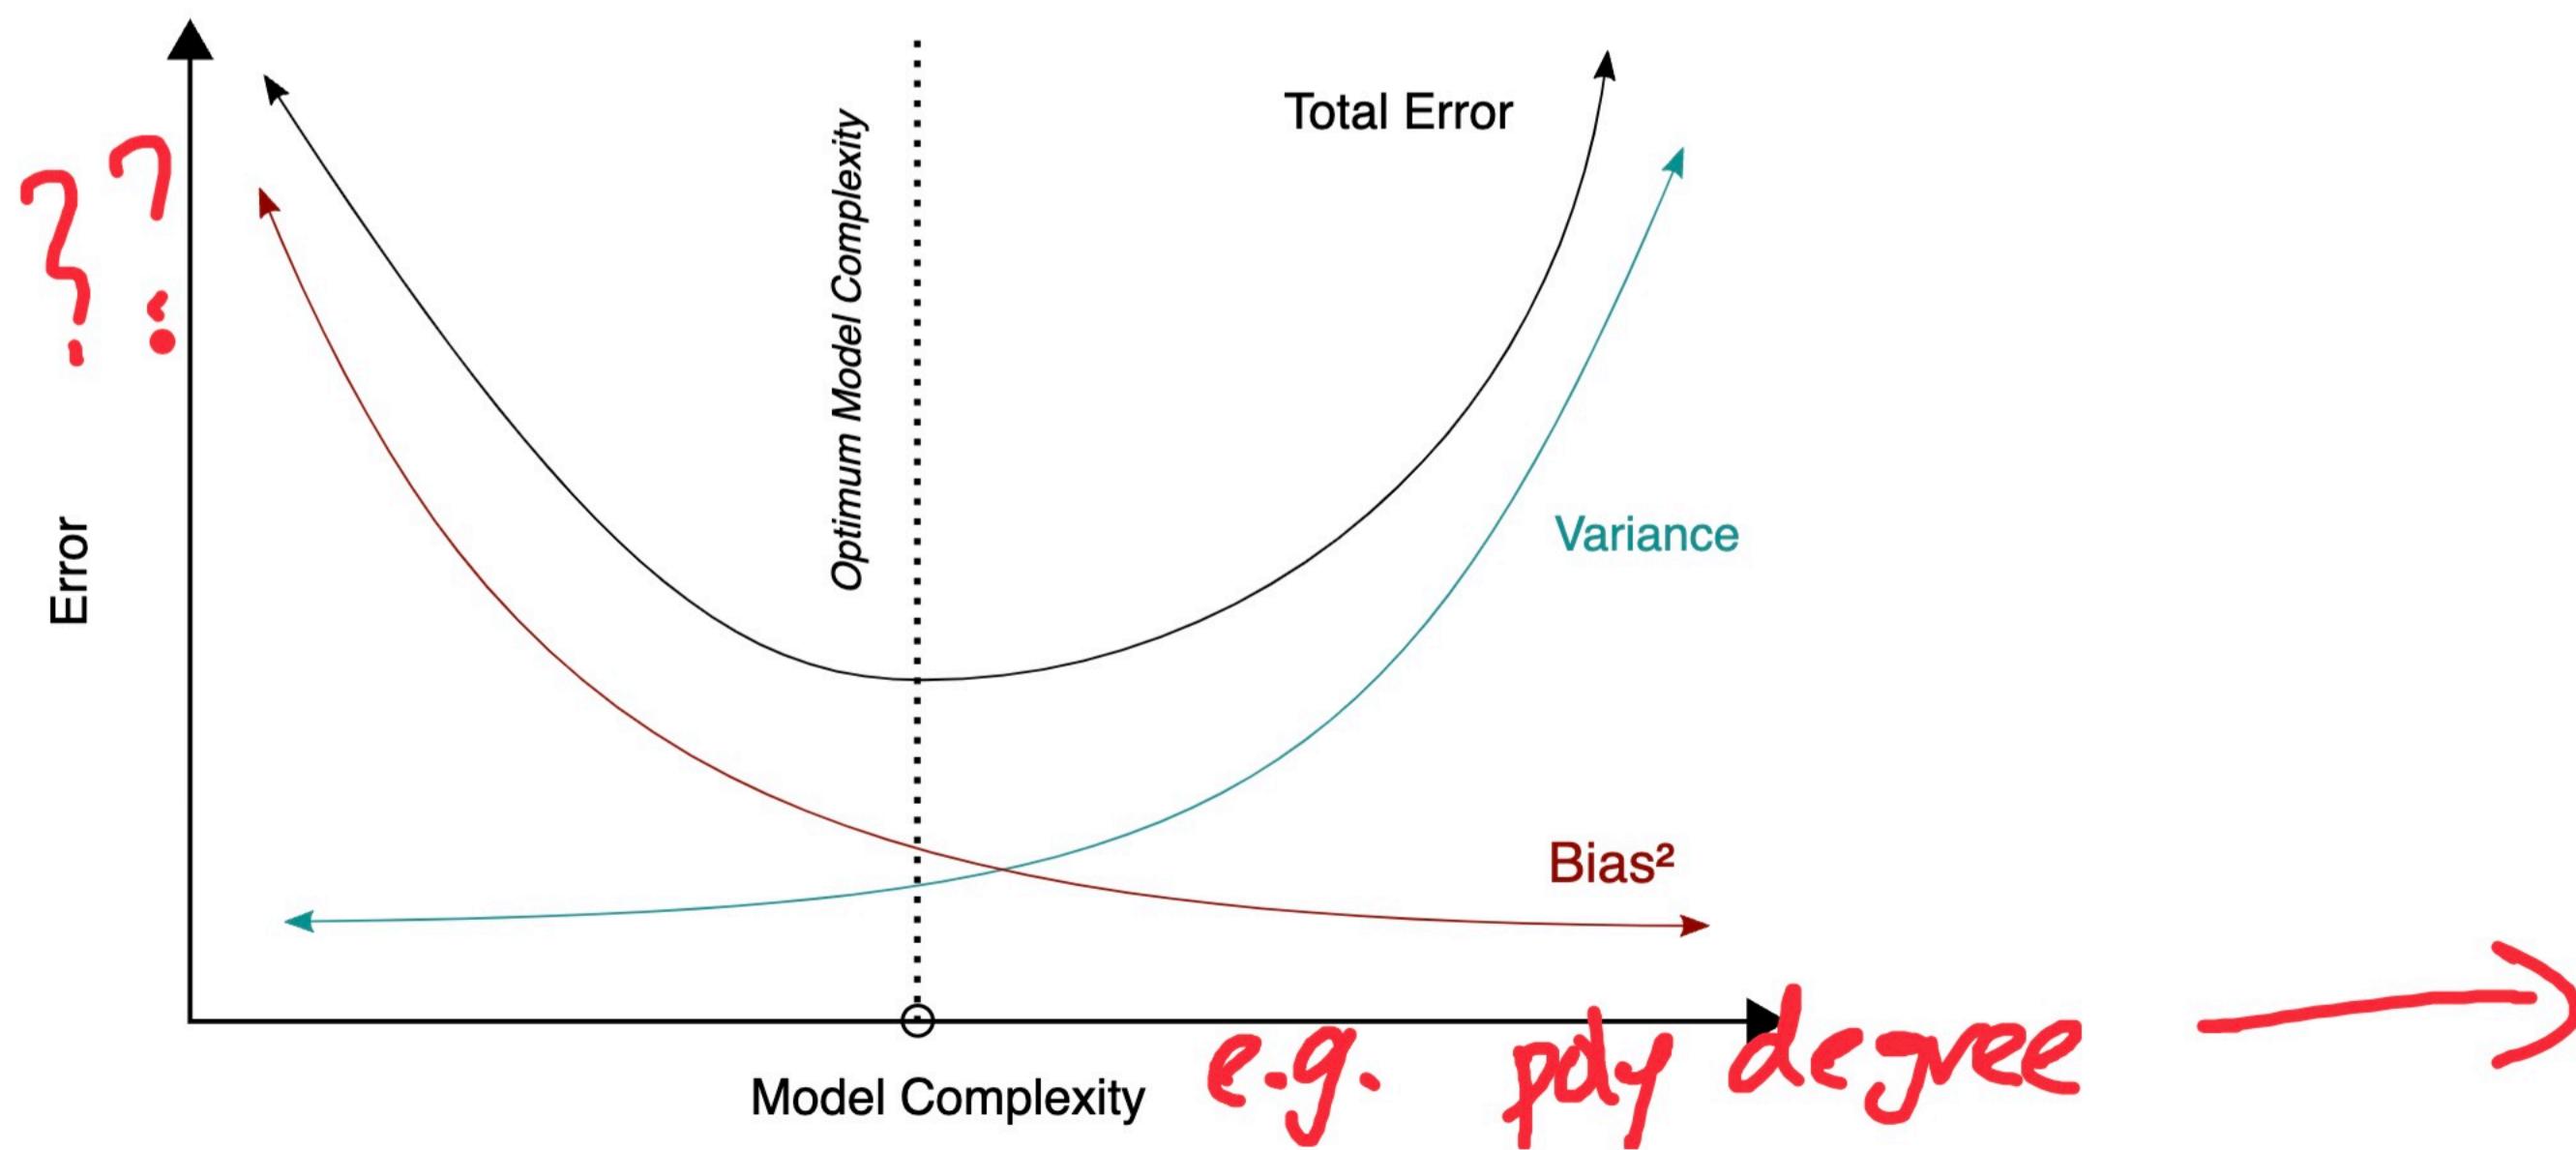
$$\mathbb{E}[y_{\text{new}} - H(\vec{x}_{\text{new}})]^2 \approx \frac{1}{n} \sum_{i=1}^n (y_i - H(\vec{x}_i))^2 = R(H)$$



$$R_{\text{sq}}(h) = \frac{1}{n} \sum_{i=1}^n (y_i - h)^2 = \underbrace{\frac{1}{n} \sum_{i=1}^n (y_i - \bar{y})^2}_{\text{variance of } y} + (\bar{y} - h)^2$$

- **Key takeaway:** If we care about minimizing (empirical) risk, we can equivalently try to minimize both model bias and model variance.
- As model variance increases, model bias tends to decrease, and vice versa.
That is, there is a **tradeoff** between bias and variance:

how do we choose model complexity ???



Out[18]:

	x	y
0	-2.00	-6.00
1	-1.95	-7.33
2	-1.90	-9.18
...
97	2.90	25.75
98	2.95	22.40
99	3.00	32.47

100 rows × 2 columns

In [19]:

```
1 X = sample_1[['x']]
2 y = sample_1['y']
3 # We don't have to choose 0.25.
4 # We also don't have to set a random_state;
5 # we've done this so that we get the same results in lecture every time.
6 train_test_split(X, y, random_state=23)
```

good for reproducibility

Out[19]: Ellipsis

```
4 # We also don't have to set a random_state;  
5 # we've done this so that we get the same results in lecture every time.  
6 X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=23)
```

- Before proceeding, let's check the sizes of `X_train` and `X_test`.

```
In [24]: 1 print('Rows in X_train:', X_train.shape[0])  
2 display(X_train.head())  
3 print('Rows in X_test:', X_test.shape[0])  
4 display(X_test.head())
```

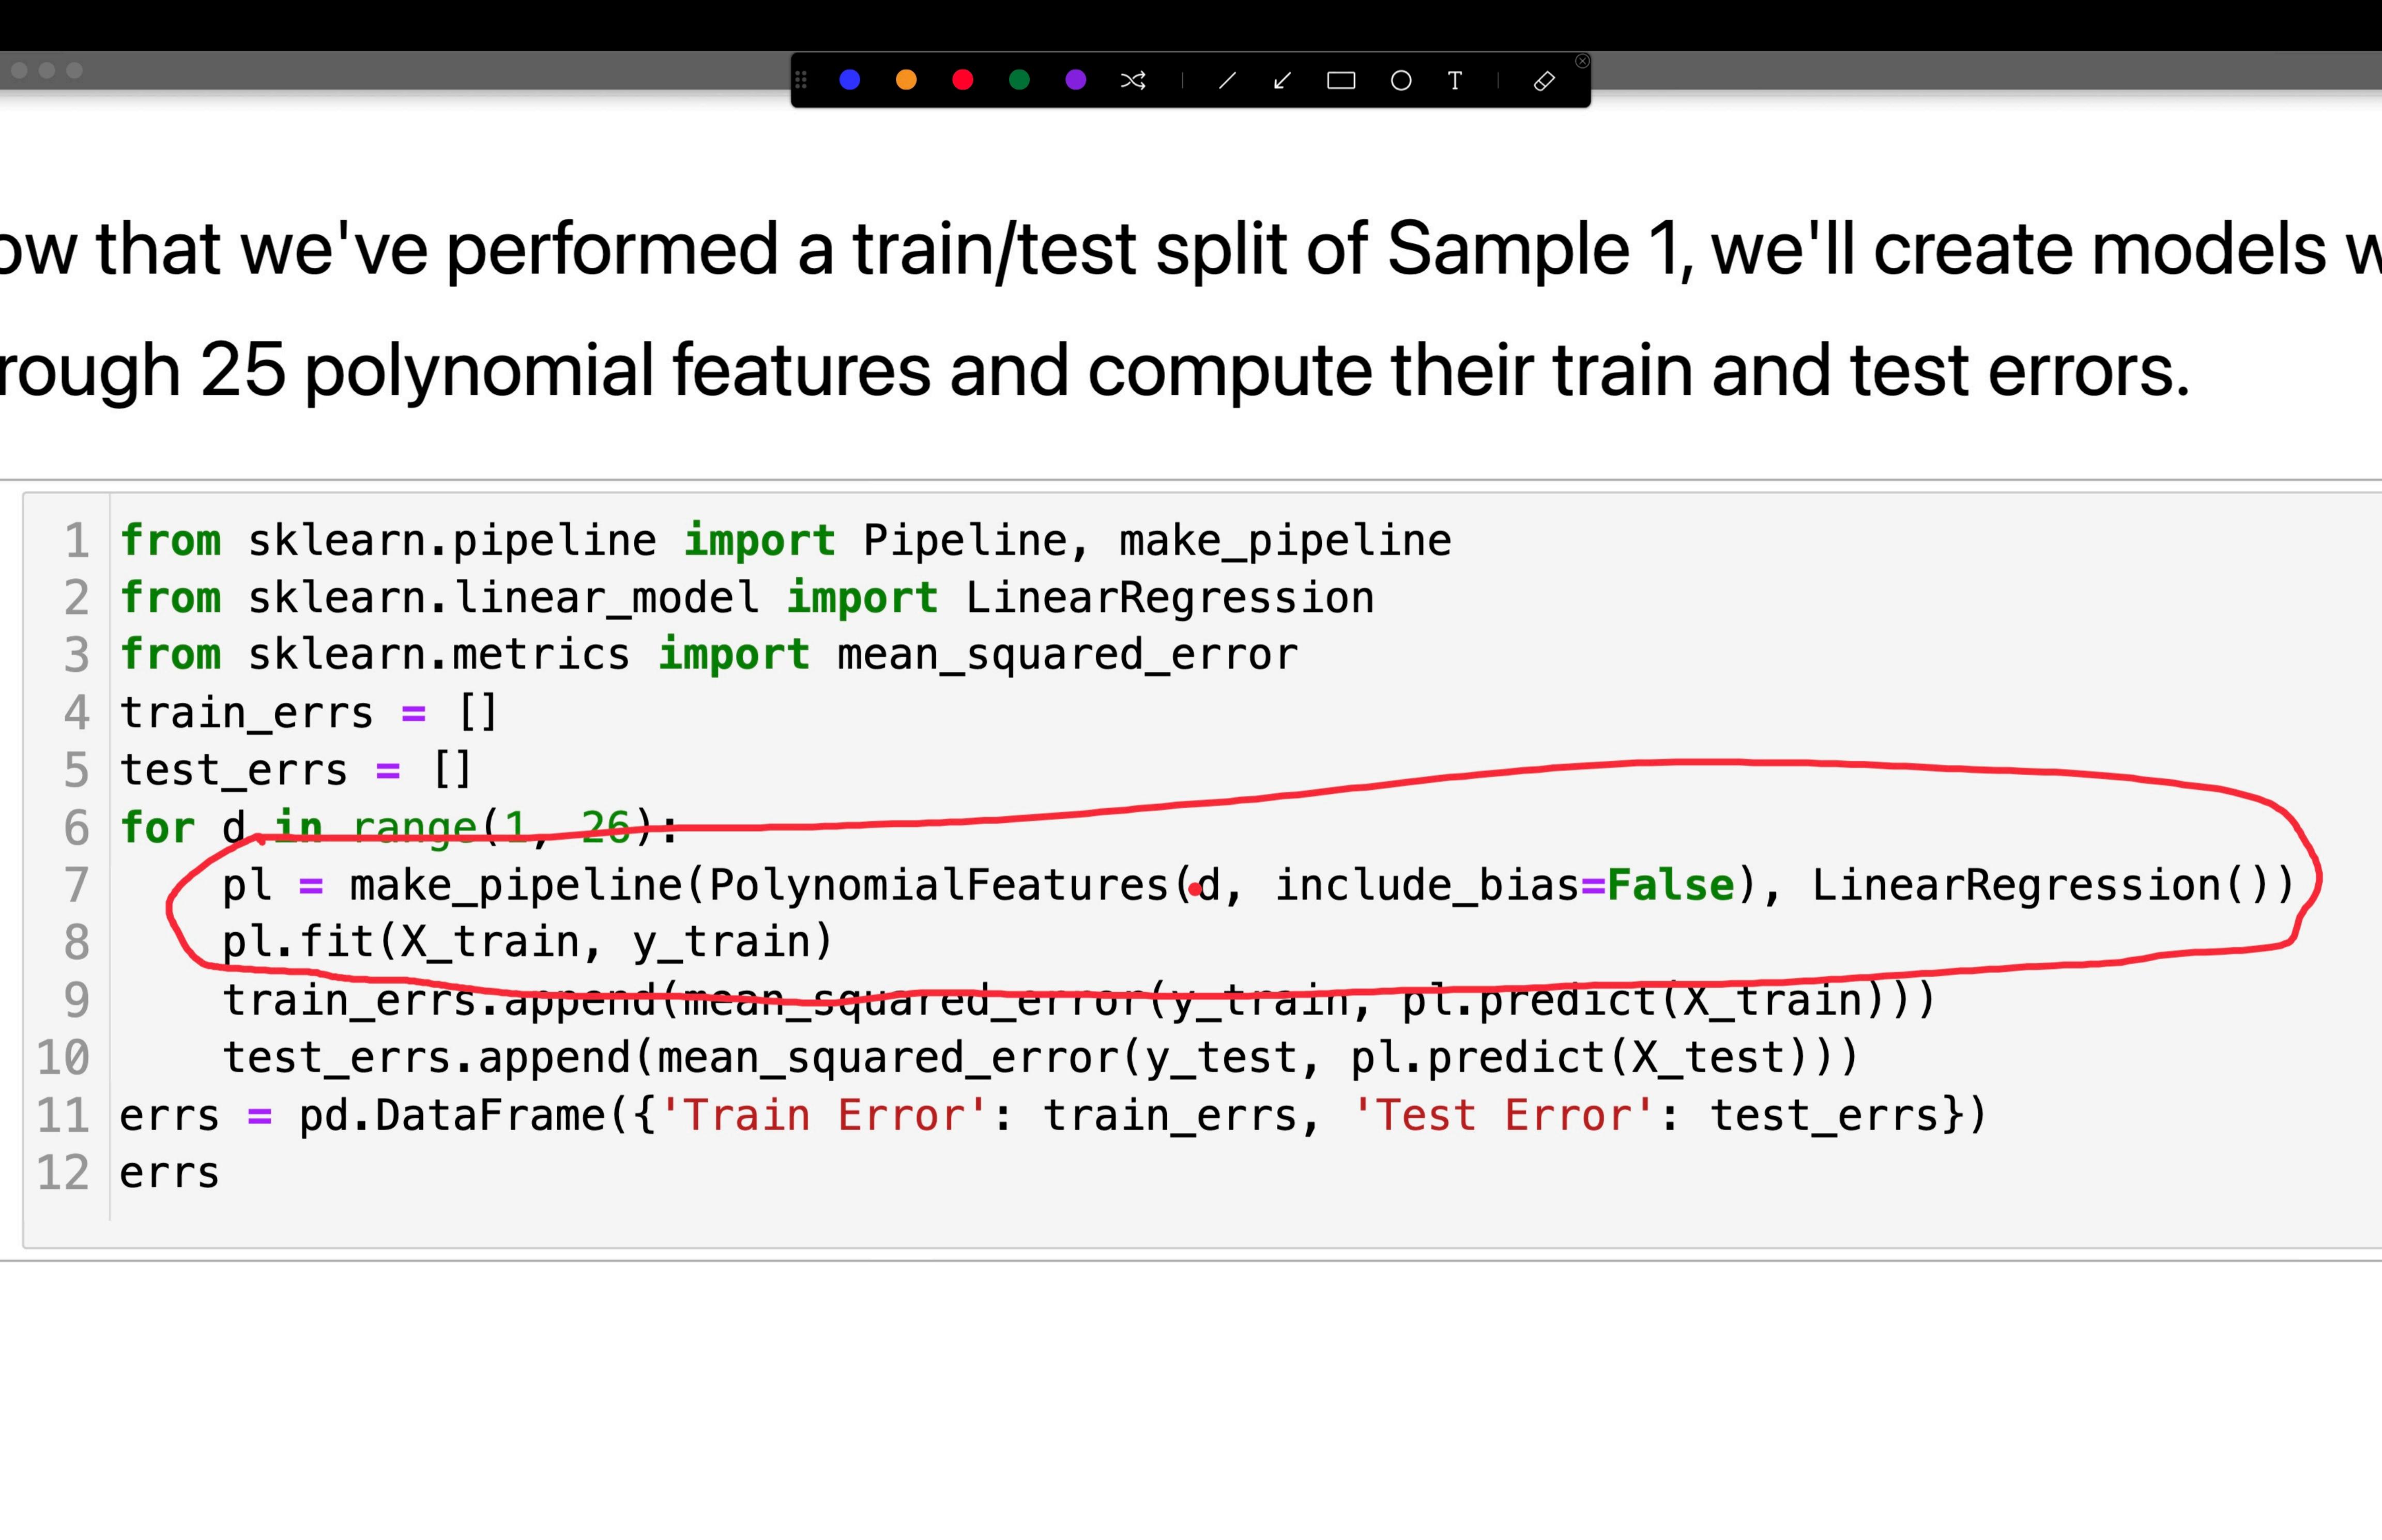
Rows in `X_train`: 75

x
4 -1.80
53 0.68
5 -1.75
81 2.09
98 2.95

randomly chosen..

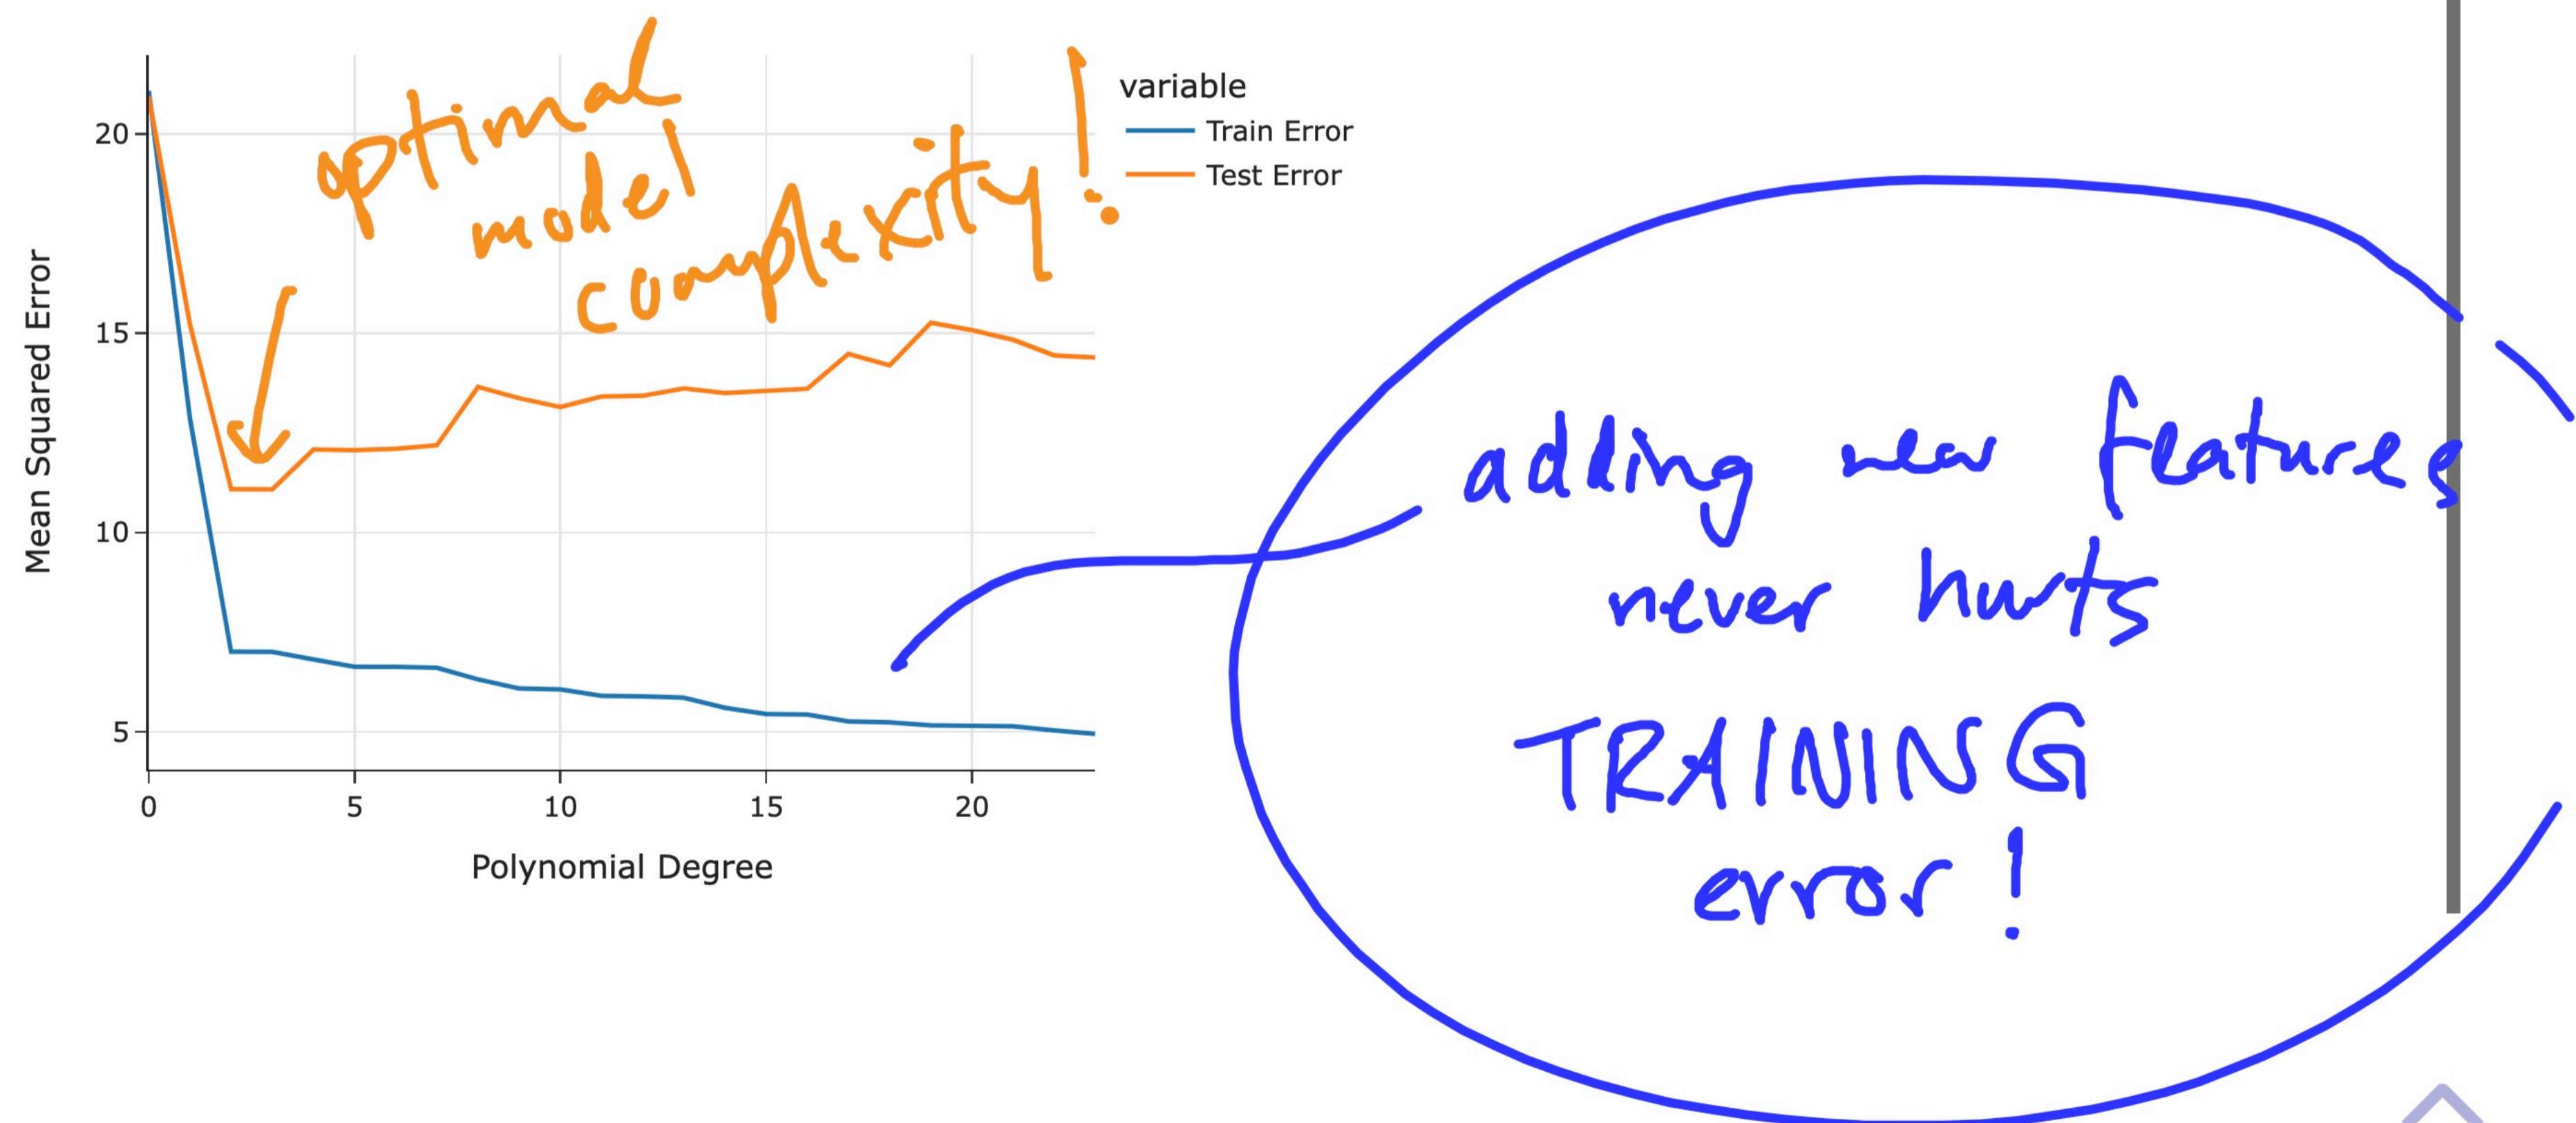
Rows in `X_test`: 25

x
26 -0.69
80 2.04
82 2.14
68 1.43
77 1.89



- Let's look at the plots of training error vs. degree and test error vs. degree.

```
In [26]: 1 fig = px.line(errs.iloc[:-1])  
2 fig.update_layout(showlegend=True, xaxis_title='Polynomial Degree', yaxis_title='Me
```



GridSearchCV

- Let's use k -fold cross-validation to choose a polynomial degree that best generalizes to unseen data.

As before, we'll choose our polynomial degree from the list [1, 2, ..., 25].

- GridSearchCV takes in:

- an **un-fit** instance of an estimator, and
- a **dictionary** of hyperparameter values to try,

and performs k -fold cross-validation to find the **combination of hyperparameters** with the best average validation performance.

```
In [29]: 1 from sklearn.model_selection import GridSearchCV  
2 GridSearchCV?
```

- Why do you think it's called "grid search"?

for each of 20 hyperparams:
train $10 \times (9 \times$ including $X.iloc[0]$,
 $\backslash x$ with n valid)

Question 🤔 (Answer at practicaldsc.org/q)

- Suppose you have a training dataset with ~~1000~~ rows.
- You want to decide between 20 hyperparameters for a particular model.
- To do so, you perform 10-fold cross-validation.
- **How many times is the first row in the training dataset (`X.iloc[0]`) used for training a model?**

$$S_t : 20 \times 9 = \boxed{180}$$

```

13 results
14 .sort_values('Average Training MSE')
15 .plot(kind='barh', barmode='group', width=1000)
16 .update_layout(xaxis_title='Mean Squared Error', yaxis_title='Model')
17 )
18 commute_models_summarized

```

