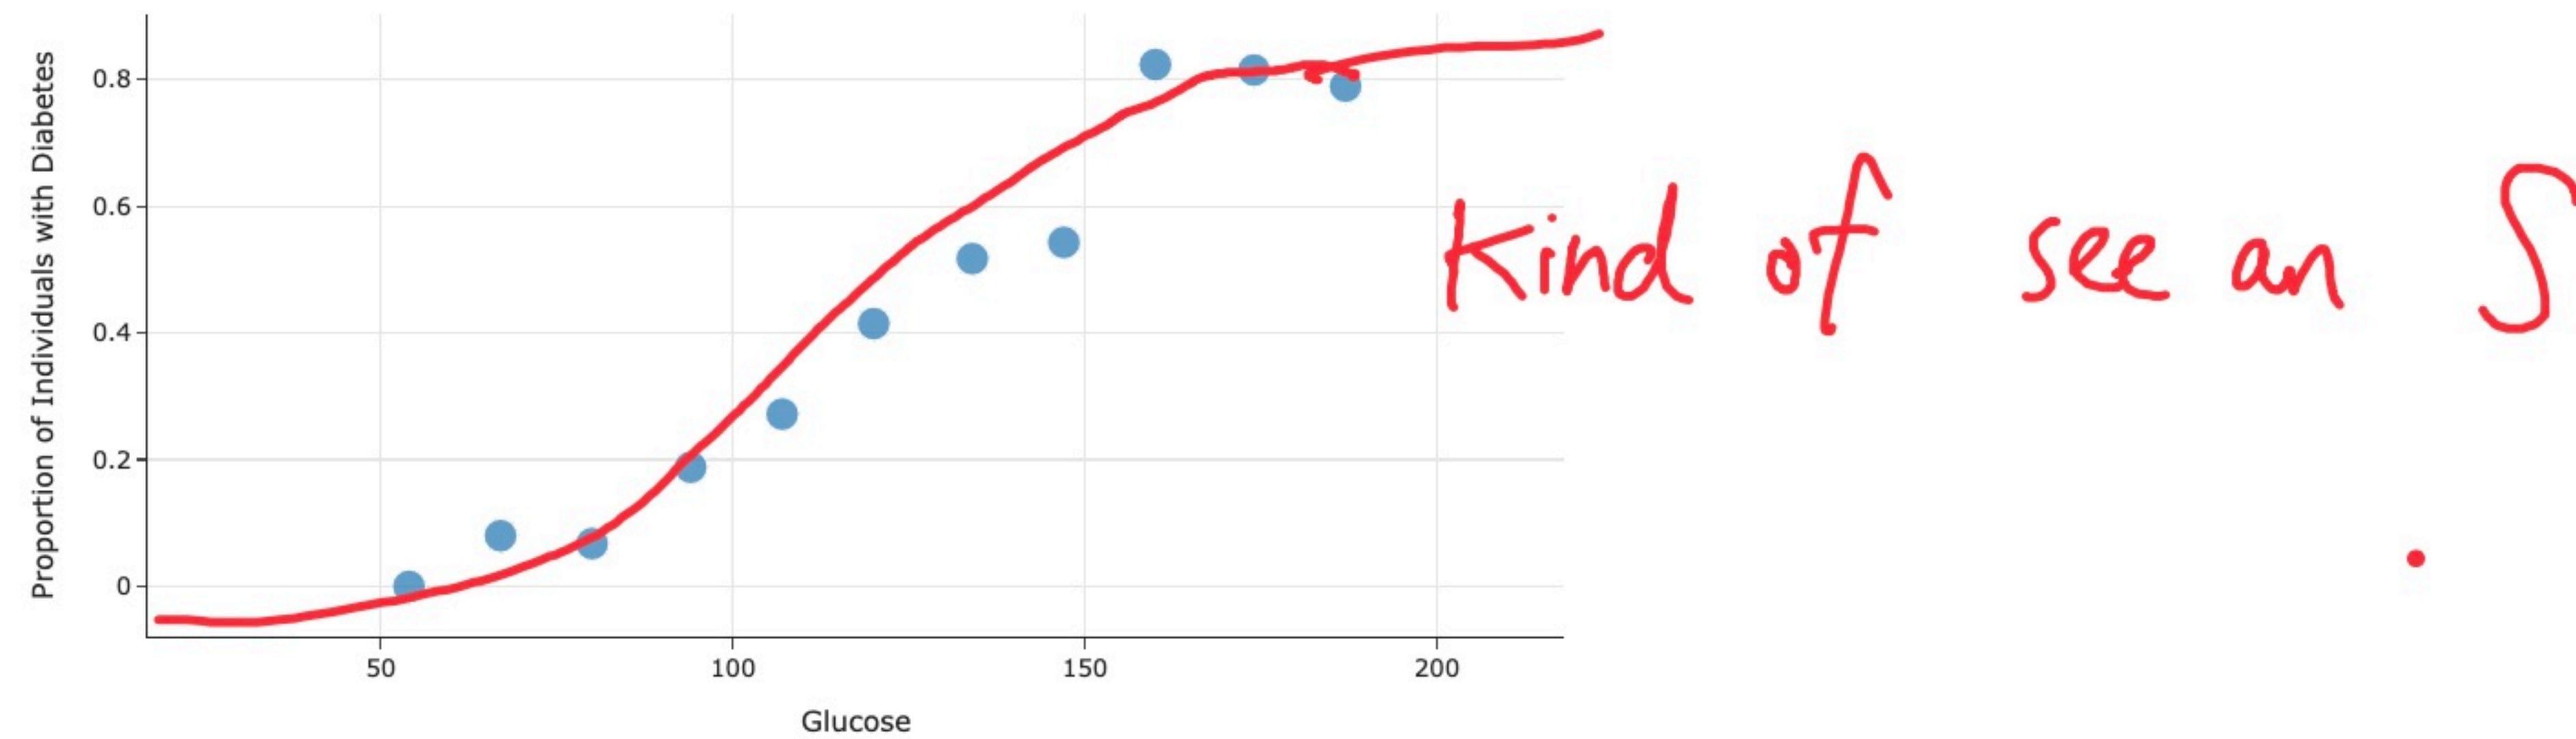


- Another approach we could try is to:

- Place 'Glucose' values into **bins**, e.g. 50 to 55, 55 to 60, 60 to 65, etc.
- Within each bin, compute the proportion of patients in the training set who had diabetes.

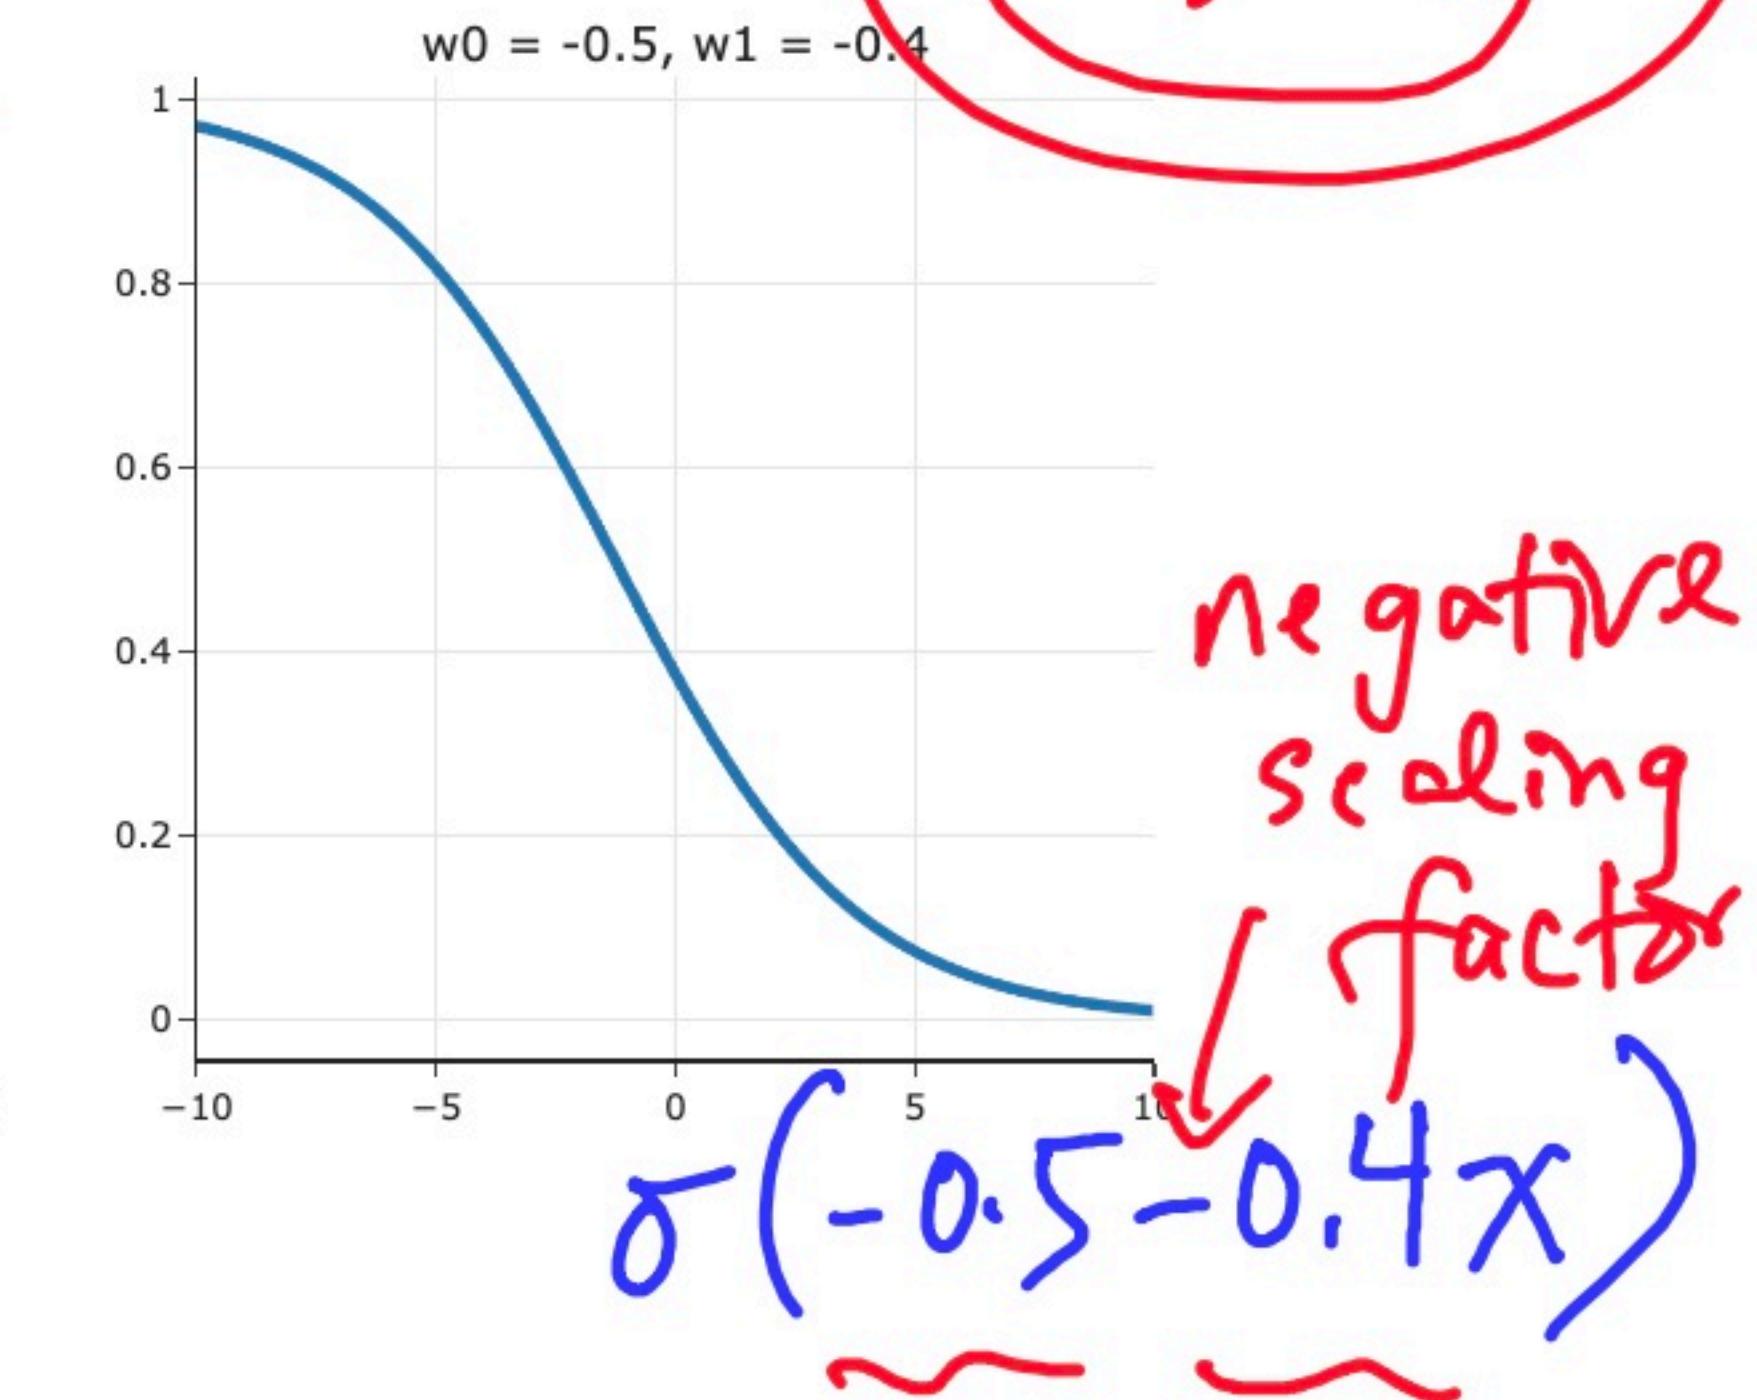
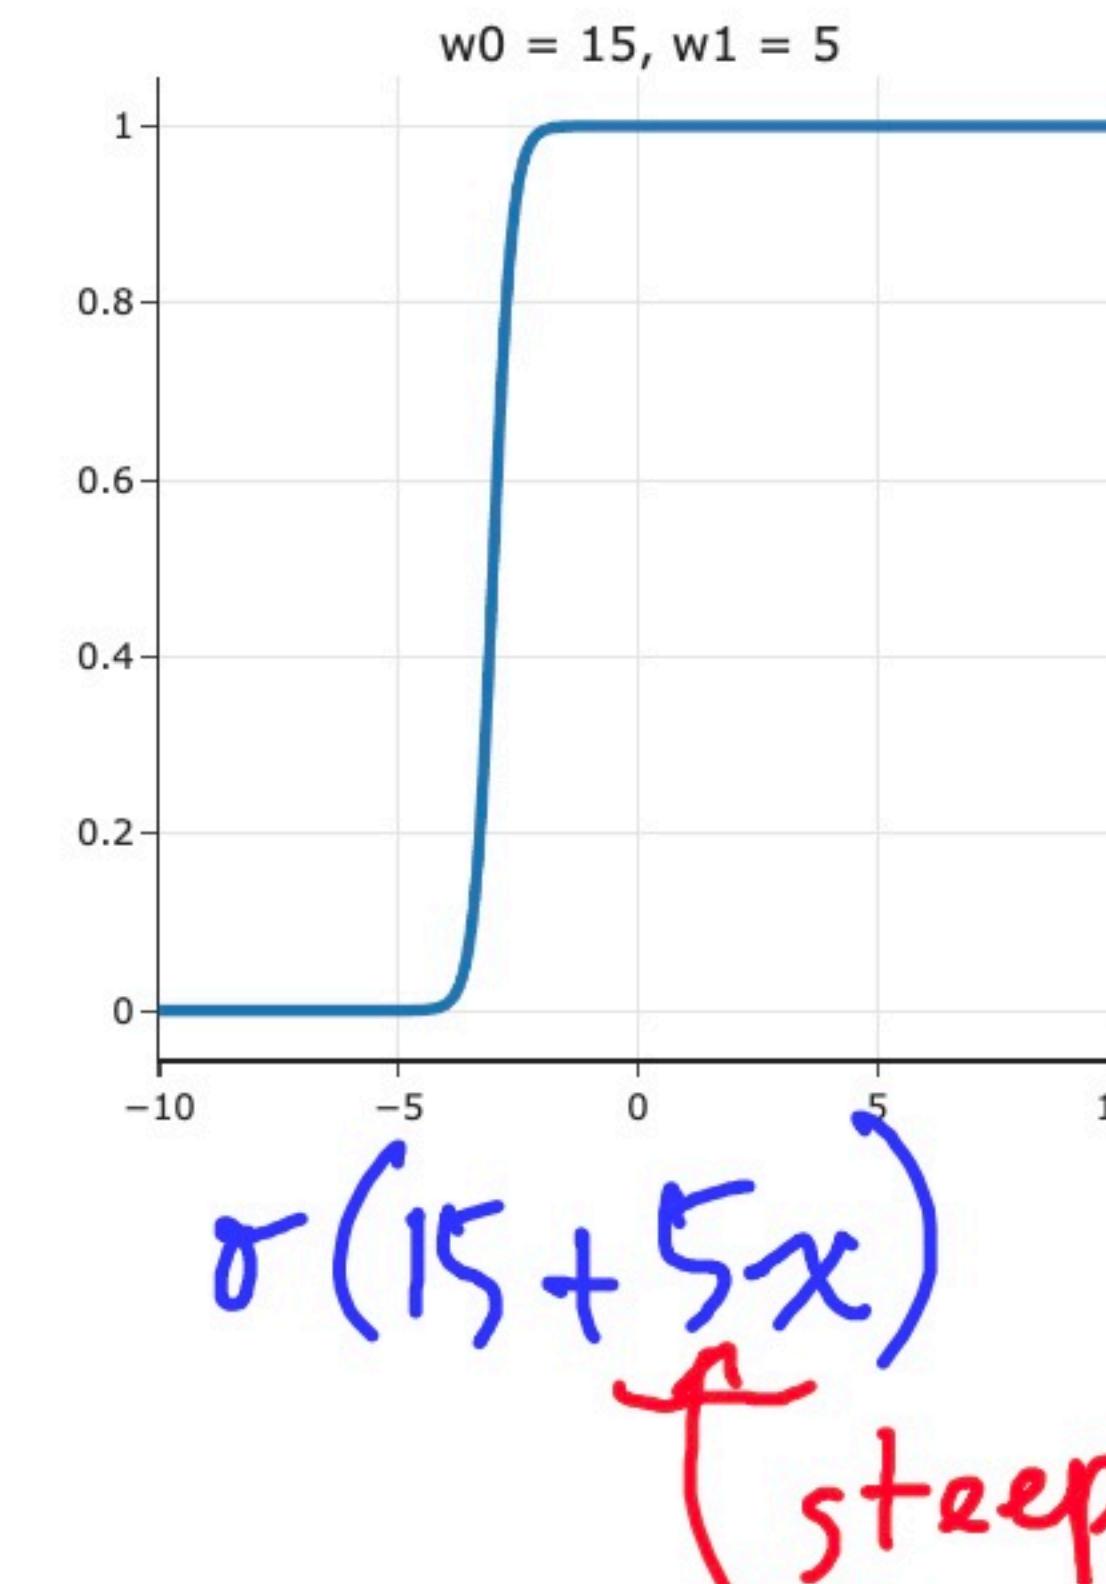
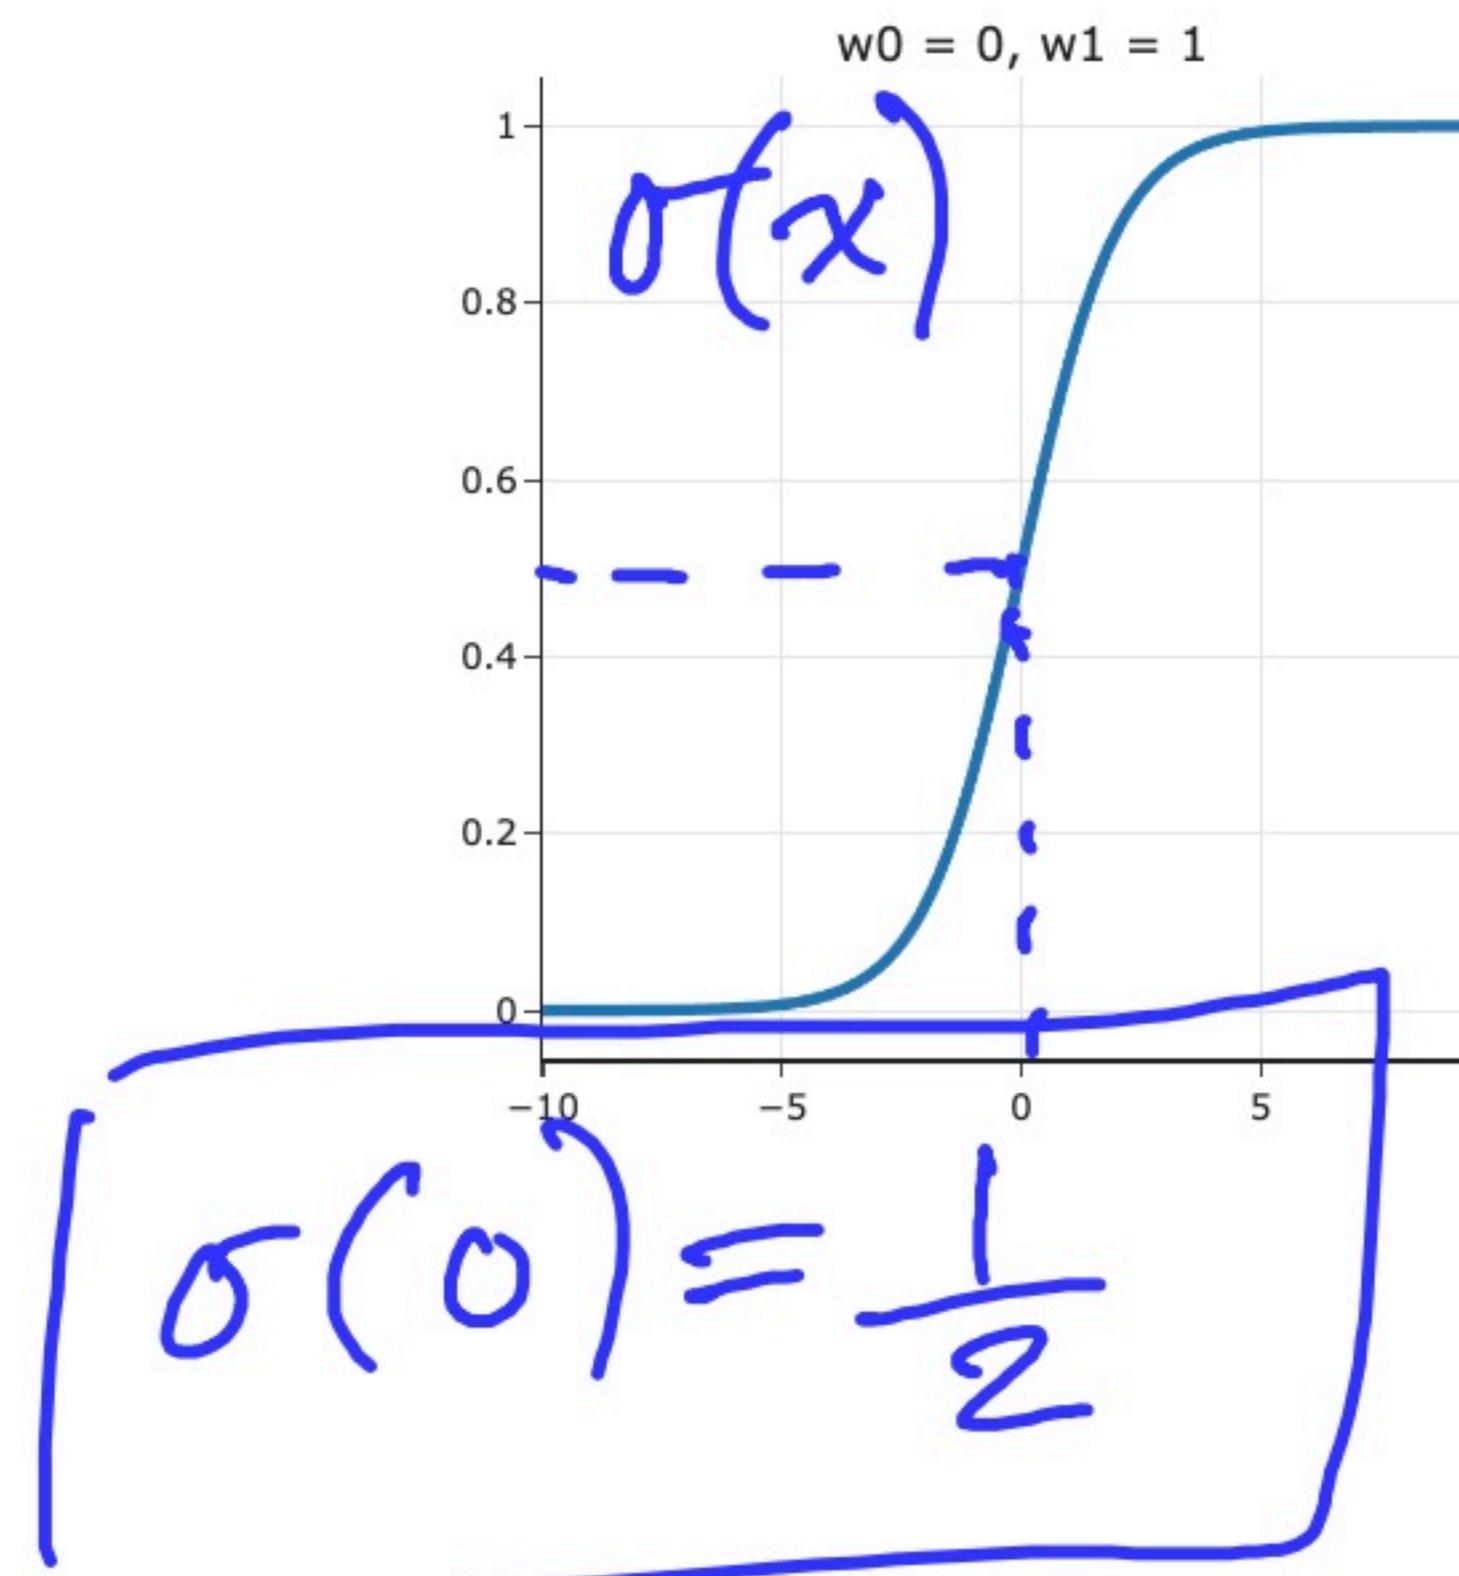
In [6]:

```
1 # Take a look at the source code in lec22_util.py to see how we did this!
2 # We've hidden a lot of the plotting code in the notebook to make it cleaner.
3 util.make_prop_plot(X_train, y_train)
```



- Below, we'll look at the shape of $y = \sigma(\underline{w_0} + \underline{w_1}x)$ for different values of w_0 and w_1 .
 - w_0 controls the position of the curve on the x -axis.
 - w_1 controls the "steepness" of the curve.

In [7]: 1 util.show_three_sigmoid()



σ : Sigmoid curve

sigmoid's logistic.

Logistic regression

- Logistic **regression** is a linear **classification** technique that builds upon linear regression.
It is **not** called logistical regression!

- It models the probability of belonging to class 1, given a feature vector:

$$P(y_i = 1 | \vec{x}_i) = \sigma(w_0 + w_1 x_i^{(1)} + w_2 x_i^{(2)} + \dots + w_d x_i^{(d)}) = \sigma(\vec{w} \cdot \text{Aug}(\vec{x}_i))$$

linear regression model

guaranteed that

ed that



Logistic regression

- Logistic **regression** is a linear **classification** technique that builds upon linear regression.
It is **not** called logistical regression!

- It models **the probability of belonging to class 1, given a feature vector:**

$$P(y_i = 1 | \vec{x}_i) = \sigma(w_0 + w_1 x_i^{(1)} + w_2 x_i^{(2)} + \dots + w_d x_i^{(d)}) = \sigma(\vec{w} \cdot \text{Aug}(\vec{x}_i))$$

linear regression model parameters

guaranteed that

$$0 < P(y_i = 1 | \vec{x}_i) < 1$$



```
In [9]: 1 from sklearn.linear_model import LogisticRegression
```

- Let's fit a `LogisticRegression` classifier. Specifically, this means we're asking `sklearn` to learn the optimal parameters w_0^* and w_1^* in:

$$P(y_i = 1 | \text{Glucose}_i) = \sigma(w_0 + w_1 \cdot \text{Glucose}_i)$$

```
In [11]: 1 model_logistic = LogisticRegression()
2 model_logistic.fit(X_train[['Glucose']], y_train)
```

```
Out[11]: LogisticRegression
```

training data

- We get a test accuracy that's roughly in line with the test accuracies of the two models we saw last class.

75% test accuracy

```
In [12]: 1 model_logistic.score(X_test[['Glucose']], y_test)
```

```
Out[12]: 0.75
```



Attempting to use squared loss

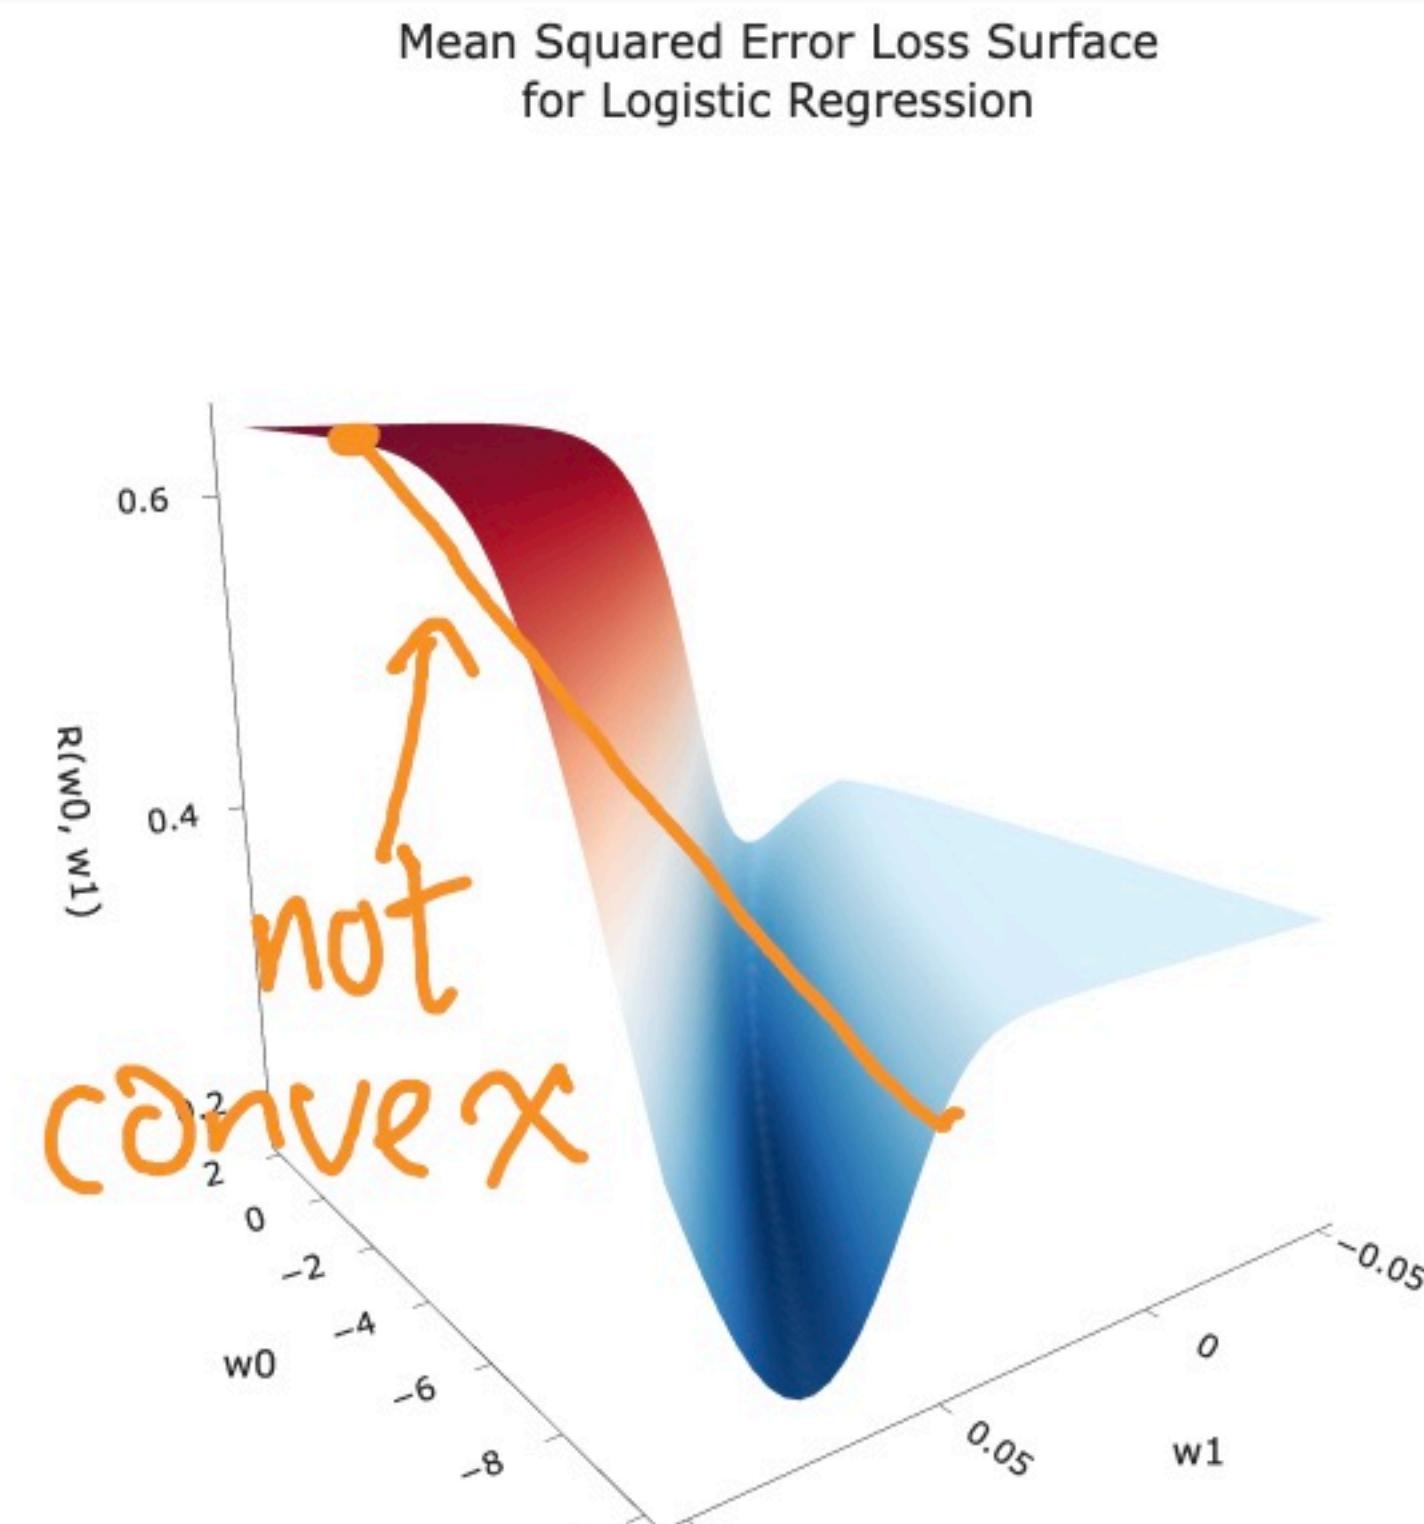
- Our default loss function has always been squared loss, so we could try and use it here.

$$R_{\text{sq}}(\vec{w}) = \frac{1}{n} \sum_{i=1}^n (y_i - \underbrace{\sigma(\vec{w} \cdot \text{Aug}(\vec{x}_i))}_{\text{logistic regression model's predictions}})^2$$



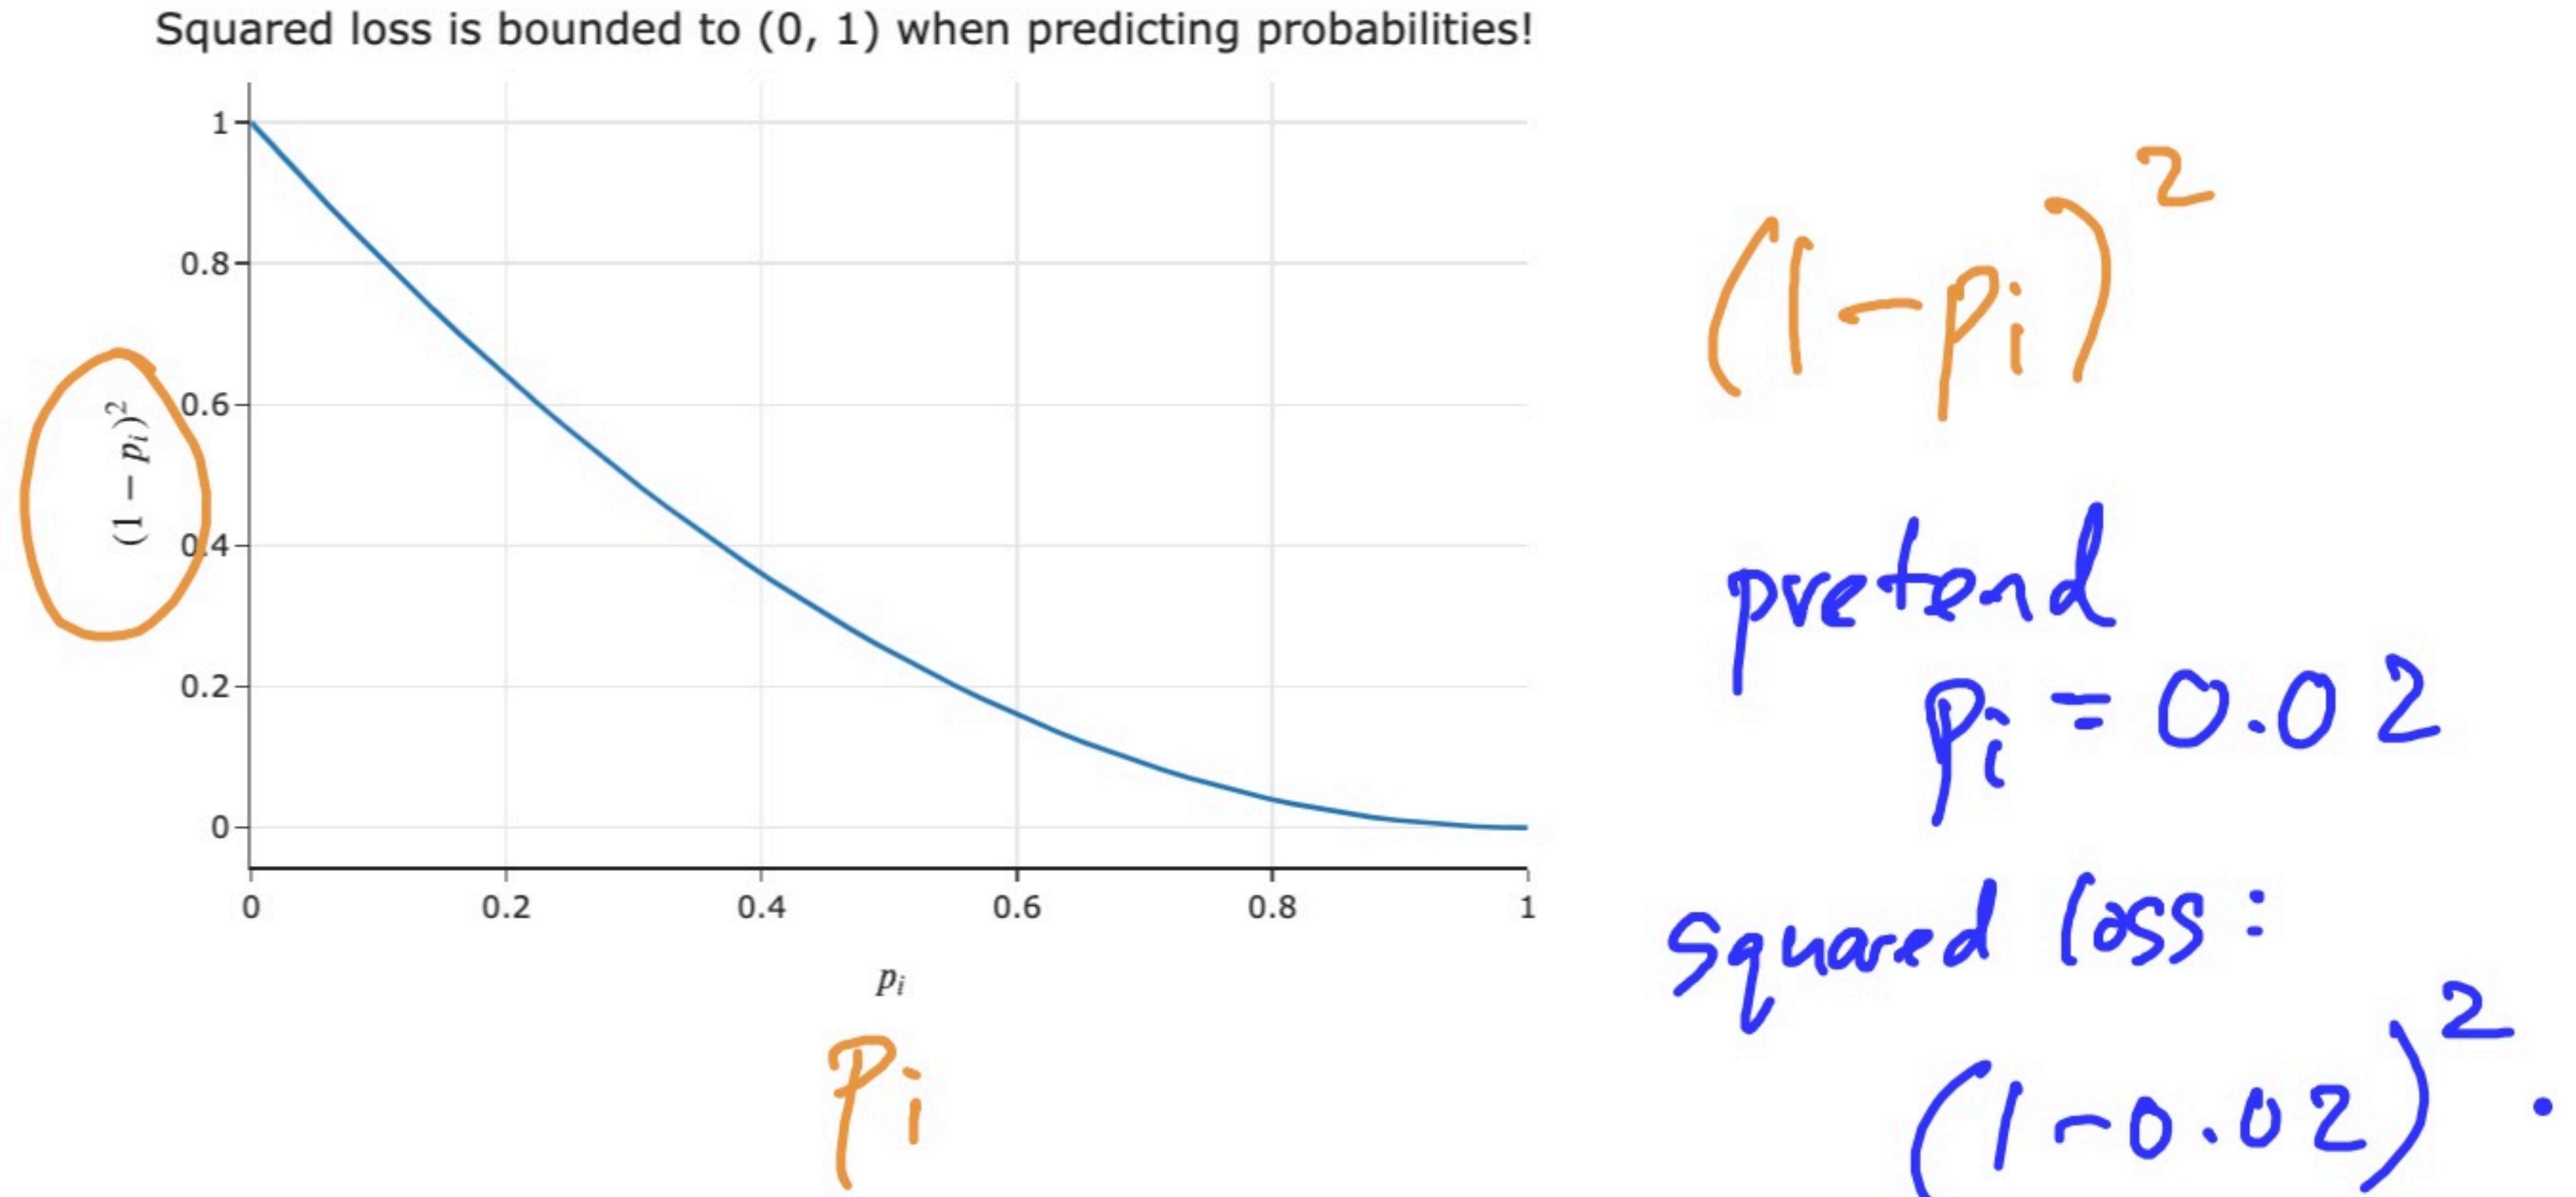
$$R_{\text{sq}}(w_0, w_1) = \frac{1}{n} \sum_{i=1}^n \left| y_i - \sigma(w_0 + w_1 \underbrace{x_i}_{\text{Glucose}_i}) \right|$$

In [18]: 1 util.show_logistic_mse_surface(X_train, y_train)



- Suppose $y_i = 1$. Then, the graph of the squared loss of the prediction p_i is below.

In [19]: 1 util.show_squared_loss_individual()

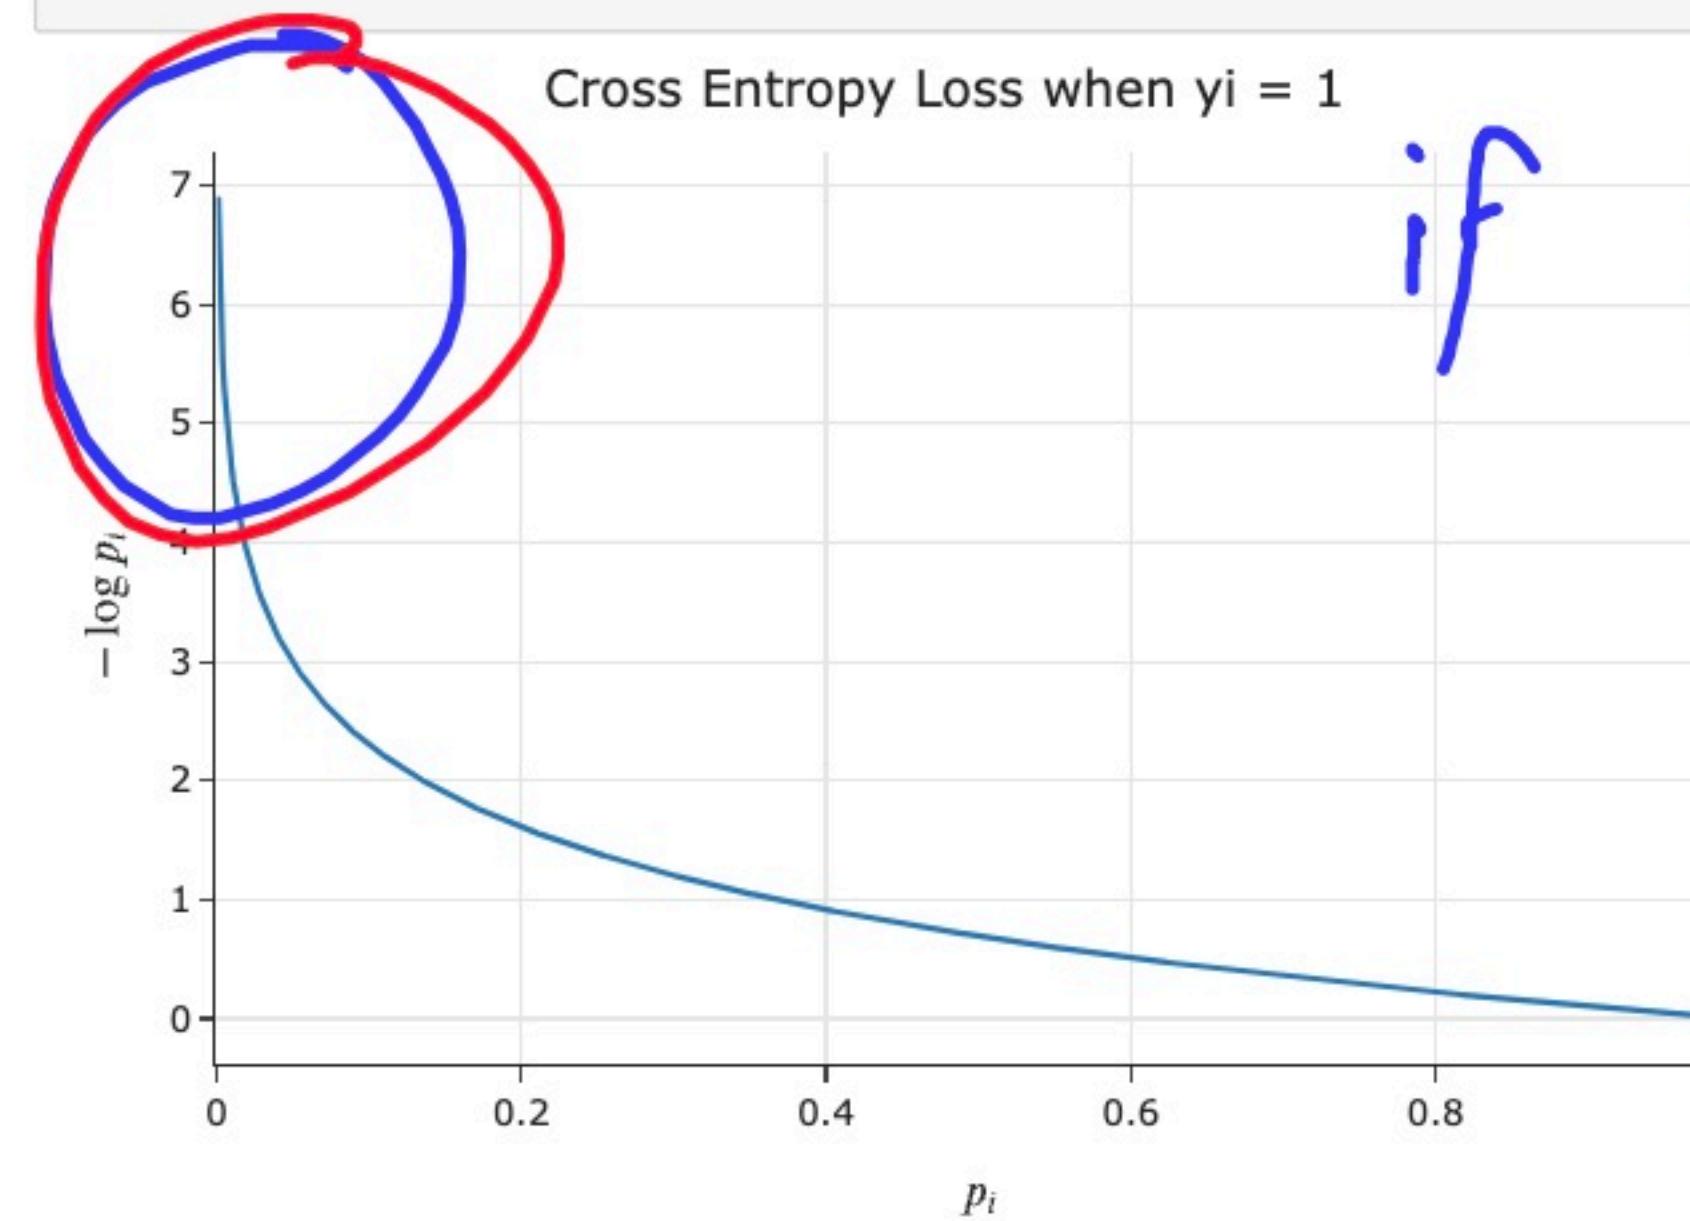


predicted probability, then:

$$L_{\text{ce}}(y_i, p_i) = \begin{cases} -\log(p_i) & \text{if } y_i = 1 \\ -\log(1 - p_i) & \text{if } y_i = 0 \end{cases}$$

- Note that in the two cases – $y_i = 1$ and $y_i = 0$ – the cross-entropy loss function resembles squared loss, but is unbounded when the predicted probabilities p_i are far from y_i .

In [20]: 1 util.show_ce_loss_individual_1()



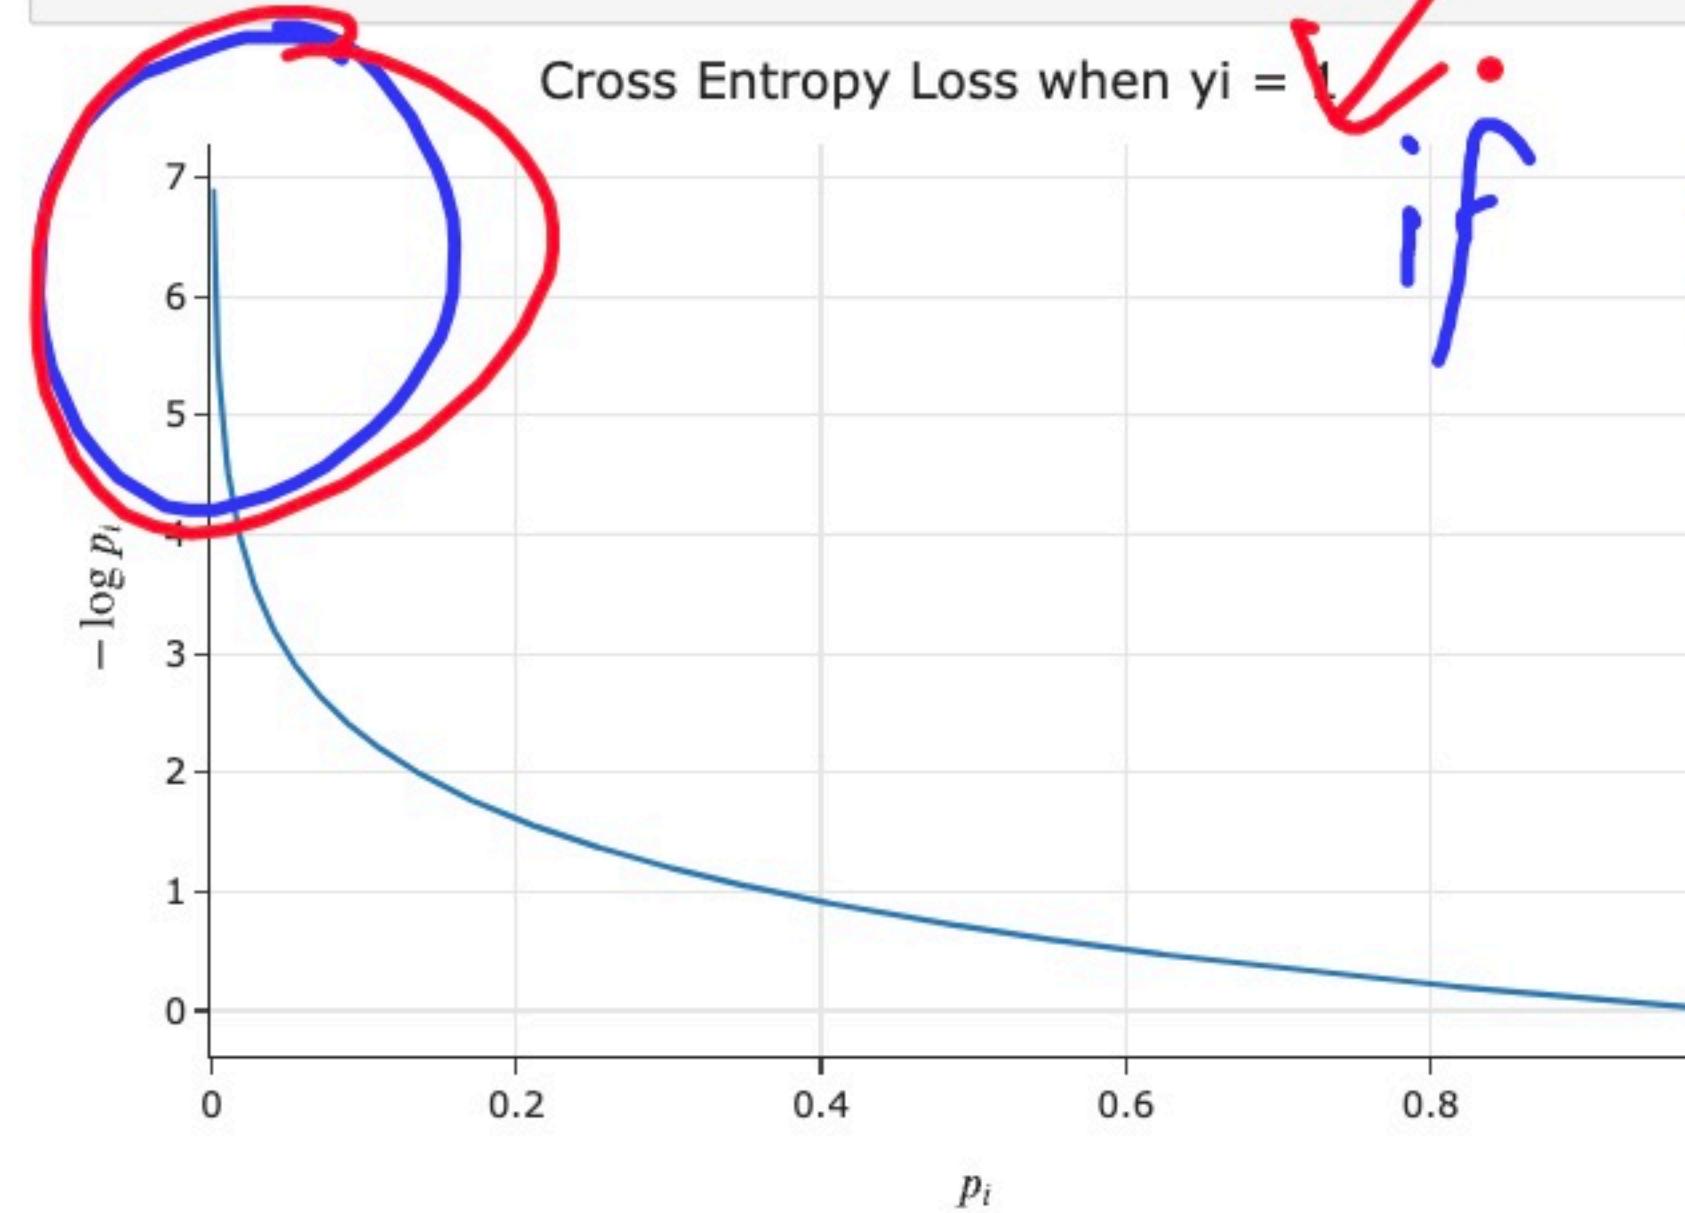
if patient i has diabetes.
but we predict
 p_i = Probability of diabetes
to loss be is small very high!

predicted probability, then:

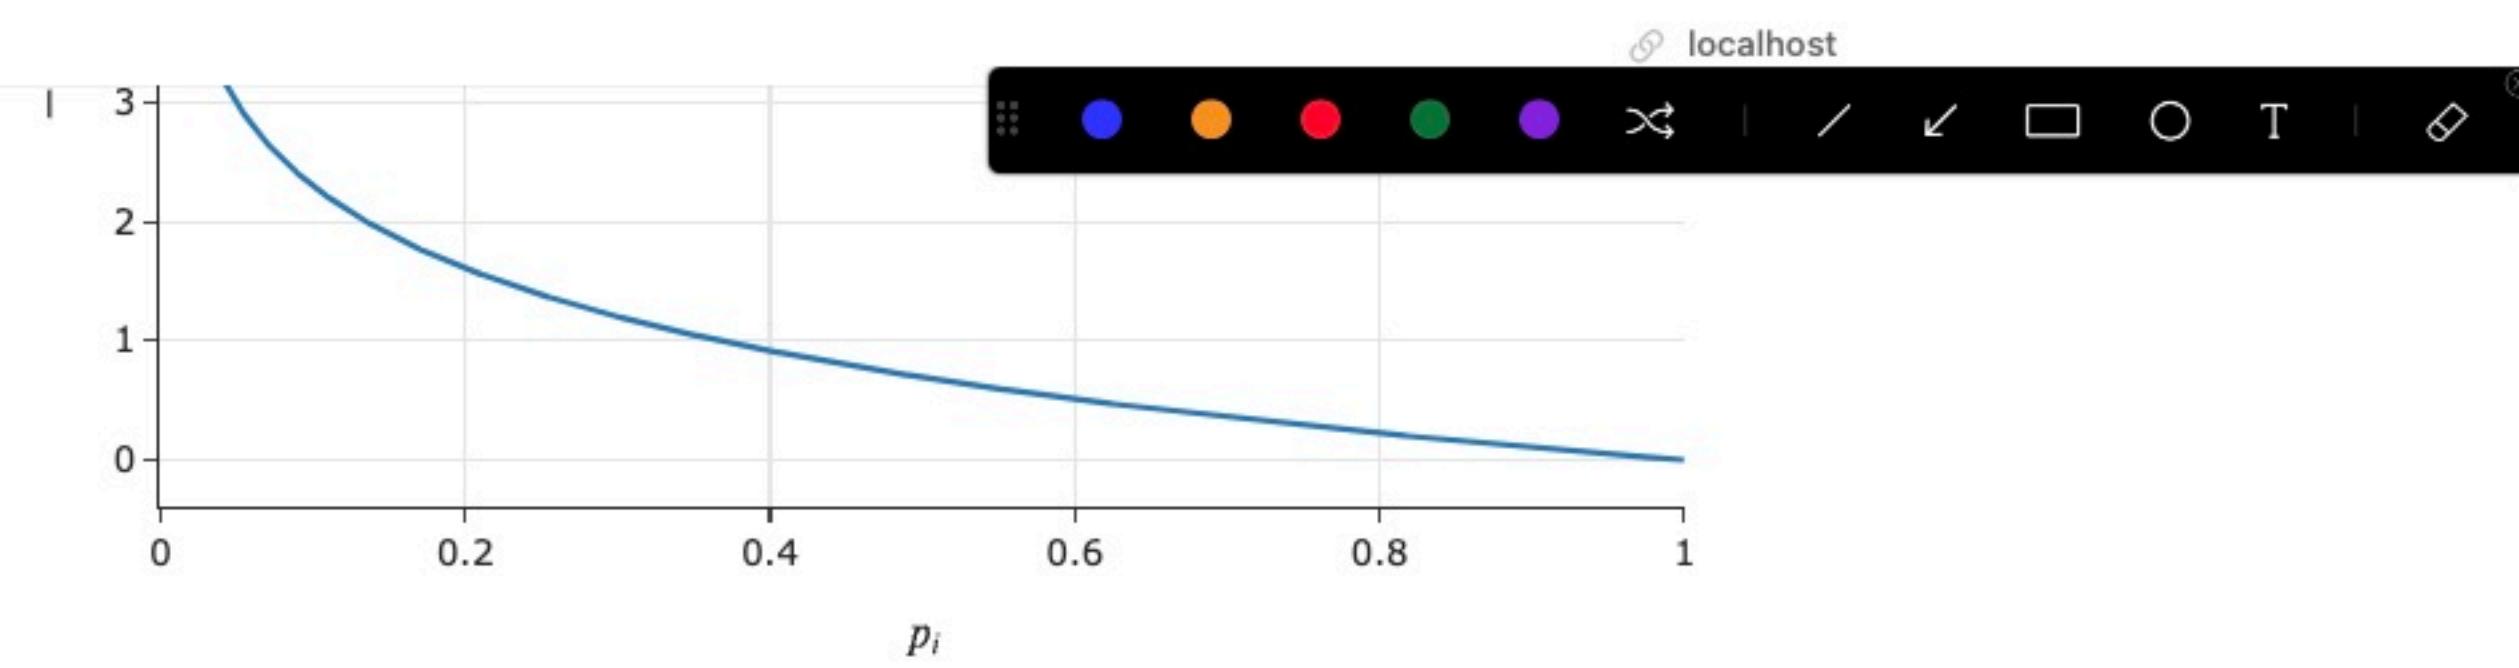
$$L_{\text{ce}}(y_i, p_i) = \begin{cases} -\log(p_i) & \text{if } y_i = 1 \\ -\log(1 - p_i) & \text{if } y_i = 0 \end{cases}$$

- Note that in the two cases – $y_i = 1$ and $y_i = 0$ – the cross-entropy loss function resembles squared loss, but is unbounded when the predicted probabilities p_i are far from y_i .

In [20]: 1 util.show_ce_loss_individual_1()



if patient i has diabetes
but we predict
 p_i = Probability of diabetes
to loss be is small very high!



$$L_{ce}(y_i, p_i) = \begin{cases} -\log(p_i) & \text{if } y_i = 1 \\ -\log(1 - p_i) & \text{if } y_i = 0 \end{cases}$$

In [21]: 1 util.show_ce_loss_individual_0()



mirror image of
each other .



A non-piecewise definition of cross-entropy loss

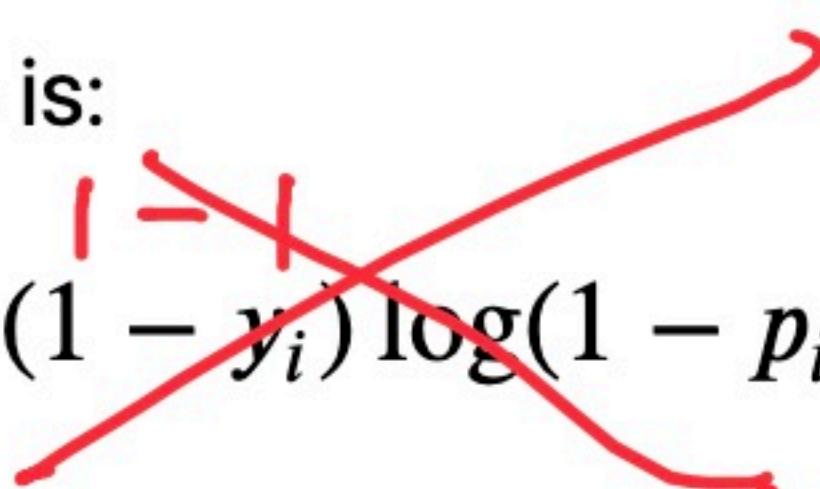
- We can define the cross-entropy loss function piecewise. If y_i is an observed value and p_i is a predicted **probability**, then:

$$L_{\text{ce}}(y_i, p_i) = \begin{cases} -\log(p_i) & \text{if } y_i = 1 \\ -\log(1 - p_i) & \text{if } y_i = 0 \end{cases}$$

- An equivalent formulation of L_{ce} that isn't piecewise is:

$$L_{\text{ce}}(y_i, p_i) = -(y_i \log p_i + (1 - y_i) \log(1 - p_i))$$

if $y_i = 1$:







A non-piecewise definition of cross-entropy loss

- We can define the cross-entropy loss function piecewise. If y_i is an observed value and p_i is a predicted **probability**, then:

$$L_{\text{ce}}(y_i, p_i) = \begin{cases} -\log(p_i) & \text{if } y_i = 1 \\ -\log(1 - p_i) & \text{if } y_i = 0 \end{cases}$$

- An equivalent formulation of L_{ce} that isn't piecewise is:

$$L_{\text{ce}}(y_i, p_i) = -(\cancel{y_i \log p_i} + (1 - y_i) \log(1 - p_i))$$

if $y_i = 0$.





Average cross-entropy loss

- **Cross-entropy loss** for an observed value y_i and predicted **probability**

$p_i = P(y = 1 | \vec{x}_i) = \sigma(\vec{w} \cdot \text{Aug}(\vec{x}_i))$ is:

$$L_{\text{ce}}(y_i, p_i) = -(y_i \log p_i + (1 - y_i) \log(1 - p_i))$$

$$- (y_i \log \sigma(\vec{w} \cdot \text{Aug}(\vec{x}_i))) + (1 - y_i) \log(1 - \sigma(\vec{w} \cdot \text{Aug}(\vec{x}_i))) - \dots$$



$$P(y_i = 1 | \vec{x}_i) = \sigma(w_0 + w_1 x_i^{(1)} + w_2 x_i^{(2)} + \dots + w_d x_i^{(d)}) = \sigma(\vec{w} \cdot \text{Aug}(\vec{x}_i))$$

2. Choose a loss function.

cross-entropy loss

$$L_{\text{ce}}(y_i, p_i) = -(y_i \log p_i + (1 - y_i) \log(1 - p_i))$$

$$\text{where } p_i = P(y = 1 | \vec{x}_i) = \sigma(\vec{w} \cdot \text{Aug}(\vec{x}_i))$$

3. Minimize average loss to find optimal model parameters.

As we've now seen, average loss could also be regularized!

$$\begin{aligned} R_{\text{ce}}(\vec{w}) &= -\frac{1}{n} \sum_{i=1}^n (y_i \log p_i + (1 - y_i) \log(1 - p_i)) \\ &= -\frac{1}{n} \sum_{i=1}^n [y_i \log(\sigma(\vec{w} \cdot \text{Aug}(\vec{x}_i))) + (1 - y_i) \log(1 - \sigma(\vec{w} \cdot \text{Aug}(\vec{x}_i)))] \end{aligned}$$

The actual minimization here is done using numerical methods, through `sklearn`.

LogisticRegression in sklearn, revisited

- The LogisticRegression class in sklearn has a lot of hidden, default hyperparameters.

```
In [24]: 1 LogisticRegression?
```

- It performs L_2 regularization ("ridge logistic regression") **by default**. The hyperparameter for regularization strength, C , is the **inverse** of λ ; by default, it sets $C = 1$.

$$C = \frac{1}{\lambda}$$

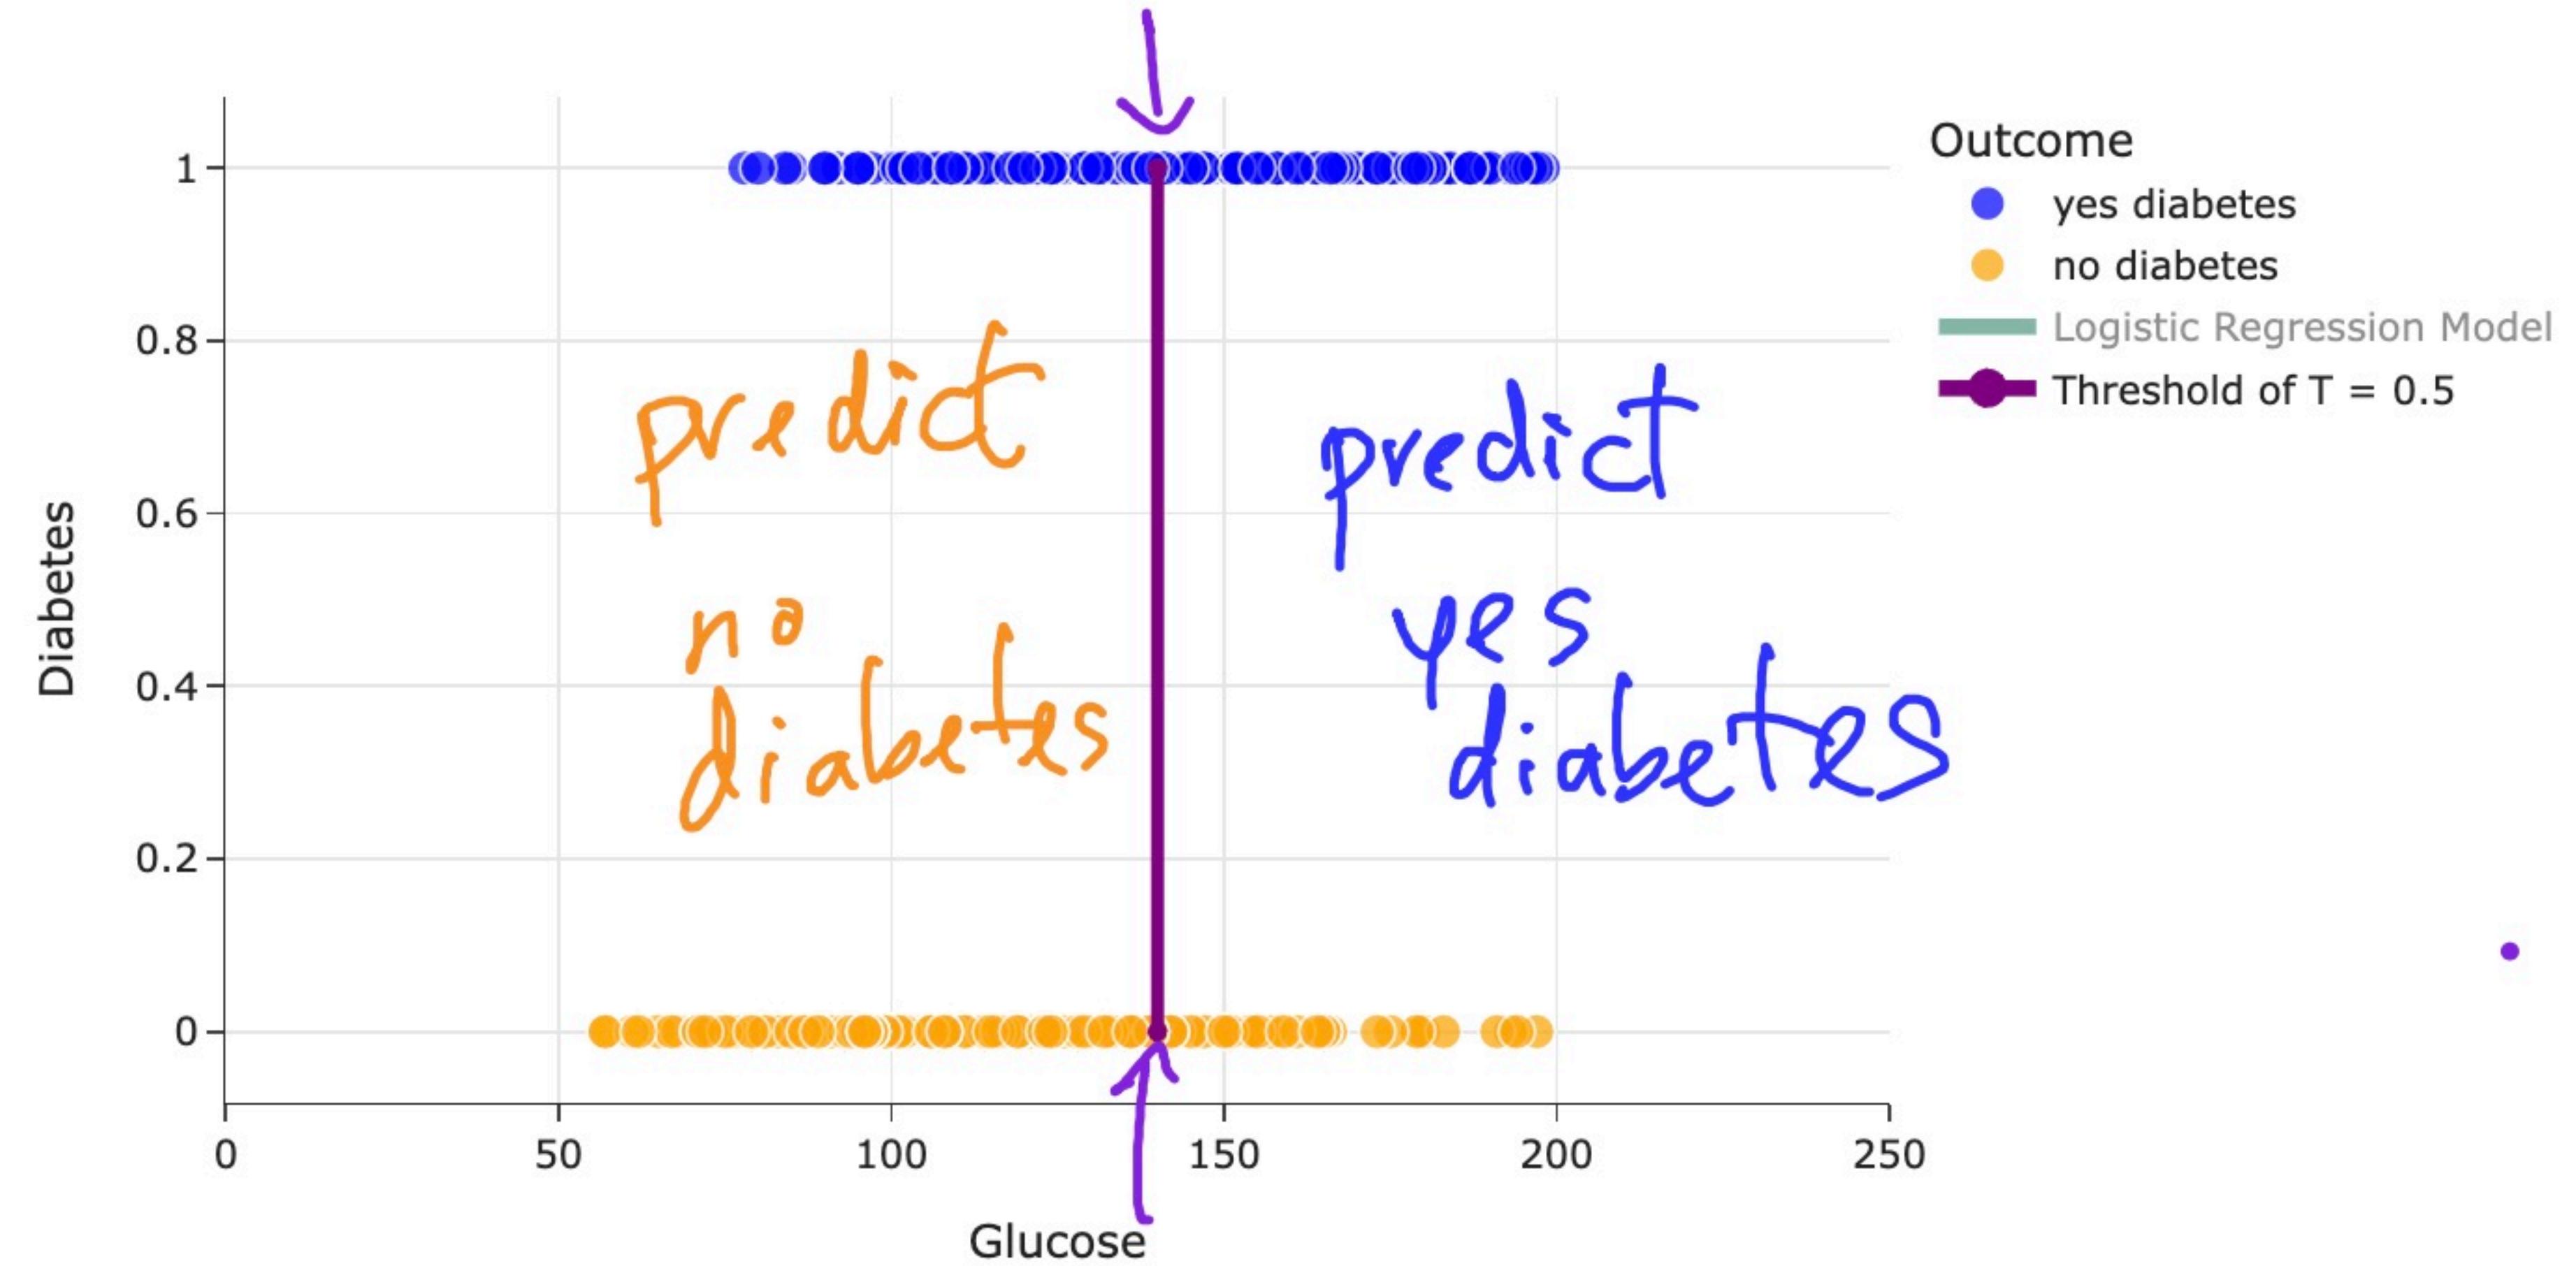
- So, for a given value of C , it minimizes:

$$R_{\text{ce-reg}}(\vec{w}) = -\frac{1}{n} \sum_{i=1}^n [y_i \log(\sigma(\vec{w} \cdot \text{Aug}(\vec{x}_i))) + (1 - y_i) \log(1 - \sigma(\vec{w} \cdot \text{Aug}(\vec{x}_i)))] + \frac{1}{C} \sum_{j=1}^d w_j^2$$

average CE loss

by default, used L_2 reg.

```
In [29]: 1 util.show_one_feature_plot_with_logistic_and_x_threshold(X_train, y_train, 0.5)
```



- How do we find the exact x -axis position of the decision boundary?

If we can, then we'd be able to predict whether someone has diabetes just by looking at their 'Glucose' value.

- In our single feature model where the feature is 'Glucose', our predicted probabilities are of the form:

$$P(y_i = 1 | \text{Glucose}_i) = \sigma(w_0^* + w_1^* \cdot \text{Glucose}_i)$$

$$\begin{aligned} f^{-1}(f(x)) \\ = x \end{aligned}$$

- Suppose we fix a threshold, T . Then, our **decision boundary** is of the form:

$$\sigma^{-1}(\sigma(w_0^* + w_1^* \cdot \text{Glucose}_T)) = T$$

- If we can invert $\sigma(t)$, then we can re-arrange the above to solve for the 'Glucose' value at the threshold:

$$\text{Glucose}_T = \frac{\sigma^{-1}(T) - w_0^*}{w_1^*}$$

$w_0^* + w_1^* \cdot (\text{Glucose}_T) = \sigma^{-1}(T)$

- In our single feature model where the feature is 'Glucose', our predicted probabilities are of the form:

$$P(y_i = 1 | \text{Glucose}_i) = \sigma(w_0^* + w_1^* \cdot \text{Glucose}_i)$$

- Suppose we fix a threshold, T . Then, our **decision boundary** is of the form:

$$\sigma(w_0^* + w_1^* \cdot \text{Glucose}_T) = T$$

- If we can invert $\sigma(t)$, then we can re-arrange the above to solve for the 'Glucose' value at the threshold:

$$\text{Glucose}_T = \frac{\sigma^{-1}(T) - w_0^*}{w_1^*}$$

$\log\left(\frac{T}{1-T}\right)$

- **Important:** If $p = \sigma(t)$, then $\sigma^{-1}(p) = \log\left(\frac{p}{1-p}\right)$ is the inverse of $\sigma(t)$.

$\sigma^{-1}(p)$ is called the **logit** function.

- Suppose an event occurs with probability p .

- The **odds** of that event are:

$$\text{odds}(p) = \frac{p}{1 - p}$$

- For instance, if there's a $p = \frac{3}{4}$ chance that Michigan wins this week, then the **odds** that Michigan wins this week are:

$$\text{odds}\left(\frac{3}{4}\right) = \frac{\frac{3}{4}}{\frac{1}{4}} = 3$$

- Interpretation: it's 3 times more likely that Michigan wins than loses.

log odds:

- We can interpret $\sigma^{-1}(p) = \log\left(\frac{p}{1-p}\right)$ as the "log odds" of p !

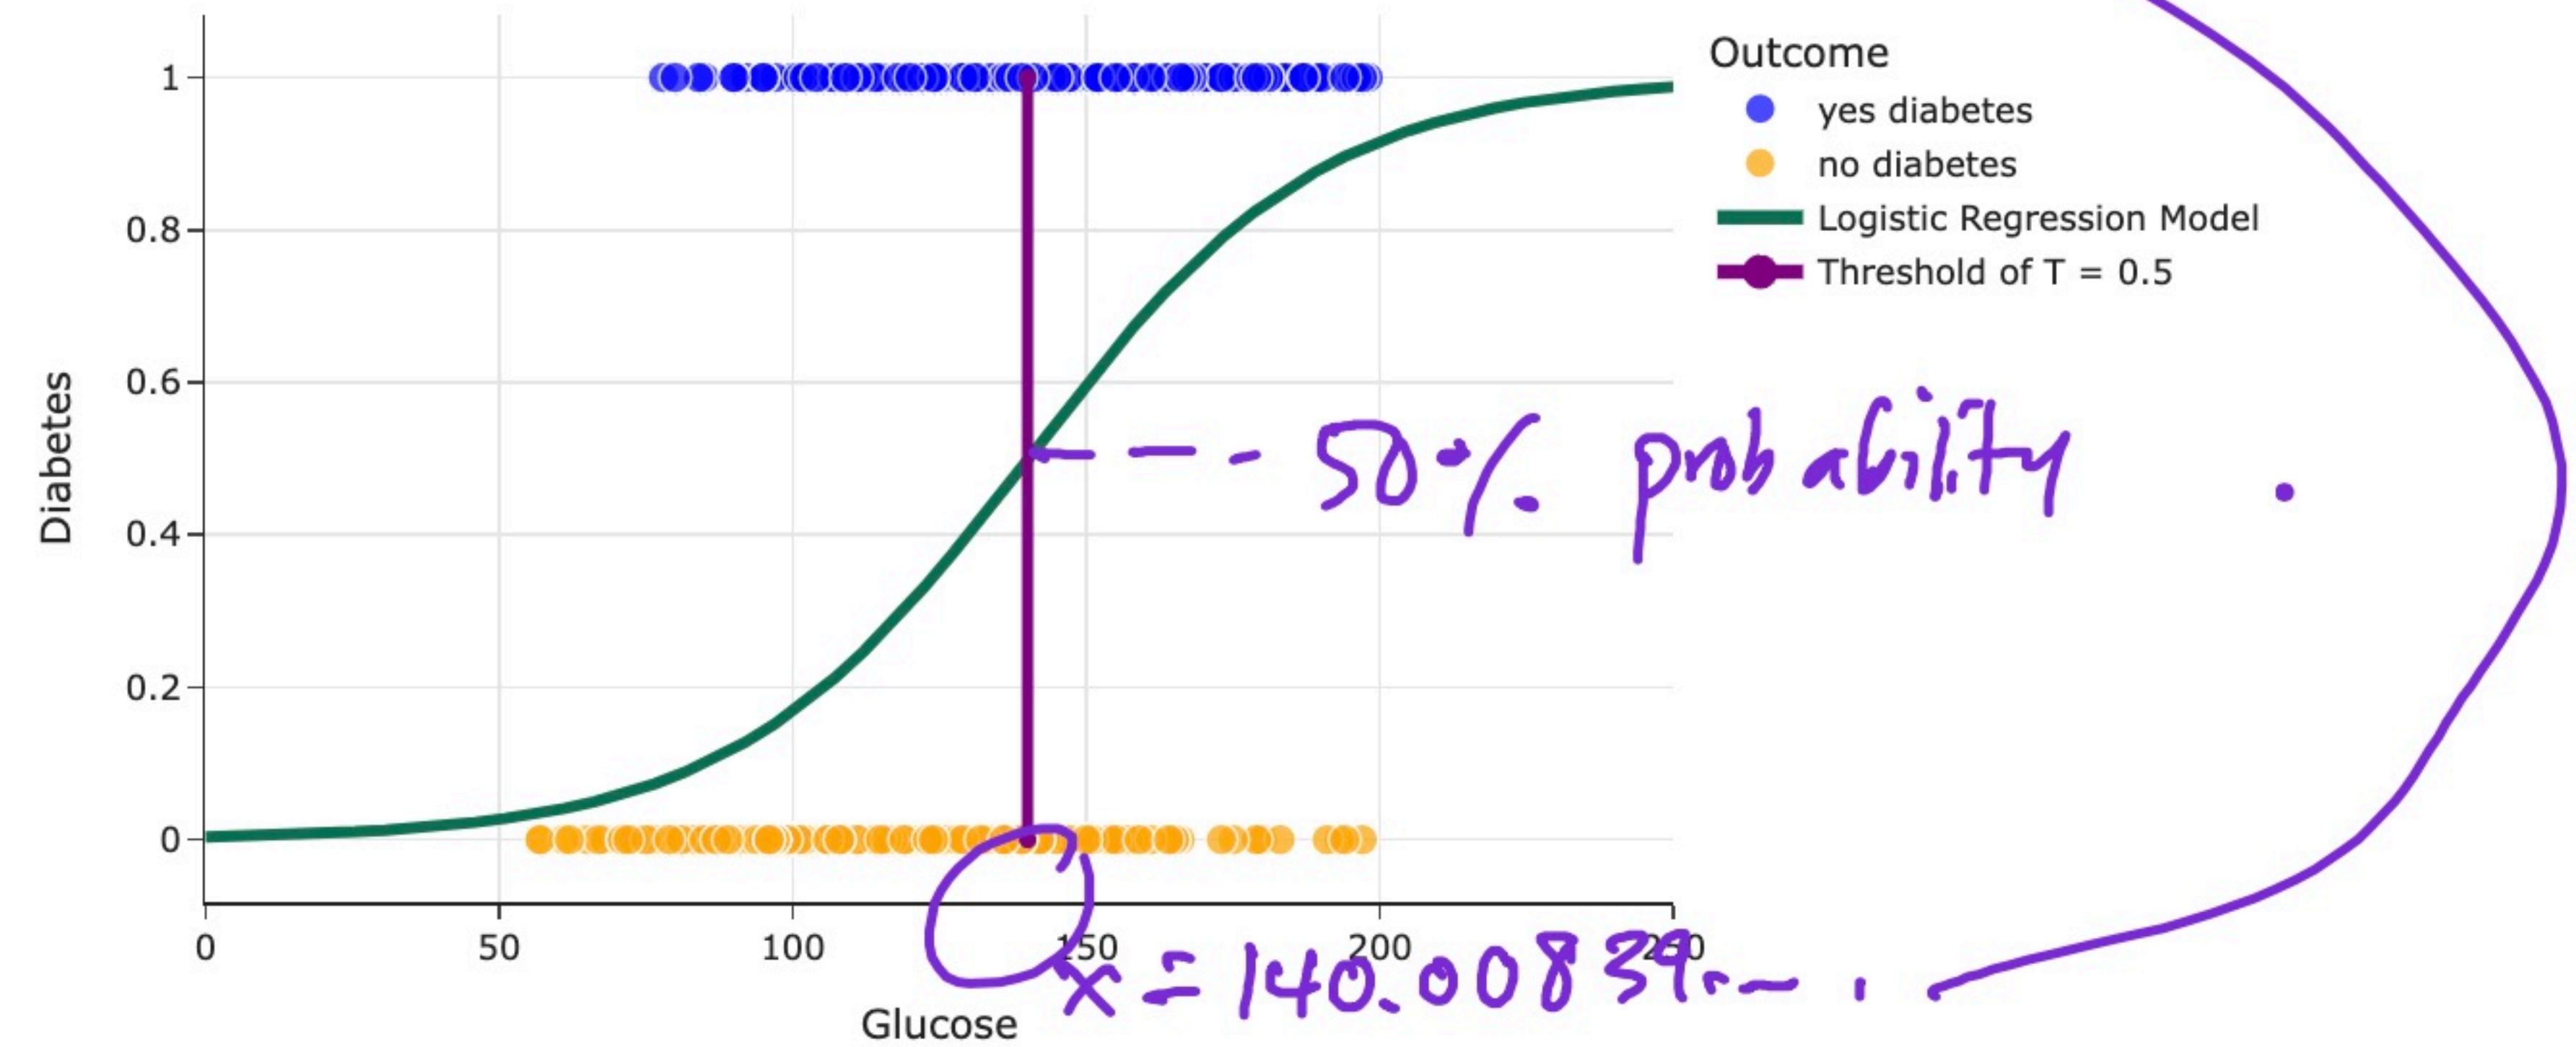
See the reference slides for more details.

```
3 T = 0.5
```

```
4 glucose_threshold = (np.log(T / (1 - T)) - w0_star) / w1_star  
5 glucose_threshold
```

```
Out[30]: 140.0083983057046
```

```
In [31]: 1 util.show_one_feature_plot_with_logistic_and_x_threshold(X_train, y_train, 0.5)
```

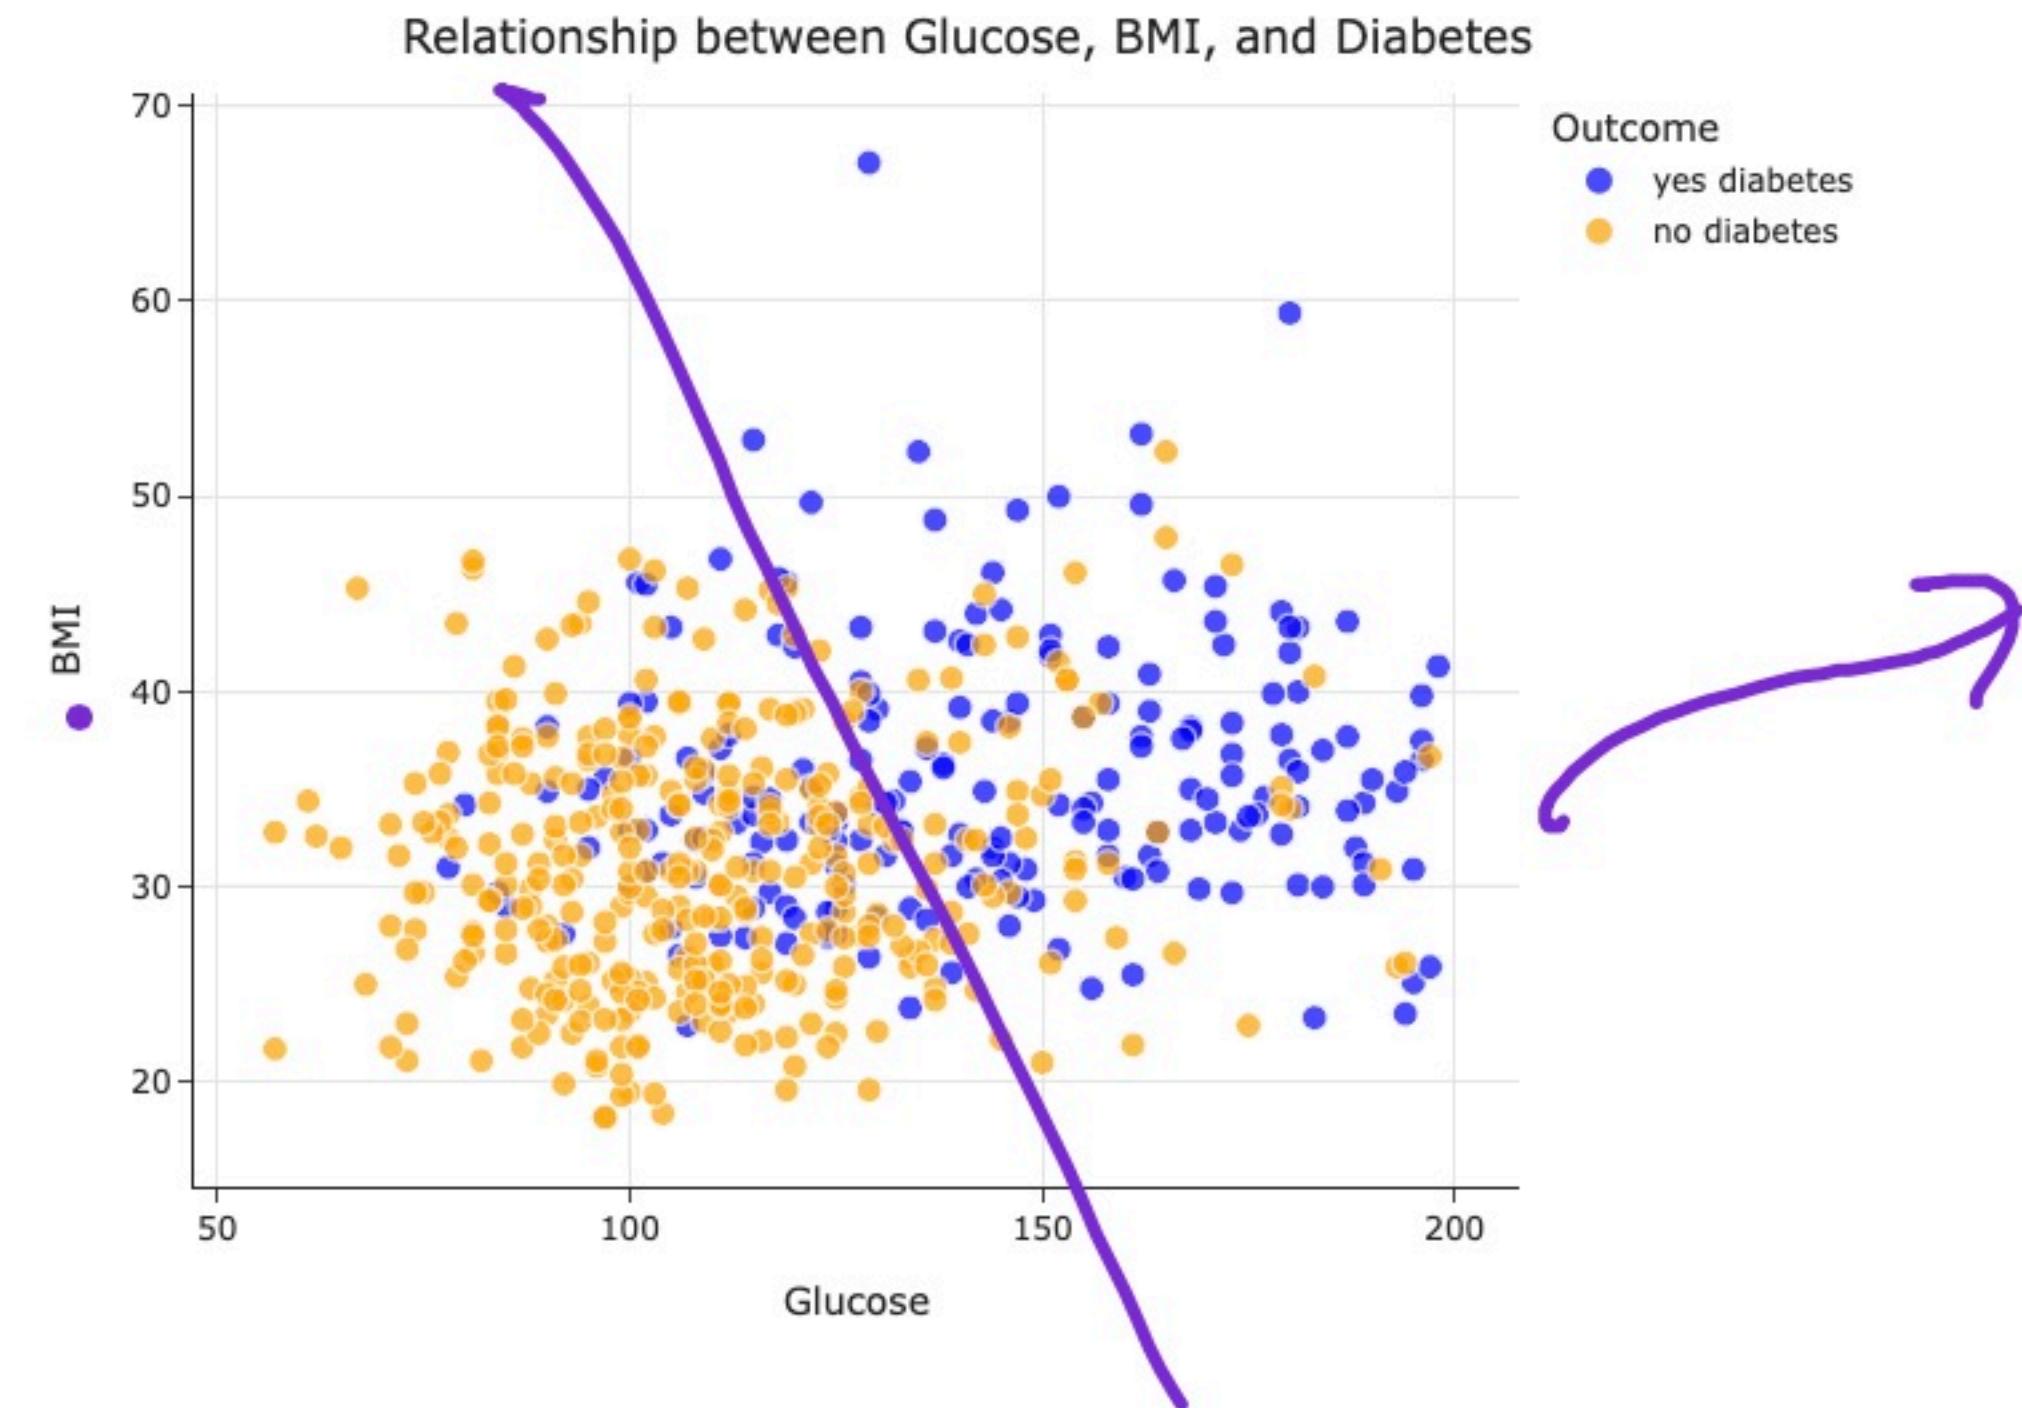


Logistic regression with multiple features

$d=2$ features

- Now, as we did last class, let's use both 'Glucose' and 'BMI' to predict diabetes.

```
In [33]: 1 util.create_base_scatter(X_train, y_train)
```



2D graph

linear decision boundary
IN THE
FEATURE SPACE

• Recall, the logistic regression model

$$P(y_i = 1 | \text{Glucose}_i, \text{BMI}_i) = \sigma(-7.85 + 0.04 \cdot \text{Glucose}_i + 0.08 \cdot \text{BMI}_i)$$

- The graph below shows the predicted probabilities of **class 1 (diabetes)** for different combinations of features.

```
In [36]: 1 util.show_logistic(model_logistic_multiple, X_train, y_train)
```

