


T 0.76

Confusion Matrix with Threshold 0.76



$TN + FP$
is constant
(362)

$FN + TP$
is constant
(202)

Precision vs. threshold


- Precision is defined as:

$$\text{precision} = \frac{TP}{\# \text{ predicted positive}} = \frac{TP}{TP + FP}$$


Here, a false positive (FP) is when we predict that someone has diabetes when they do not.

- How does the model's training **precision** change as the threshold changes?

```
In [ ]: 1 util.plot_vs_threshold(X_train, y_train, 'Precision')
```





In [22]: `1 util.plot_vs_threshold(X_train, y_train, 'Precision')`



rare case
can go down if, e.g.
$$\frac{3}{3+1} \rightarrow \frac{2}{2+1}$$

$$3/4 \quad \quad 2/3$$

- If the "bar" is higher to predict 1, then we will have fewer positives in general, and thus fewer false positives.

- As the **threshold increases** , the denominator in precision $= \frac{TP}{TP+FP}$ will decrease, and so **precision tends to increase** .

There are some cases where a slightly higher threshold led to a slightly lower precision; why?

$$\frac{TP}{TP+FP}$$

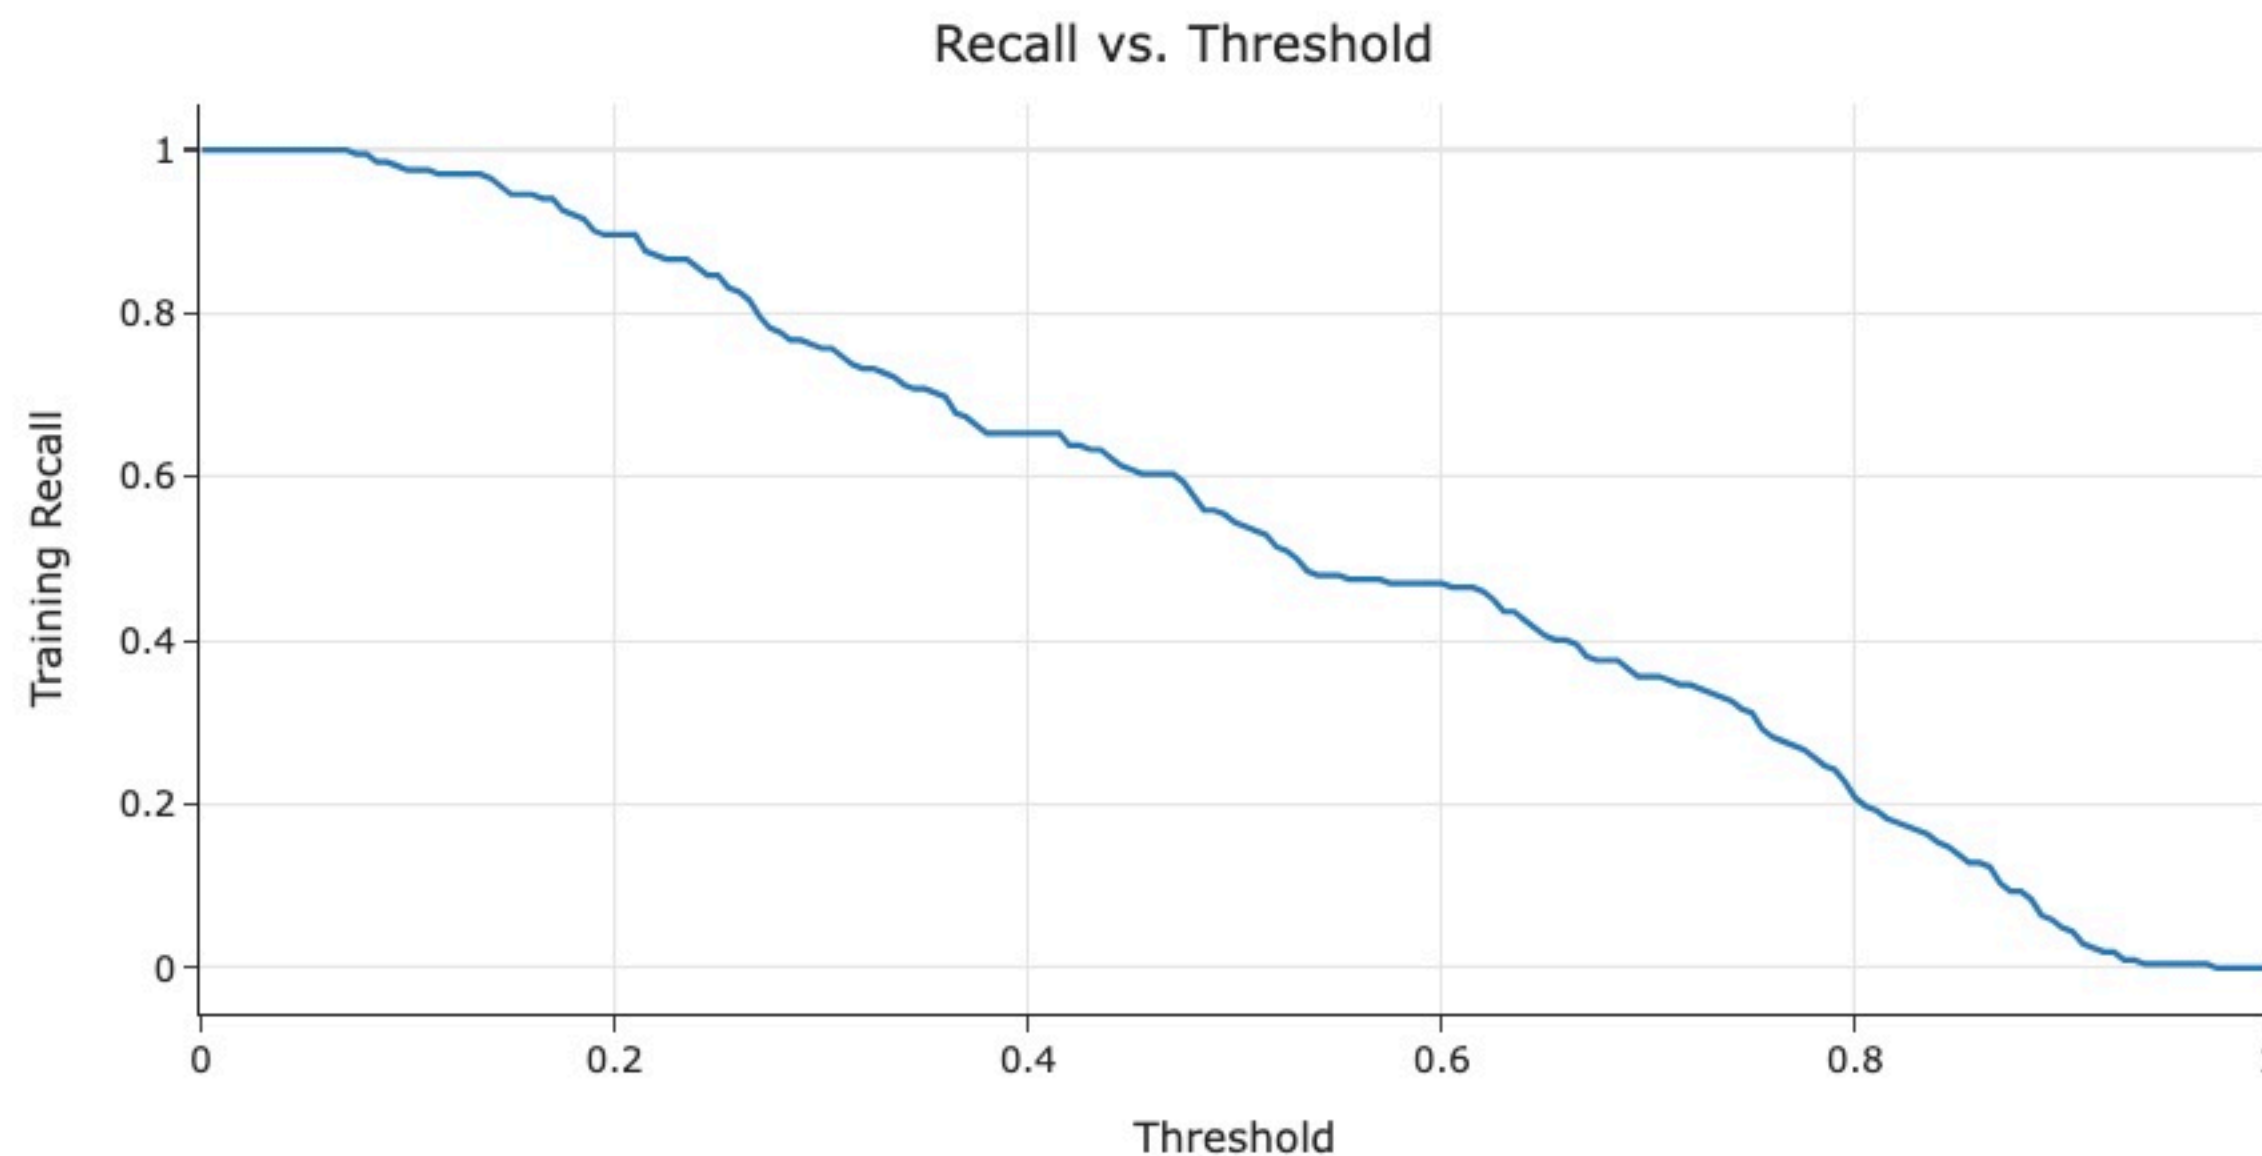
(The entire fraction is circled in blue, with a blue arrow pointing down from the denominator.)

$$\text{recall} = \frac{TP}{\# \text{ actually positive}} = \frac{TP}{TP + FN} = \text{constant.}$$

Here, a false negative (FN) is when we predict that someone does not have diabetes, when they really do.


- How does the model's training **recall** change as the threshold changes?

In [23]: `1 util.plot_vs_threshold(X_train, y_train, 'Recall')`



- We can visualize how precision and recall vary **together**.


In [24]: `1 util.pr_curve(X_train, y_train)`



- The curve above is called a **PR curve**.

- We can visualize how precision and recall vary **together**.

In [24]: `1 util.pr_curve(X_train, y_train)`



pick close to ideal!


ideal

- The curve above is called a **PR curve**.



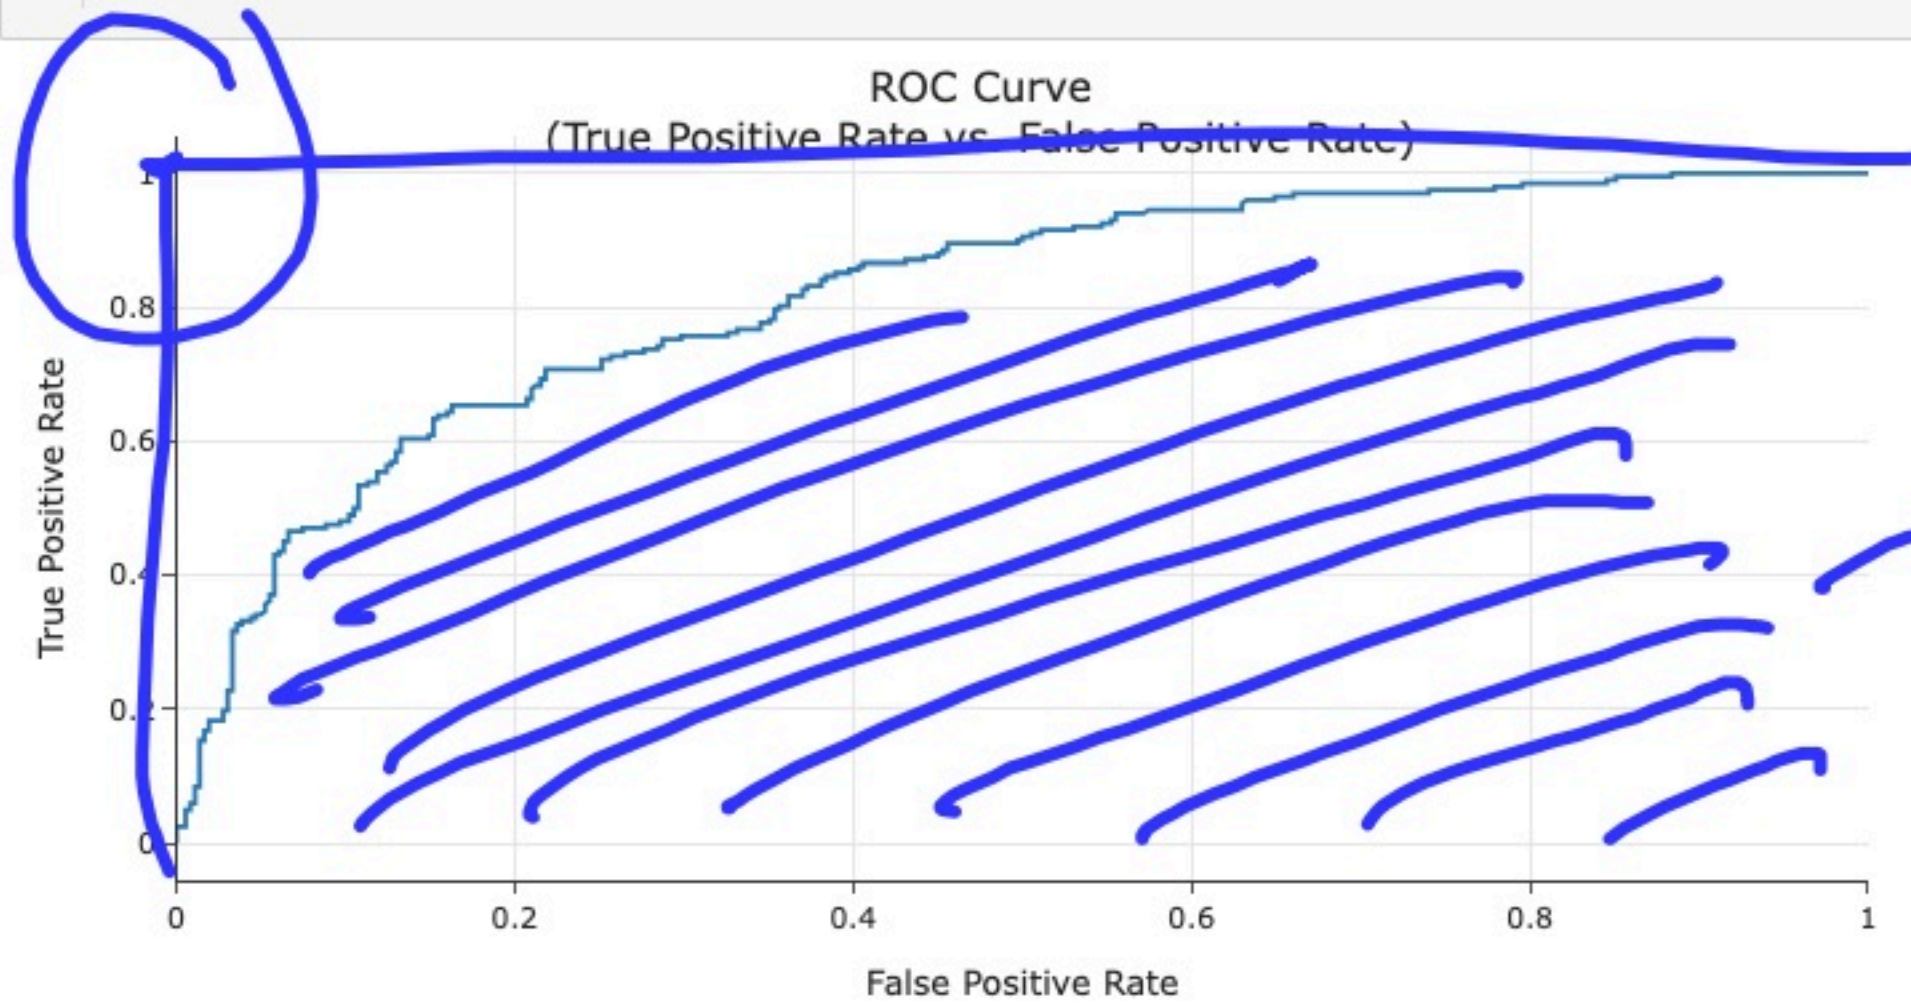
The ROC curve for our classifier looks like:

```
In [25]: 1 util.draw_roc_curve(X_train, y_train)
```



The ROC curve for our classifier looks like:

```
In [25]: 1 util.draw_roc_curve(X_train, y_train)
```



~ AUC : want close to 1.

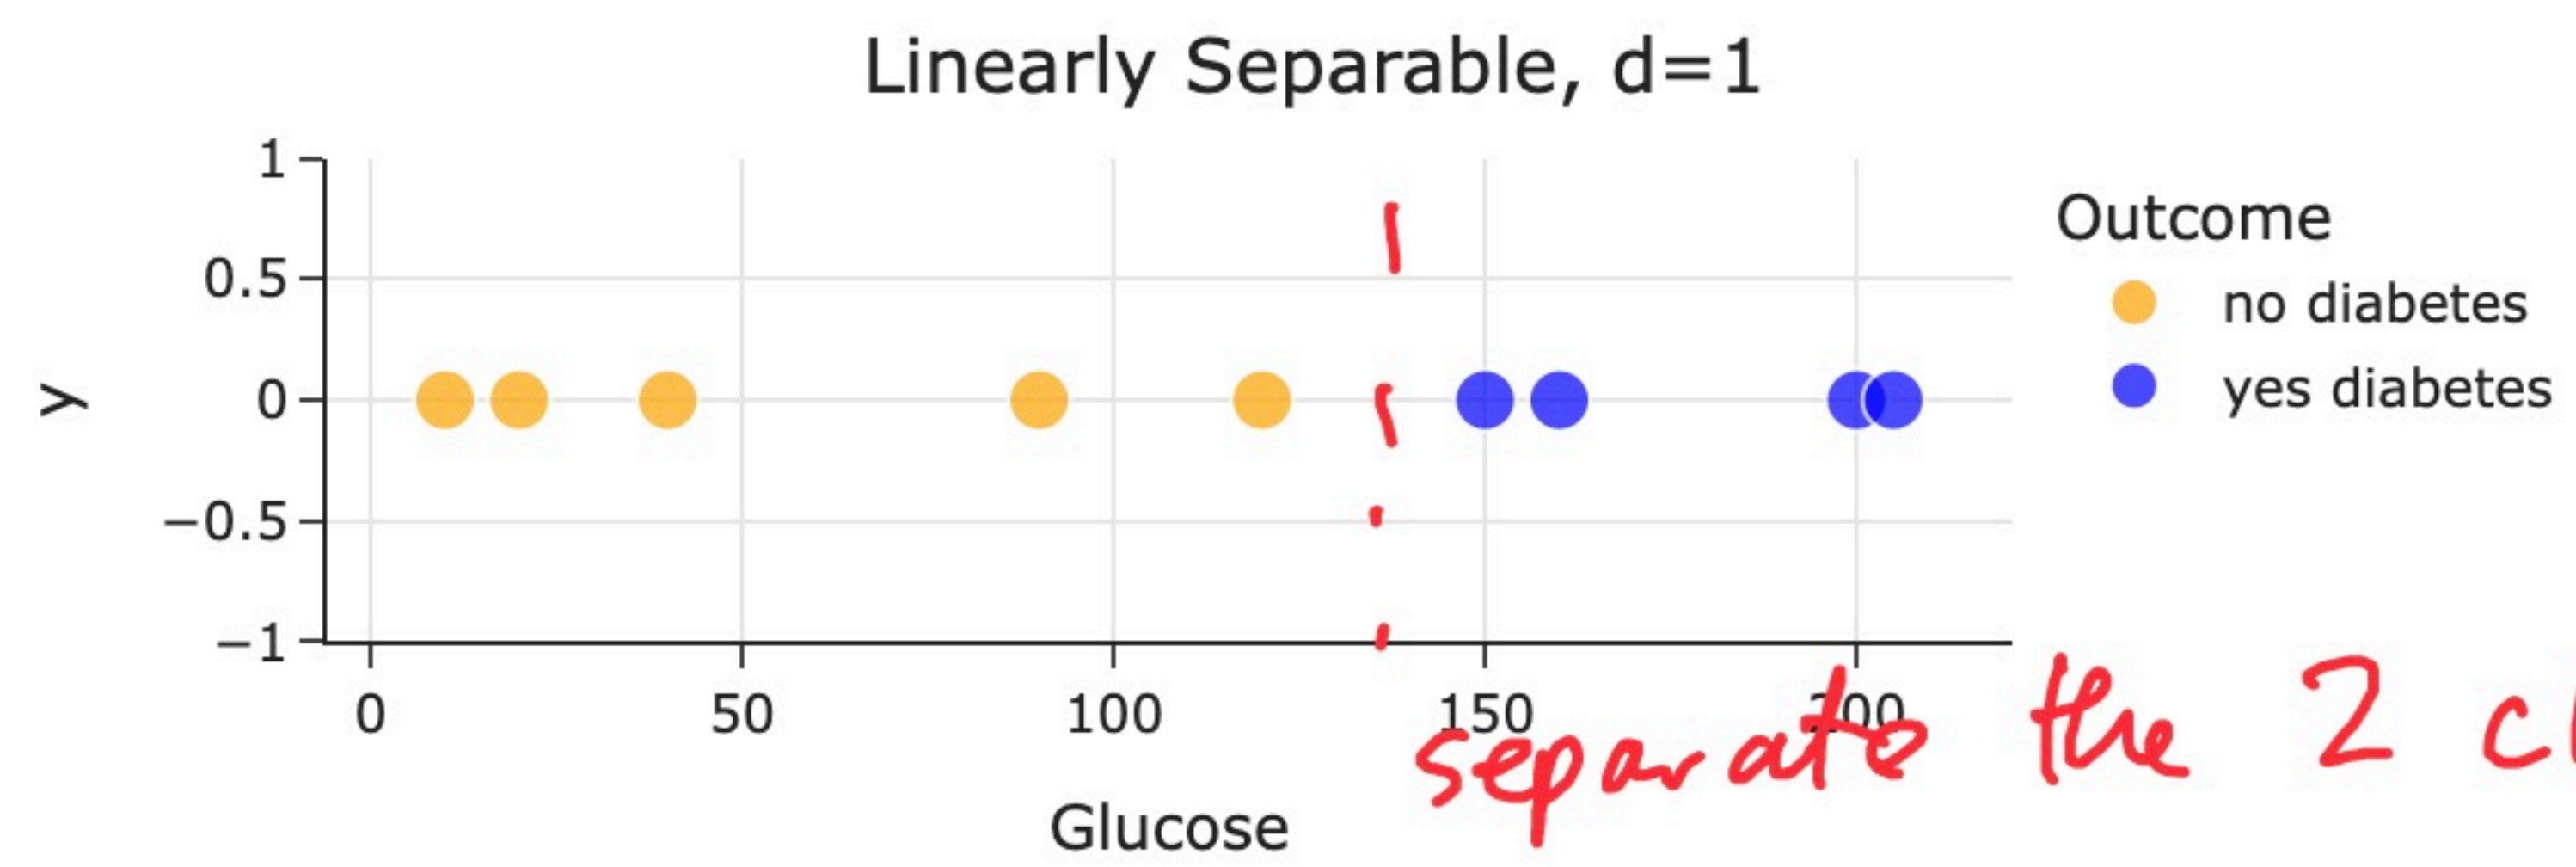
- If we care about TPR and FPR equally, the best threshold is the one whose point is closest to the **top left corner** in the plot above.

Why? The top left corner is where $TPR = 1$ and $FPR = 0$, and we want TPR to be high and FPR to be low.

- A common metric for the quality of a binary classifier is the **area under curve (AUC)** for the ROC curve.


Larger values are better!

In [29]: 1 util.lin_sep_1D()




separate the 2 classes!

In [28]: 1 util.non_lin_sep_1D()



```
32]: 1 util.bad_example_1D()
```


NOT Linearly Separable, $d=1$



not in the feature space!

- Why is the dataset below **not** linearly separable?

In [32]: 1 util.bad_example_1D()

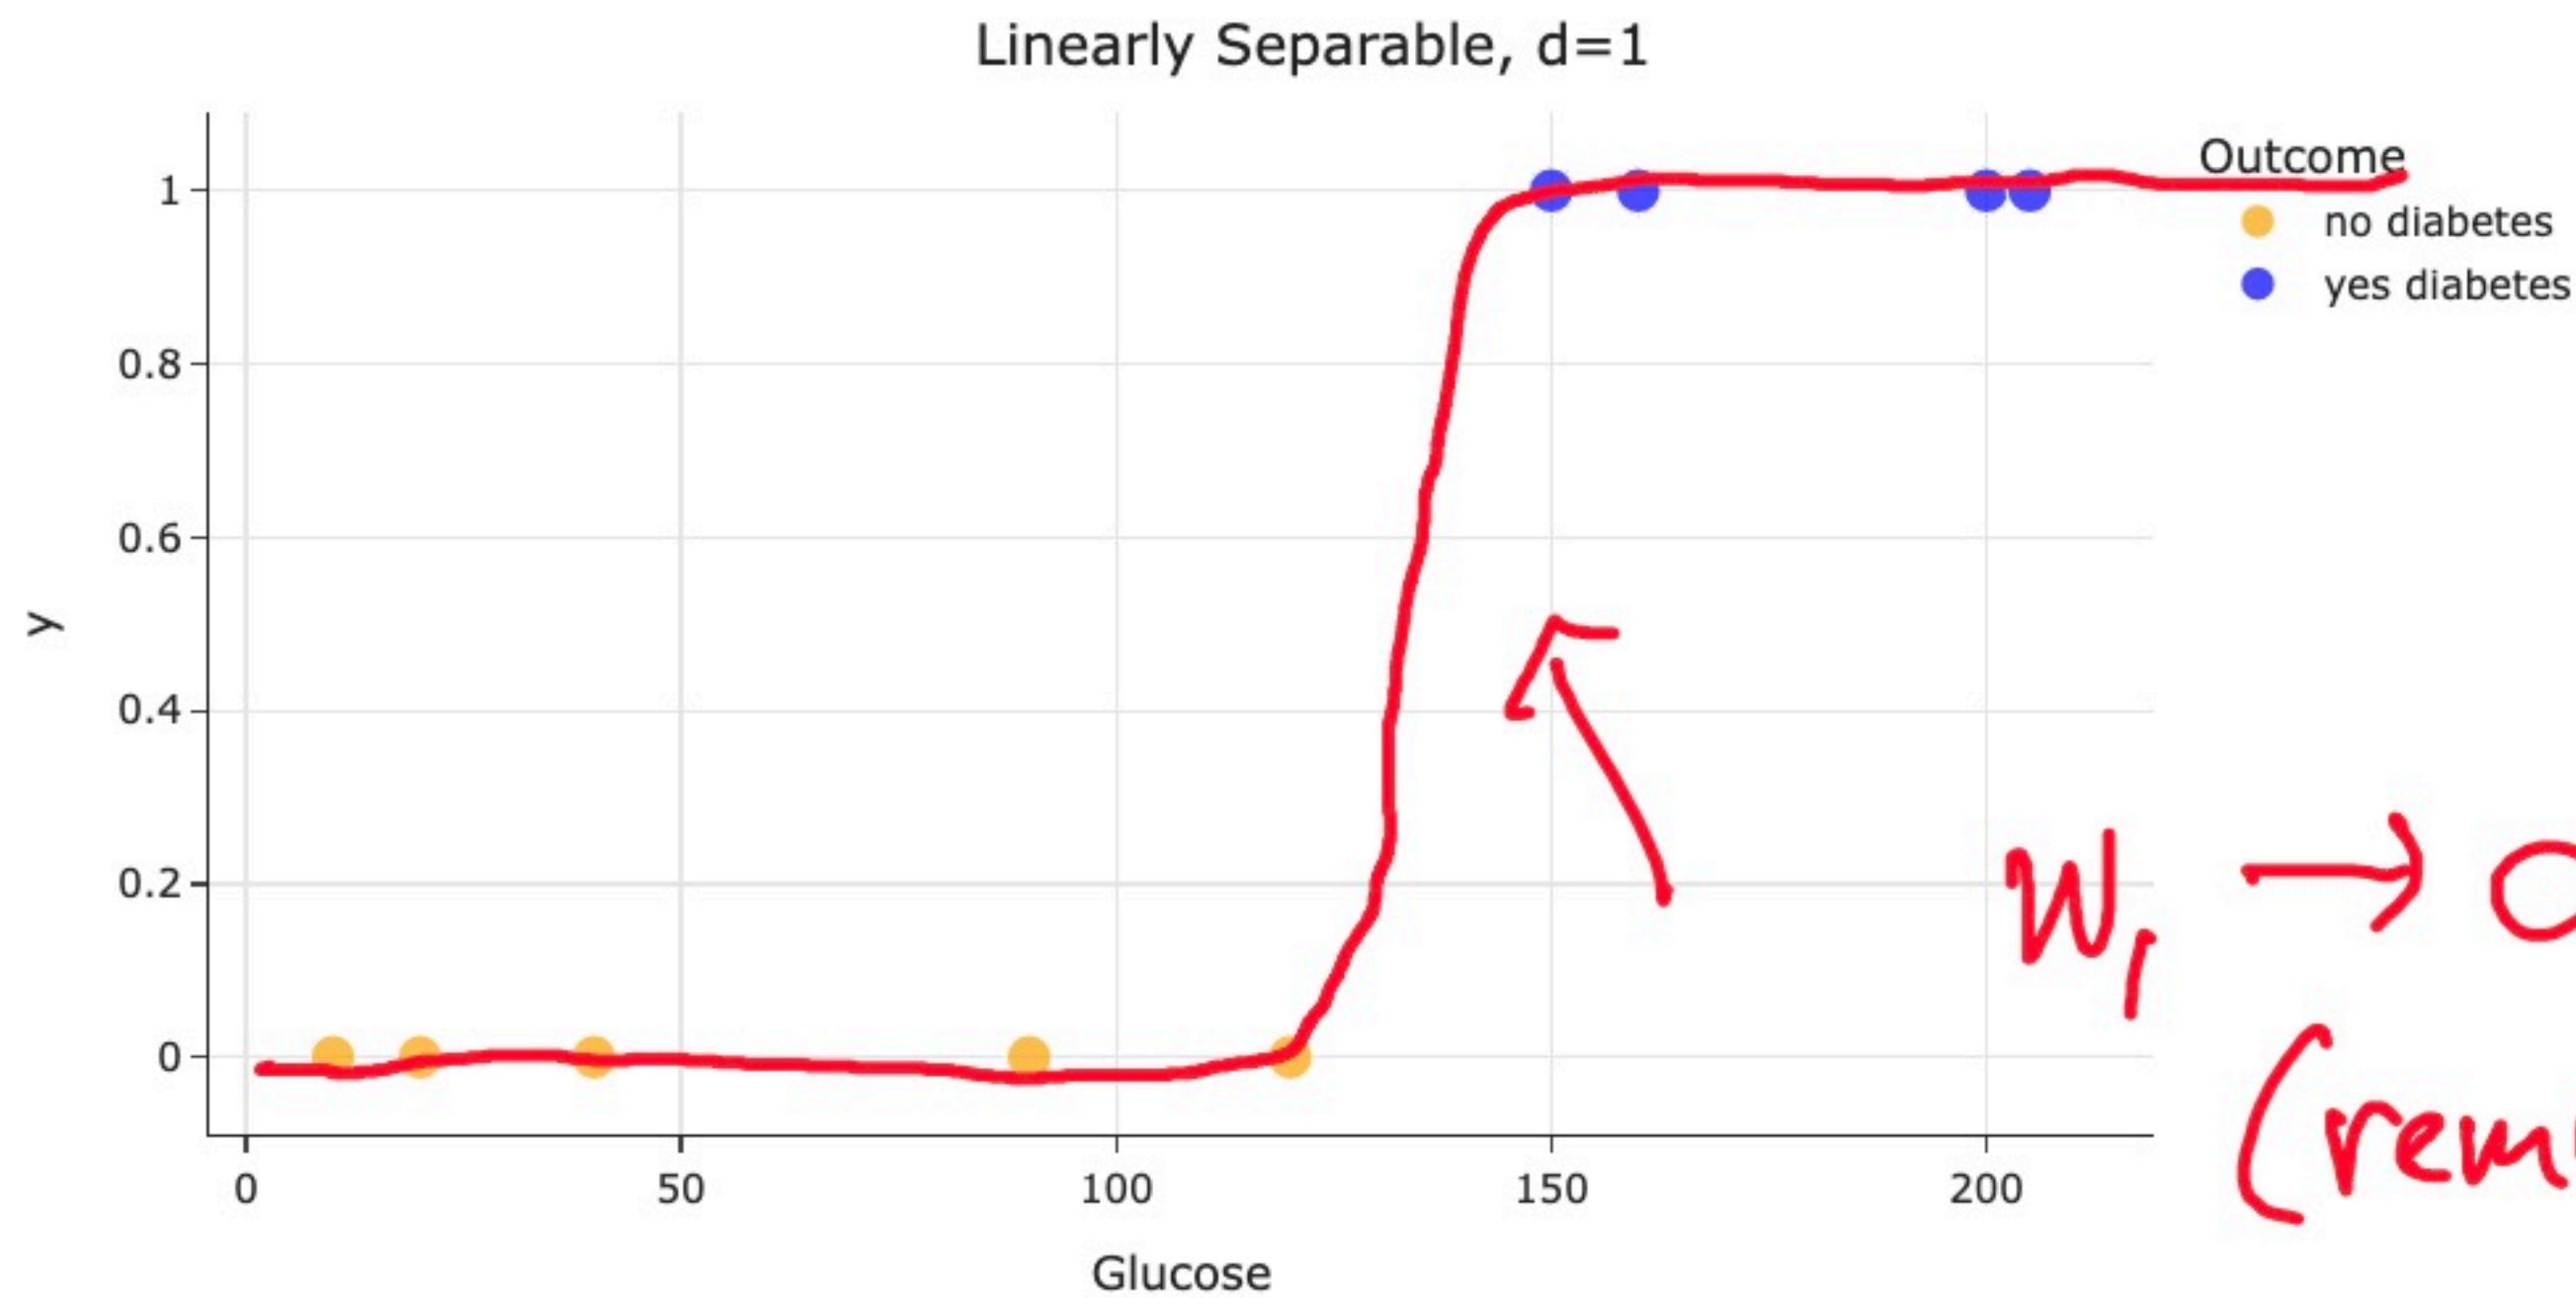


I doesn't exist.

See the annotated slides for more details.

$$P(y_i = 1 | \text{Glucose}_i) = \sigma(w_0 + w_1 \cdot \text{Glucose}_i) = \frac{1}{1 + e^{-(w_0 + w_1 \cdot \text{Glucose}_i)}}$$

```
In [35]: 1 util.lin_sep_1D_elevated()
```



$w_1 \rightarrow \infty$
(remember, w_1 controls steepness)

cover rest
on tuesday →.

Logistic regression for multiclass classification